

Biology 305: Biostatistics

Dr. Jacob C. Cooper & Dr. Melissa Wuellner

2025-01-01

Table of contents

Introduction	8
Class authors	9
1 Installing R	10
1.1 Setup	10
1.1.1 Installing <i>R</i>	10
1.1.2 Installing <i>RStudio</i>	11
1.1.3 Setting up the “native pipe” operator	13
1.1.4 Customizing <i>R</i>	14
2 Intro to R	16
2.1 Creating an <i>RMarkdown</i> document	17
2.1.1 Setup	17
2.1.2 Coding	20
2.1.3 Plotting	25
2.2 Stop touching your mouse	26
2.2.1 Tab complete	27
2.2.2 Help	27
3 Running your first analysis	28
3.1 Create a new document	28
3.2 Working with data	28
3.2.1 Libraries	28
3.2.2 Downloading data	30
3.2.3 Subsetting data	32
3.3 Creating a new column	35
3.4 Your turn!	35
4 Descriptive Statistics	36
4.1 Purposes of descriptive statistics	36
4.2 Sample data and preparation	36
4.3 Downloading the data	36
4.4 Descriptive statistics	37
4.4.1 Notation	37
4.4.2 Mean	38

4.4.3	Range	39
4.4.4	Median	40
4.4.5	Other quartiles and quantiles	42
4.4.6	Mode	44
4.4.7	Variance	47
4.4.8	Standard deviation	49
4.4.9	Standard error	49
4.4.10	Coefficient of variation	50
4.4.11	Outliers	51
4.5	Homework: Descriptive Statistics	52
4.5.1	Homework instructions	52
4.5.2	Data for homework problems	53
5	Diagnosing data visually	57
5.1	The importance of visual inspection	57
5.2	Sample data and preparation	57
5.3	Histograms	57
5.3.1	Histograms on numeric vectors	57
5.3.2	Histograms on frequency counts	59
5.3.3	ggplot histograms	60
5.4	Boxplots	65
5.5	Skewness	66
5.6	Kurtosis	68
5.7	Cumulative frequency plot	70
5.8	Homework: Visualizing data	72
5.8.1	Helpful hint	72
5.8.2	Problem 1	73
5.8.3	Problem 2	74
5.8.4	Problem 3	74
5.8.5	Problem 4	75
5.9	Acknowledgment	75
6	Normality & hypothesis testing	76
6.1	Normal distributions	76
6.1.1	Example in nature	77
6.1.2	Effect of sampling	81
6.2	Hypothesis testing	83
6.2.1	Critical Values - α	84
6.2.2	Introduction to p values	84
6.2.3	Calculating a Z score	89
6.2.4	Calculating the p value	90
6.2.5	Workthrough Example	91
6.3	Confidence Intervals	97

6.4	Homework: Normal distributions	98
6.4.1	Z to p	98
6.4.2	P to z	99
6.4.3	Confidence intervals	99
6.4.4	Italian climate - spring	99
6.4.5	Italian climate - Summer	100
7	Transforming data	101
7.1	Importance of normality Data	101
7.2	Testing if data are normal	102
7.2.1	Histograms	102
7.2.2	QQ Plots	103
7.2.3	Shapiro-Wilk test	105
7.3	Transforming data	106
7.3.1	Log transformation example	106
7.4	Homework: Transforming data	108
7.4.1	Analyzing data	109
7.4.2	Transforming Data	109
7.4.3	Submitting assignment	109
8	Probability distributions	110
8.1	Probability distributions	110
8.2	Review: Basic statistics	110
8.3	Binomial distribution	111
8.3.1	Binomial examples	111
8.3.2	Binomial exact tests	114
8.4	Poisson distribution	116
8.4.1	Poisson example	116
8.4.2	Poisson test	117
8.5	Cumulative Probabilities	117
8.5.1	Binomial cumulative	118
8.5.2	Poisson cumulative probability	119
8.6	Homework	119
8.6.1	Problem 1	120
8.6.2	Problem 2	120
8.6.3	Problem 3	120
8.6.4	Problem 4	120
8.6.5	Problem 5	120
8.6.6	Problem 6	120
8.6.7	Problem 7	120
8.6.8	Problem 8	121
8.6.9	Problem 9	121

9	2 (Chi-squared) tests	122
9.1	χ^2 -squared distribution	122
9.1.1	Calculating the test statistic	123
9.1.2	χ^2 goodness-of-fit test	123
9.1.3	χ^2 test of independence	126
9.2	Fisher's exact test	131
9.2.1	Chapter 7	132
10	Means testing	133
10.1	Introduction	133
10.1.1	Parametric and Non-parametric tests	133
10.1.2	A little history	133
10.2	Dataset	134
10.3	Parametric tests	134
10.3.1	t -distribution	134
10.3.2	t -tests	137
10.3.3	One-sample t -tests	138
10.3.4	Two-sample t -tests	139
10.3.5	Paired t -tests	140
10.4	Non-parametric tests	141
10.4.1	Wilcoxon tests	141
10.4.2	Mann-Whitney U	142
10.4.3	Wilcoxon signed rank test	142
10.5	Confidence intervals	143
10.6	Homework: Chapter 9	143
10.7	Homework: Chapter 10	143
11	ANOVA: Part 1	144
11.1	Introduction	144
11.2	ANOVA: By hand	147
11.2.1	Calculate group means and Grand Mean	150
11.2.2	Total sum of squares	151
11.2.3	Within-group sum of squares	152
11.2.4	Among-group sum of squares	153
11.2.5	Calculate degrees of freedom	153
11.2.6	Calculate mean squares	154
11.2.7	Get F statistic	154
11.2.8	Get p value	154
11.3	ANOVA: By R	155
11.4	Post-hoc Tukey test	155
11.4.1	Tukey test by hand	155
11.4.2	Tukey test in R	157
11.5	Plotting our ANOVA results	160

11.6	Kruskal-Wallis tests	163
11.6.1	By hand	166
11.6.2	Using <i>R</i>	168
11.7	Homework: Chapter 11	169
12	ANOVA: Part 2	170
12.1	Two-way ANOVA	170
12.2	Designs	171
12.2.1	Randomized block design	171
12.2.2	Repeated measures	176
12.2.3	Factorial ANOVA	181
12.2.4	ANOVA with interaction	181
12.3	Friedman's test	185
12.3.1	Using <i>R</i>	185
12.4	Homework: Chapter 12	188
13	Correlation & regression	189
13.1	Introduction	189
13.2	Correlation	189
13.2.1	Pearson's	192
13.2.2	Spearman's	194
13.2.3	Other non-parametric methods	194
13.3	Regression	195
13.3.1	Parametric	195
13.3.2	Non-parametric	198
13.4	Homework	198
13.4.1	Chapter 13	198
13.4.2	Chapter 14	198
14	Multivariate methods	199
14.1	Multivariate analyses	199
14.2	Principal components	199
14.3	Clustering	203
14.3.1	K-means clustering	203
14.3.2	Hierarchical clustering	203
14.4	Homework	204
14.4.1	Disentangling microbiomes	204
15	Pick the test	205
15.1	Picking the test	205
15.2	Exceptions	205
15.3	Overview of picking the test	205

15.4	Another method - checklist	208
15.4.1	Explanatory Variable	208
15.4.2	Continuous explanatory variable	208
15.4.3	Discrete explanatory variable	209
16	Conclusions	211
16.1	Parting thoughts	211
16.2	211
17	Functions & Glossary	212
17.1	Common Commands	212
17.1.1	The basics	212
17.1.2	Code basics	212
17.1.3	Code formatting	213
17.1.4	Boolean logic	214
17.2	Basic statistics	215
17.3	Null hypotheses	217
17.3.1	χ^2 tests	217
17.3.2	One-sample tests	217
17.3.3	Multi-sample tests	217
17.3.4	Correlation	218
17.3.5	Regression	218
17.4	Custom functions from class and elsewhere	218
17.4.1	Normal distributions	220
17.4.2	ANOVA	220
	References	228

Introduction

Welcome to Biology 305 at the University of Nebraska at Kearney!

In this class, you will learn:

1. The basics of study design, the importance of understanding your research situation before embarking on a full study, and practice creating research frameworks based on different scenarios.
2. The basics of data analysis, including understanding what kind of variables are being collected, why understanding variable types are important, and basic tests to understand univariate distributions.
3. Basic multivariate statistics, including ANOVA, correlation, and regression, for comparing multiple different groups.
4. The basics of coding and working in *R* for performing statistical analyses.

This site will help you navigate different homework assignments to perform the necessary *R* tests.

Welcome to class!

Class authors

- [Dr. Melissa Wuellner](#)

Primary lecture designer, author of many of the example problems.

- [Dr. Jacob C. Cooper](#)

Primary website designer, primary *R* coding for site

- [Caleb Rother](#)

Teaching assistant, website designer, works with homework and coding

1 Installing R

In this class, we will be using *R* to perform statistical analyses. *R* is a free software program designed for use in a myriad of statistical and computational scenarios. It can handle extremely large datasets, can handle spatial data, and has wrappers for compatibility with *Python*, *Bash*, and other programs (even *Java*!).

1.1 Setup

First, we need to download *R* onto your machine. We are also going to download *RStudio* to assist with creating *R* scripts and documents.

1.1.1 Installing *R*

First, navigate to the [R download and install page](#). Download the appropriate version for your operating system (Windows, Mac, or Linux). **Note** that coding will be formatted slightly different for Windows than for other operating systems. If you have a Chromebook, you will have to [follow the online instructions for installing both programs on Chrome](#).

Follow the installation steps for *R*, and verify that the installation was successful by searching for *R* on your machine. You should be presented with a coding window that looks like the following, but may not be an exact match:

```
R version 4.4.1 (2024-06-14) -- "Race for Your Life"
Copyright (C) 2024 The R Foundation for Statistical Computing
Platform: aarch64-apple-darwin20
```

```
R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.
```

```
Natural language support but running in an English locale
```

```
R is a collaborative project with many contributors.
Type 'contributors()' for more information and
```

'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or

'help.start()' for an HTML browser interface to help.

Type 'q()' to quit R.

>

If that screen appears, congratulations! *R* is properly installed. If the install was not successful, please talk to your instructor and teaching assistant(s) for help with installation.

1.1.2 Installing *RStudio*

RStudio is a GUI (graphics user interface) that helps make *R* easier to use. Furthermore, it allows you to create documents in *R*, including websites (such as this one), PDFs, and even presentations. This can greatly streamline the research pipeline and help you publish your results and associated code in a quick and efficient fashion.

Head over to the [RStudio download website](#) and download “*RStudio* Desktop”, which is free. Be sure to pick the correct version for your machine.

Open *RStudio* on your machine. You should be presented with something like the following:

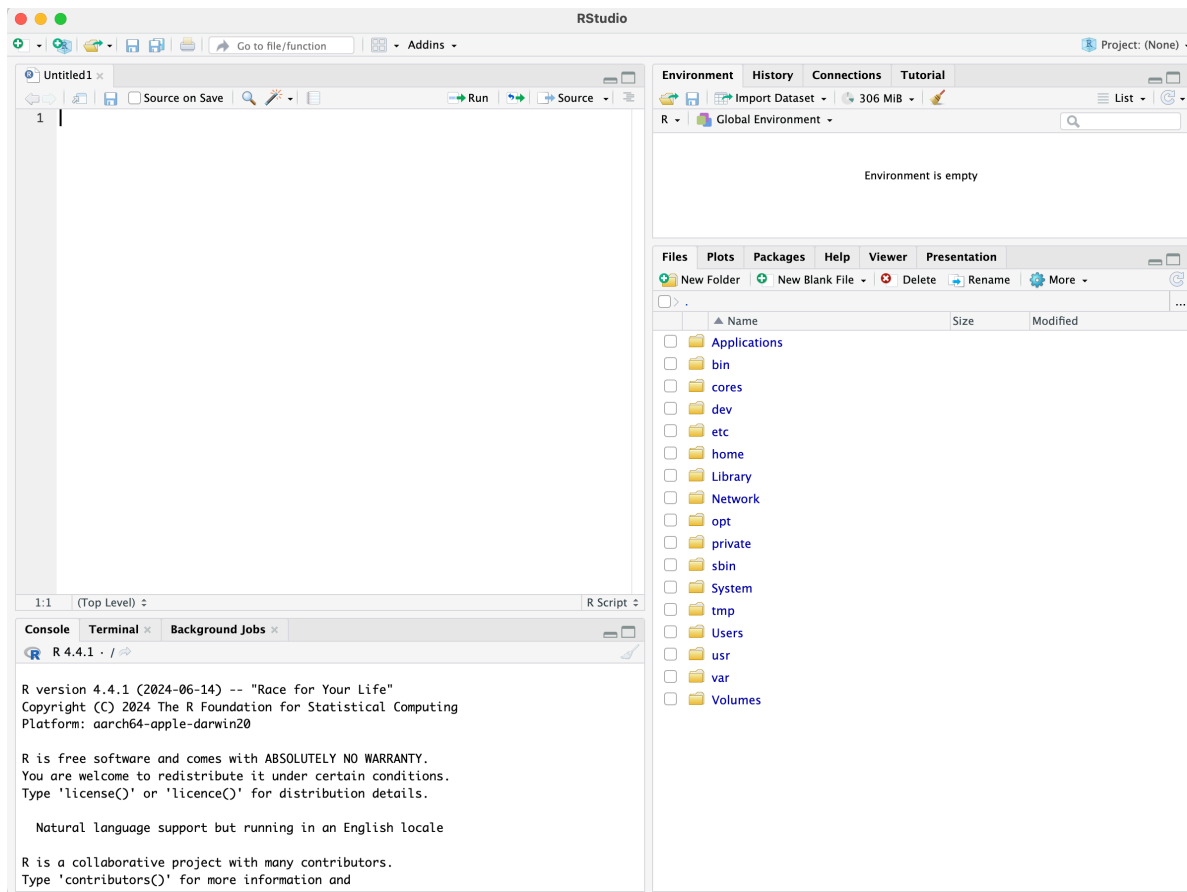


Figure 1.1: *RStudio* start window. Note that the screen is split into four different quadrants. Top left: *R* documents; bottom left: *R* program; top right: environment window; bottom right: plots, help, and directories.

In *RStudio*, the top left window is always going to be our coding window. This is where we will type all of our code and create our documents. In the bottom left we will see *R* executing the code. This will show what the computer is “thinking” and will help us spot any potential issues. The top right window is the “environment”, which shows what variables and datasets are stored within the computers’ memory. (It can also show some other things, but we aren’t concerned with that at this point). The bottom right window is the “display” window. This is where plots and help windows will appear if they don’t appear in the document (top left) window itself.

1.1.3 Setting up the “native pipe” operator

There are two ways to type commands in *R*: in a “nested” fashion and in a “piped” fashion. We will be using the “piped” fashion for ease of use.

In *RStudio*, click on the tabs at the top of the screen to go to Tools > Global Options, and in the pop up screen select Code. On this screen, you should see a tick box for “Use native pipe operator”. Make sure this box is checked. Now, we can use CTRL + SHIFT + M (a.k.a., ^ + SHIFT + M on Mac) to insert the “pipe” operator |>.




Figure 1.2: Screenshot showing where to activate the native pipe operator within RStudio’s settings (i.e., options).

1.1.3.1 Using |> pipes

Let’s say that we have a random series of 100 numbers from a random normal distribution, which we can get by copying and running the following in the lower left coding window:

```
# make repeatable  
set.seed(42)
```

```
# get 100 values from normal
x <- rnorm(100)
```

We want to get the mean of these values *and* round these values to two decimal places for our final answer. To do this in a *nested* fashion, we would do the following:

```
# round mean of x
round(mean(x), 2)
```

```
[1] 0.03
```

Note that we call this the “nested” method because one command (`mean`) is located *inside* another command (`round`). This is fine for short series of commands, but can get confusing when too many commands are combined. Thus, we can use the `|>` pipe method instead.

The pipe operator takes whatever came previously and puts it through the next command. For example, `mean(x)` could also be written as `x |> mean()`, where `x` would be placed inside of `mean()`.

```
# take x
x |>
  # get mean
  mean() |>
  # round answer
  round(2)
```

```
[1] 0.03
```

As you can see, the above breaks things down into a more step-by-step fashion, and makes code easier to follow. We will be using this extensively in this course.

1.1.4 Customizing R

RStudio’s interface can be customized within the global options under “appearance”. You can explore different fonts, different font sizes, and different color schemes. Please choose one that makes it easy for you to:

1. Has a color scheme that makes the difference between different commands, numbers, and objects easy to see and is easily readable *for you*.
2. Has a font that is a type and size that is easily readable *for you*.




Figure 1.3: Visual settings on Dr. Cooper’s machine. Note that the colors make the numbers and commands stand out, and the font and color scheme are easily legible.

CONGRATS! You now have *R* installed on your computer.

2 Intro to *R*

Now that we have installed *RStudio*, we will need to create an *RMarkdown* document that we can use for answering questions and turning in assignments. These documents will also be used for appendices and other materials in future publications you create - *RStudio* can easily knit documents into `.pdf`, `.docx` (i.e., “Word” documents), and `.html`. In this class, we will use almost exclusively `.html` documents, as these are the most compatible with *Canvas*.

- **Quick note:** Mac keyboards differ from Windows keyboards in that they have more keys; as such, shortcuts will differ!
- For both operating systems, I will list keys that have to be pushed together with a “+” between them. (If you need to push the “+” key, I will offset it like + so; for example, pushing CONTROL and PLUS on Windows at the same time makes a window’s font bigger; I would write this as CTRL + +).
- If you need to hit a single key that differs between operating systems, I will always list the Windows/Linux key first and the Mac key second; for example, Enter/return for inserting a new line.

2.0.0.1 Windows/Linux keys

- CTRL: Control
- ALT: Alt

2.0.0.2 Mac keys

- ⌘ : Command
- ⌥ : Option
- ⌃ : Control

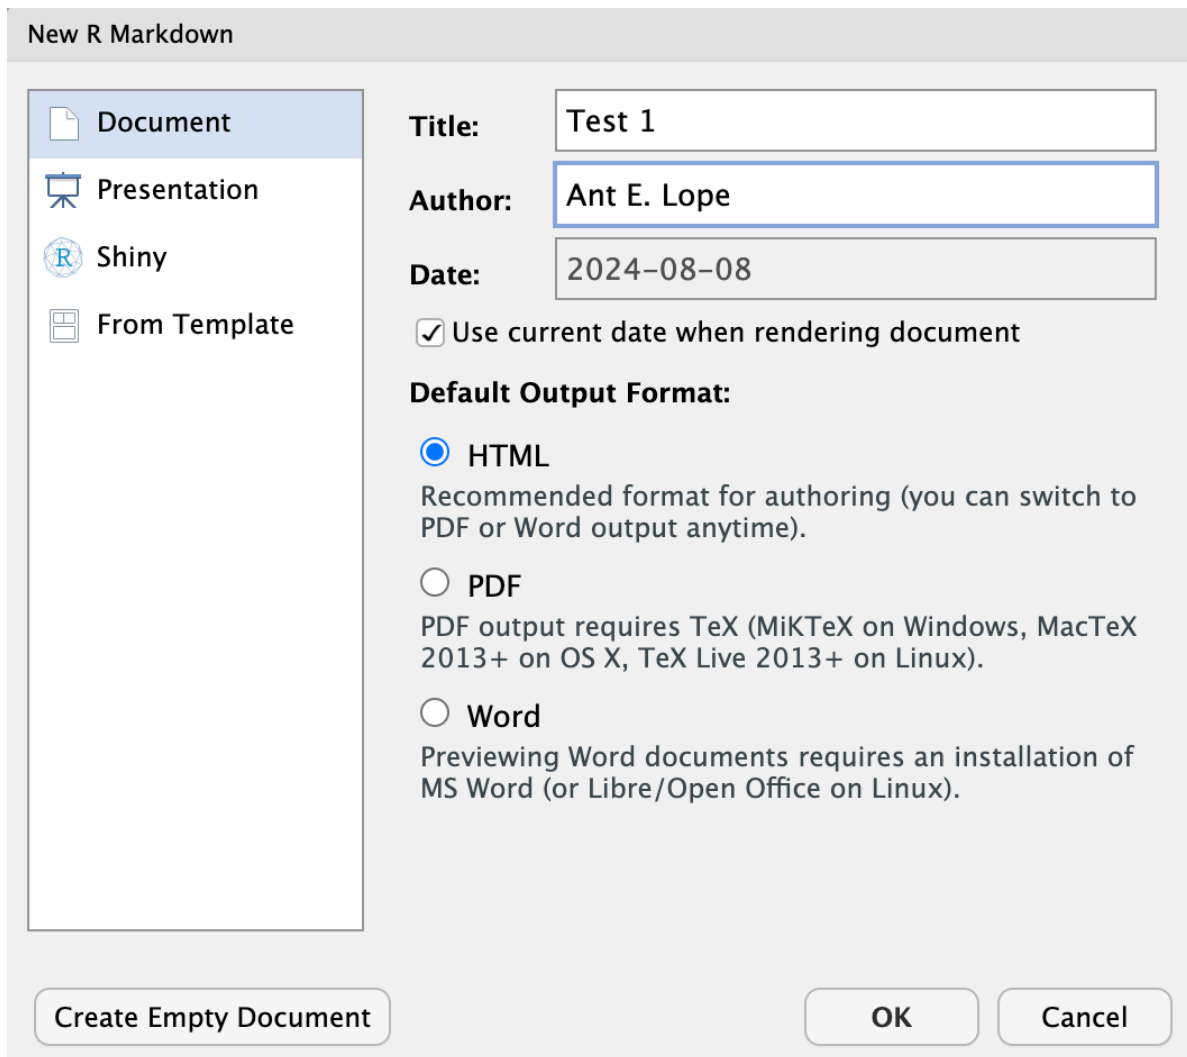
2.1 Creating an *RMarkdown* document

2.1.1 Setup

2.1.1.1 Creating a new document

In this class, we will be creating documents in what is called *RMarkdown*. This is a rich-text version of *R* that allows us to create documents with the code embedded. In *RStudio*, click the “+” button in the far top left to open the New Document menu. Scroll down this list and click on R Markdown.

A screen like this will appear:



The image shows a 'New R Markdown' dialog box. On the left is a vertical menu with four options: 'Document' (selected), 'Presentation', 'Shiny', and 'From Template'. Each option has a small icon to its left. To the right of the menu are three text input fields: 'Title:' with the value 'Test 1', 'Author:' with the value 'Ant E. Lope', and 'Date:' with the value '2024-08-08'. Below these fields is a checked checkbox labeled 'Use current date when rendering document'. Further down is the section 'Default Output Format:' with three radio button options: 'HTML' (selected), 'PDF', and 'Word'. Each radio button option has a descriptive text block below it. At the bottom of the dialog are three buttons: 'Create Empty Document', 'OK', and 'Cancel'.

New R Markdown

Document

Presentation

Shiny

From Template

Title: Test 1

Author: Ant E. Lope

Date: 2024-08-08

☒ Use current date when rendering document

Default Output Format:

☒ **HTML**
Recommended format for authoring (you can switch to PDF or Word output anytime).

☐ **PDF**
PDF output requires TeX (MiKTeX on Windows, MacTeX 2013+ on OS X, TeX Live 2013+ on Linux).

☐ **Word**
Previewing Word documents requires an installation of MS Word (or Libre/Open Office on Linux).

Create Empty Document

OK

Cancel

Figure 2.1: A new file window for an *RMarkdown* file.

After entering a title and your name and selecting **document** in the left hand menu, click **OK**.

2.1.1.2 Document setup

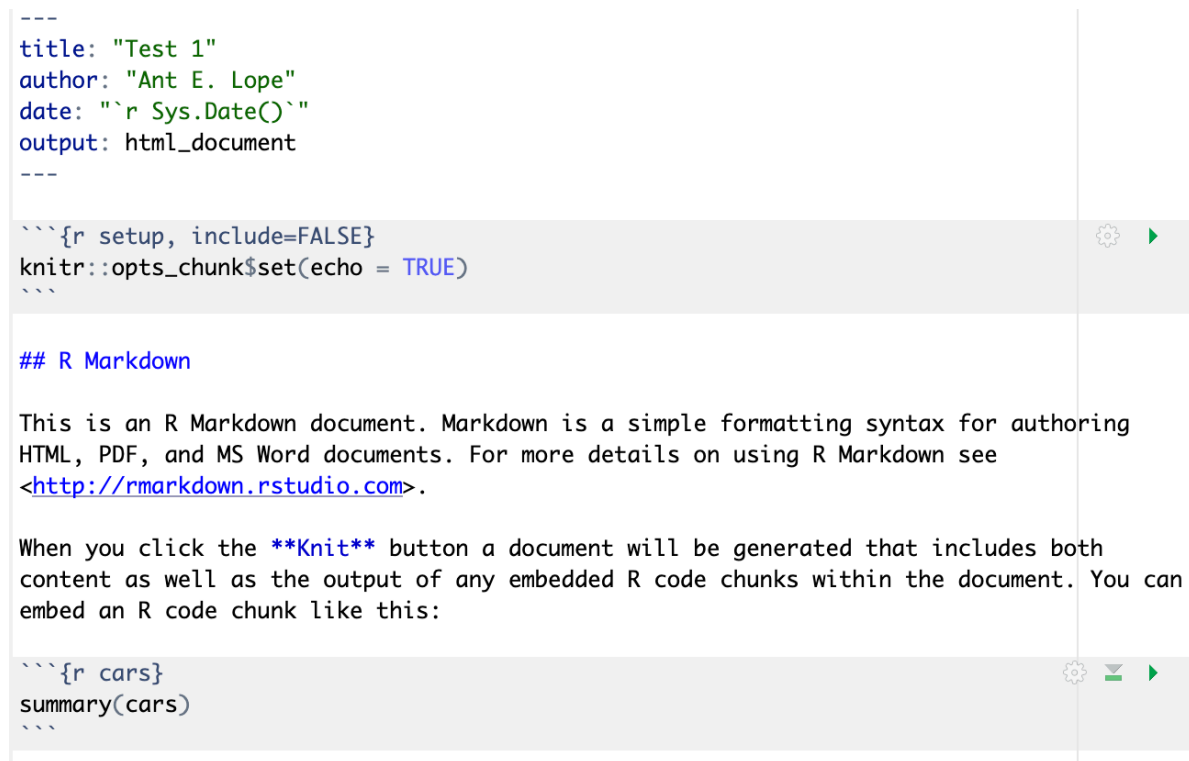


Figure 2.2: An example of a markdown script.

In the image above, we can see what a “default” *RMarkdown* script looks like after creating the file. At the top of the document, between all of the dashes, we have the `yaml` header that tells *R* what kind of document will be created, who the author is, and what the date is for the document (by default, this is 2025-08-25, which just prints the data on which the code was run). In this class, we will be saving documents as `html` as they are the easiest documents to create and save. These documents will include all of your code, text, and even any plots you may create!

- **The first thing you should do is save your document.** Since this is a “first go”, feel free to save it as `test.rmd` or something to that effect, *but save the document in the folder you want to use for assignments!*
- After the document is saved, you can hit the `knit` button at the top of the screen and see what happens. It should create a document in the same location as your saved file that is an `html` format.

- A new document will appear in a pop up window as well. The information in this document will be the default text loaded above.
- After running the document once, delete all text below the --- at the bottom of the `yaml` header so we can create our own document.

2.1.1.3 Writing text

Plain text in your *RStudio* window will be rendered as plain text in the final document. (i.e., whatever you type normally will become “normal text” in the finished document).

- **Important shortcut:** An important shortcut is the “double click”. Write a word, any word (like `potato`) in your document. If you double click, it will select the whole word. If you have a word or sentence selected, you can hit `"`, `(`, `[`, `{`, or `*` and *RStudio* will automatically put the closing equivalent of the symbol at the end of your selection.
- **Headers:** Lines preceded with `#` will become headers; the number of `#` determines the level of the header, with `#` being main headers, `##` being a second level header, and `###` being a third level header, etc.
- **Formatting:** Words can be italicized by putting an asterisk on each side of the word (`*italic*`) and bold by putting two asterisks on each side (`**bold**`).
- **Hyperlinks:** URLs are also supported, with `<>` on each side of a URL making it clickable. To turn a word into a hyperlink, type `[words to show](https://target_URL.com)`. It is important that the words are in brackets and the URL is in parentheses, or it will not format correctly.

2.1.2 Coding

In your computer, you should have erased all information except for the `yaml` header between the dashes at the top of the document. Make sure to **save** your file in a folder where you want your assignment to be located. It is important you do this step up front as the computer will sometimes save in random places if you don’t specify a file location at the beginning. *Don’t forget to save your work frequently!*

2.1.2.1 Code chunks

Code chunks are sections within the document where we can write code and it will be executed. These chunks are always preceded by a series of three “hashes” and the code language specified in brackets (i.e., `{r}`) and are followed by a set of three hashes as well. In your window, these will appear with a greyed-out background to show that the area is a code chunk.

We will create an example set of text, and then insert a code chunk to see how this works.

2.1.2.2 Test text and chunks

- Type `This is a test of the *Rmarkdown* code.` as plaintext in your document.
- Hit `Enter` twice.
- Insert a code chunk by using the correct shortcut:
 - **Windows/Linux:** `CTRL + ALT + I`
 - **Mac:** `⌘ + ⌥ + I`
- Push the down arrow until you are between the sets of hashes in the code chunk.
- Type `print("hello world")`
 - `print`: a command that returns a **string** of characters. A string is any set of characters that should be read together as a single object or line.
 - `"hello world"`: placing the letters inside quotes `"` tells *R* that these are *character data* that should be read as a **string**.
 - The full command should return an output of `"hello world"`.

This is a test of the **RMarkdown** code.

```
```${r}
print("Hello world!")
```
```

Figure 2.3: Text to type in your *Rmarkdown* document.

After typing this into the document, hit `knit` near the top of the upper left window. *R* will now create an HTML document that should look like this:

Test 1

Ant E. Lope

2024-08-07

This is a test of the *R*Markdown code.

```
print("Hello world!")
```

```
## [1] "Hello world!"
```

Figure 2.4: The output from the above code knitted into a document.

We can see now that the HTML document has the `yaml` header at the top of the document for the title of the document, the author's name, the date on which the code was run. The plain text we used has been printed, and the greyed-out box with color coded *R* code is followed by the output we expected.

- Let's try something a little more complex. Create a new code chunk and type the following:
 - *NOTE*: Mac users can use `⌘ + -` to automatically insert `<-`, saving a lot of time. Unfortunately, there is no equivalent shortcut for Windows.
 - We are avoiding using `=` to create objects as this can lead to issues in saving objects to *R*'s memory in some cases and can lead to confusion between objects and parts of functions.

```
x <- 1:10
```

- `x`: putting a plain letter, character, or word in *R* asks it to refer to something in its memory. If it does not exist in its memory, it can be created by using something like `<-`. When we define a set of characters as numbers or as something else, we are creating an **object**.
- `1:10`: this will print all whole integers between the listed numbers, inclusive of those numbers. In this case, `1:10` is the same as `c(1,2,3,4,5,6,7,8,9,10)`, where `c` *concatenates* and puts together the numbers in a single **string**.
- Run your code by placing your cursor on the same line as `x <- 1:10` and pressing **CTRL + Enter** / **⌘ + return**.

- After running, you will see the following appear in your top right window, showing you that `x` was saved:

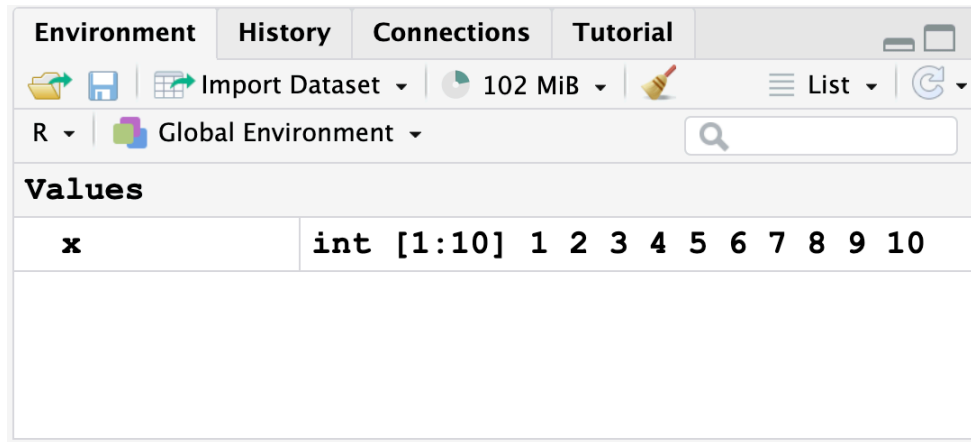


Figure 2.5: The object `x` now being saved to the environment, as a set of ten integers.

- Hit `Enter` / `return` to create a new line within the code chunk. Type the letter `x` and hit `CTRL + Enter` / `+ return`. You should see the following:

```
x
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

- Create a new object, `y`, that is the square of `x`, as follows:
 - `^`: this is the “power” symbol; thus, `x^2` is x^2 , `x^3` is x^3 , etc.
 - Don’t forget to use `CTRL + Enter` / `+ return` to run the line of code.

```
y <- x^2
```

After running the code, you should have both objects in your memory, like so:

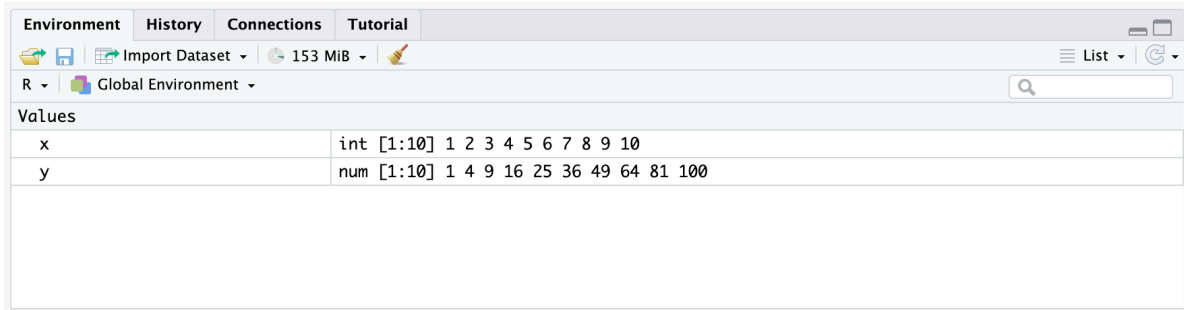


Figure 2.6: *RStudio* environment window showing saved objects. These are in the computer’s memory.

We can now print `y`, just as we did for `x`, to see what the object contains in the document:

```
y
```

```
[1] 1 4 9 16 25 36 49 64 81 100
```

- *Note* that you can multiply objects together as well. When we have strings, the strings will be multiplied together. Try the following:

```
x*y
```

```
[1] 1 8 27 64 125 216 343 512 729 1000
```

```
y-x
```

```
[1] 0 2 6 12 20 30 42 56 72 90
```

- **Note:** Since I didn’t “store” this value as a variable in *R* using `<-`, the value is not in the environment.

2.1.2.3 Taking notes

We can take notes and annotate code within a code chunk by using the character `#` before anything we wish to annotate. *Note* that the color of the text after `#` inside the code chunk will change to show us that the code will not run.


```
x
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

```
# y
```

Above, we can see only `x` printed, as `y` was blocked with a `#`. We can use this to add text inside the code as well:

```
# subtract x from y  
y - x
```

```
[1] 0 2 6 12 20 30 42 56 72 90
```

You are *highly encouraged* to add notes to your code, both to help others understand your logic and to remind yourself of what you did in the future!


2.1.3 Plotting

Now, let's try creating a plot. This is easy in *R*, as we just use the command `plot`. At minimum, we must define an `x` variable and a `y` variable for the command `plot`. It will calculate and make everything else automatically.

- **Note:** Within commands, we must use `=`.


Type and run the following:

```
plot(x = x, y = y)
```



Now let's make the plot with some new visual parameters. I describe what these arguments do using `#` to make notes within the code chunk. *Note* that you can split commands over multiple lines after typing `,` but that the command must all be enclosed within parentheses `()`.

```
plot(x = x, # specify x values
     y = y, # specify y values
     ylab = "Y Values", # specify Y label
     xlab = "X Values", # specify X label
     main = "Plot Title", # specify main title
     pch = 19, # adjust point style
     col = "red") # make points red
```



2.2 Stop touching your mouse

RStudio has lots of keyboard shortcuts that will make your life easier. You will be a lot faster if you rely on these shortcuts. I outline some here, but see the [Glossary](#).

- Use arrow keys to move around your document
- Highlight code by holding `shift` and pushing the arrow keys
- Copy text using `CTRL + C` / `⌘ + C`
- Paste text using `CTRL + V` / `⌘ + V`
- Cut text using `CTRL + X` / `⌘ + X`

- Undo last edit or mistake using **CTRL + Z / ⌘ + Z**
- Save your document using **CTRL + S / ⌘ + S**

These shortcuts are worth learning; they will greatly increase the speed at which you get things done.

2.2.1 Tab complete

RStudio allows for “tab-completing” while typing code. Tab-completing is a way of typing the first part of a command, variable name, or file name and hitting “tab” to show all options with that spelling. You should use tab completing because it:

- reduces spelling mistakes
- reduces filepath mistakes
- increases the speed at which you code
- provides help with specific functions

2.2.2 Help

At any point in *R*, you can look up “help” for a specific function by typing `?functionname`. Try this on your computer with the following:

```
?mean
```

NOTE that you cannot **knit** a document with a help command in it! So you need to remove any lines with them before creating your document or *R* will be very unhappy.

3 Running your first analysis

Here, we are going to go through a work through with our first set of data. This will be your first assignment using *RMarkdown* and *RStudio*.

3.1 Create a new document

Refer to the previous section to create a new document. Name it `LastName_first_analysis.rmd`, with `LastName` replaced by your surname. Save this document in your assignments folder, and make sure it will `knit` as an `html` document.

- You should *always* have your last name in your file name when you submit documents in this class.

3.2 Working with data

Throughout this course, we are going to have to work with datasets that are from our book or other sources. Here, we are going to work through an example dataset.

- First, we need to install *libraries*. A *library* is a collated, pre-existing batch of code that is designed to assist with data analysis or to perform specific functions. These libraries make life a lot easier, and create short commands for completing relatively complex tasks.

3.2.1 Libraries

In this class, there is one major library that you will need *almost every week*! **Even if I don't declare this library, you should load it in your documents.**

- Libraries are declared at the very beginning of the document
- Libraries *must* be in a code chunk at the top of the document or else they will not be loaded before the code requiring them, preventing your document from being `knit`.

- Make sure you read your error codes - many times it will say that your command is not found, and that often means your library is not loaded. That, or you misspelled something.

1. **tidyverse**: this package is actually a [group of packages](#) designed to help with data analysis, management, and visualization.

NOTE: If you leave the install prompts in your RMarkdown document, *it will not knit!* Install the following using the bottom right coding window of your *RStudio* session.

```
# run this code the first time ONLY
# DO NOT INCLUDE IN RMD FILE
# does not need to be run every time you use R!

# tidyverse has a bunch of packages in it!
# great for data manipulation
install.packages("tidyverse")

# if you ever need to update:
# leaving brackets open means "update everything"
update.packages()
```

- After packages are installed, we will need to load them into our *R* environment. While we only need to do `install.packages` once on our machine, **we need to load libraries every time we restart the program!**

NOTE: The following is required in EVERY DOCUMENT that uses the tidyverse commands!

```
### MUST BE RUN AT THE TOP OF EVERY DOCUMENT ###
# Load tidyverse into R session / document
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.4      v readr      2.1.5
v forcats    1.0.0      v stringr    1.5.1
v ggplot2    3.5.2      v tibble     3.3.0
v lubridate  1.9.4      v tidyr      1.3.1
v purrr      1.1.0

-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

You should have an output like the above. What this means is:

1. The core packages that comprise the **tidyverse** loaded successfully, and version numbers for each are shown.
2. The conflicts basically means that certain commands will not work as they used to because *R* has “re-learned” a particular word.

3.2.1.1 Conflicts

To clarify the conflicts, pretend that you can only know one definition of a word at a time, in this case, the word “cola”.

- **English:** Cola is a type of soda pop or soft drink.
- **Spanish:** Cola refers to a line or a tail.

While we can figure out which definition is being used based on context, *R* can’t. It will always use the most recent definition, such that it may interpret something as “Mi perro movió la *refresco*” or “I bought a *tail* from the vending machine”. To avoid this confusion in *R*, we specify which “cola” we are referring to. In the above, this would look like “Mi perro movió la `español::cola`” and “I bought a `english::cola` from the vending machine”.

We should not have many conflicts in this class, but be aware they may exist.

3.2.2 Downloading data

Now, we need to download our first data set. These datasets are stored on GitHub. We are going to be looking at data from Dr. Cooper’s dissertation concerning Afrotropical bird distributions (Cooper 2021). This website is in the data folder on this websites’ GitHub page, [accessible here](#).

```
# read comma separated file (csv) into R memory
# reads directly from URL
ranges <- read_csv("https://raw.githubusercontent.com/jacobccooper/biol305_unk/main/datasets,
```

```
Rows: 12 Columns: 10
-- Column specification -----
Delimiter: ","
chr (1): species
dbl (9): combined_current_km2, consensus_km2, bioclim_current_km2, 2050_comb...
```

```
i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

Alternatively, we can use the operator `|>` to simplify this process. `|>` means “take whatever you got from the previous step and *pipe* it into the next step”. So, the following does the exact same thing:

```
ranges <- "https://raw.githubusercontent.com/jacobccooper/biol305_unk/main/datasets/lacustrii"
read_csv()
```

```
Rows: 12 Columns: 10
-- Column specification -----
Delimiter: ","
chr (1): species
dbl (9): combined_current_km2, consensus_km2, bioclim_current_km2, 2050_comb...

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

Using the `|>` is preferred as you can better set up a workflow and because it more closely mimics other coding languages, such as `bash`.

Let’s view the data to see if it worked. We can use the command `head` to view the first few rows:

```
head(ranges)
```

```
# A tibble: 6 x 10
  species                combined_current_km2 consensus_km2 bioclim_current_km2
  <chr>                  <dbl>          <dbl>          <dbl>
1 Batis_diops            25209.          6694.          19241.
2 Chamaetylas_poliophrys 68171.          1106.          68158.
3 Cinnyris_regius        60939.         13305.         53627.
4 Cossypha_archeri       27021.          6409.          11798.
5 Cyanomitra_alinae      78680.         34320.         63381.
6 Graueria_vittata       8770.           861.           8301.
# i 6 more variables: `2050_combined_km2` <dbl>, `2050_consensus_km2` <dbl>,
#   `2070_combined_km2` <dbl>, `2070_consensus_km2` <dbl>,
#   alltime_consensus_km2 <dbl>, past_stable_km2 <dbl>
```

We can perform a lot of summary statistics in *R*. Some of these we can view for multiple columns at once using `summary`.

```
summary(ranges)
```

| species | combined_current_km2 | consensus_km2 | bioclim_current_km2 |
|-----------------------|----------------------|-------------------|---------------------|
| Length:12 | Min. : 8770 | Min. : 861.3 | Min. : 3749 |
| Class :character | 1st Qu.: 24800 | 1st Qu.: 4186.2 | 1st Qu.: 10924 |
| Mode :character | Median : 43654 | Median : 7778.1 | Median : 31455 |
| | Mean : 68052 | Mean : 18161.8 | Mean : 42457 |
| | 3rd Qu.: 70798 | 3rd Qu.: 18558.7 | 3rd Qu.: 62835 |
| | Max. : 232377 | Max. : 79306.6 | Max. : 148753 |
| 2050_combined_km2 | 2050_consensus_km2 | 2070_combined_km2 | 2070_consensus_km2 |
| Min. : 1832 | Min. : 0.0 | Min. : 550.3 | Min. : 0.0 |
| 1st Qu.: 6562 | 1st Qu.: 589.5 | 1st Qu.: 6583.8 | 1st Qu.: 311.4 |
| Median : 26057 | Median : 6821.9 | Median : 24281.7 | Median : 2714.6 |
| Mean : 33247 | Mean : 14418.4 | Mean : 31811.0 | Mean : 8250.5 |
| 3rd Qu.: 40460 | 3rd Qu.: 18577.1 | 3rd Qu.: 38468.9 | 3rd Qu.: 10034.4 |
| Max. : 132487 | Max. : 79236.2 | Max. : 129591.0 | Max. : 53291.8 |
| alltime_consensus_km2 | past_stable_km2 | | |
| Min. : 0.0 | Min. : 0.0 | | |
| 1st Qu.: 790.9 | 1st Qu.: 0.0 | | |
| Median : 8216.8 | Median : 0.0 | | |
| Mean : 15723.3 | Mean : 127.3 | | |
| 3rd Qu.: 19675.0 | 3rd Qu.: 0.0 | | |
| Max. : 82310.5 | Max. : 1434.8 | | |

As seen above, we now have information for the following statistics for each variable:

- Min = minimum
- 1st Qu. = 1st quartile
- Median = middle of the dataset
- Mean = average of the dataset
- 3rd Qu. = 3rd quartile
- Max. = maximum

We can also calculate some of these statistics manually to see if we are doing everything correctly. It is easiest to do this by using predefined functions in *R* (code others have written to perform a particular task) or to create our own functions in *R*. We will do both to determine the average of `combined_current_km2`.

3.2.3 Subsetting data

First, we need to select only the column of interest. In *R*, we have two ways of subsetting data to get a particular column.

- `var[rows,cols]` is a way to look at a particular object (`var` in this case) and choose a specific combination of `row` number and `column` number (`col`). This is great if you know a specific index, but it is better to use a specific name.
- `var[rows,"cols"]` is a way to do the above but by using a specific column name, like `combined_current_km2`.
- `var$colname` is a way to call the specific column name directly from the dataset.

```
# using R functions
```

```
ranges$combined_current_km2
```

```
[1] 25209.4 68171.2 60939.2 27021.3 78679.9 8769.9 232377.2 17401.4
[9] 51853.5 35455.1 23570.3 187179.1
```

As shown above, calling the specific column name with `$` allows us to see only the data of interest. We can also save these data as an object.

```
current_combined <- ranges$combined_current_km2
```

```
current_combined
```

```
[1] 25209.4 68171.2 60939.2 27021.3 78679.9 8769.9 232377.2 17401.4
[9] 51853.5 35455.1 23570.3 187179.1
```

Now that we have it as an object, specifically a numeric vector, we can perform whatever math operations we need to on the dataset.

```
mean(current_combined)
```

```
[1] 68052.29
```

Here, we can see the mean for the entire dataset. However, we should always round values to the same number of decimal points as the original data. We can do this with `round`.

```
round(mean(current_combined),1) # round mean to one decimal
```

```
[1] 68052.3
```

Note that the above has a nested set of commands. We can write this exact same thing as follows:

```
# pipe mean through round
current_combined |>
  mean() |>
  round(1)
```

```
[1] 68052.3
```

Use the method that is easiest for you to follow!

We can also calculate the mean manually. The mean is $\frac{\sum_{i=1}^n x}{n}$, or the sum of all the values within a vector divided by the number of values in that vector.

```
# create function
# use curly brackets to denote function
# our data goes in place of "x" when finally run
our_mean <- function(x){
  sum_x <- sum(x) # sum all values in vector
  n <- length(x) # get length of vector
  xbar <- sum_x/n # calculate mean
  return(xbar) # return the value outside the function
}
```

Let's try it.

```
our_mean(ranges$combined_current_km2)
```

```
[1] 68052.29
```

As we can see, it works just the same as `mean`! We can round this as well.

```
ranges$combined_current_km2 |>
  our_mean() |>
  round(1)
```

```
[1] 68052.3
```

3.3 Creating a new column

The package `tidyverse` allows for easily creating new columns within a dataset using the command `mutate`.

For example, we can create a column of `combined_current_km2 + 1`.

```
ranges <- ranges |>
  mutate(combined_plus_one = combined_current_km2 + 1)
head(ranges$combined_plus_one)
```

```
[1] 25210.4 68172.2 60940.2 27022.3 78680.9 8770.9
```

3.4 Your turn!

Please complete the following:

1. Create an *RMarkdown document* that will save as an `.html`.
2. Load the data, as shown here, and print the summary statistics in the document.
3. Calculate the value of `combined_current_km2` divided by `2050_combined_km2` and print the results.
 - *Hint:* you can divide and multiply objects in *R*, like `a + b`, `a/b`, etc.
4. `knit` your results, with your name(s) and date, as an HTML document.

Let me know if you have any issues.

4 Descriptive Statistics

4.1 Purposes of descriptive statistics

Descriptive statistics allow us to quickly get an idea of how our datasets are “behaving”. This allows us to identify potential errors, trends, or particular outliers for further analysis. While you are likely familiar with many of the calculations we will go over, these remain powerful calculations for initial analysis. In this page, we will go over how to complete these calculations in *R*. You will then complete an assignment based on these topics.

4.2 Sample data and preparation

As with every week, we will need to load our relevant packages first. This week, we are using the following:

```
# enables data management tools
library(tidyverse)
# DescTools can also be used
# see section on MODE
library(DescTools)
```

REMEMBER: An *R* package must be installed before it can be used with the command `install.packages` or via the *RStudio* interface.

4.3 Downloading the data

For this week, we will be using a hypothetical dataset on the number of drinks sold by hour at Starbucks.

```
starbucks <- read_csv("https://raw.githubusercontent.com/jacobccooper/biol105_unk/main/datas
```

4.4 Descriptive statistics

Descriptive statistics are statistics that help us understand the shape and nature of the data on hand. These include really common metrics such as *mean*, *median*, and *mode*, as well as more nuanced metrics like *quartiles* that help us understand if there is any *skew* in the dataset. (*Skew* refers to a bias in the data, where more data points lie on one side of the distribution and there is a long *tail* of data in the other direction).




Figure 4.1: Examples of skew compared to a symmetrical, non-skewed distribution. Source: machinelearningparatodos.com

Note above that the relative position of the *mean*, *median*, and *mode* can be indicative of skew. Please also note that these values will rarely be exactly equal “in the real world”, and thus you need to weigh differences against the entire dataset when assessing skew. There is a lot of nuance like this in statistics; it is not always an “exact” science, but sometimes involves judgment calls and assessments based on what you observe in the data.

Using the `starbucks` dataset, we can look at some of these descriptive statistics to understand what is going on.

4.4.1 Notation

We use Greek lettering for *populations* and Roman lettering for *samples*. For example:

- σ is a population, but s is a sample (both these variables refer to *standard deviation*).
- μ is a population, but \bar{x} is a sample (both of these variables refer to the *mean*).

4.4.2 Mean

The mean is the “average” value within a set of data, specifically, the sum of all values divided by the length of those values: $\frac{\sum_{i=1}^n x}{n}$.

Click to see how to calculate *Mean*

```
head(starbucks)
```

```
# A tibble: 6 x 2
  Hour      Frap_Num
  <chr>      <dbl>
1 0600-0659      2
2 0700-0759      3
3 0800-0859      2
4 0900-0959      4
5 1000-1059      8
6 1100-1159      7
```

Here, we are specifically interested in the number of frappuccinos.

```
# get vector of frappuccino number
fraps <- starbucks$Frap_Num

# get mean of vector
mean(fraps)
```

```
[1] 6.222222
```

Note that the above should be rounded to a whole number, since we were given the data in whole numbers!

```
mean(fraps) |>
  round(0)
```

```
[1] 6
```

```
# OR

round(mean(fraps),0)
```

```
[1] 6
```

We already covered calculating the average manually in our previous tutorial, but we can do that here as well:

```
# sum values
# divide by n, length of vector
# round to 0 places
round(sum(fraps)/length(fraps),0)
```

```
[1] 6
```

4.4.3 Range

The range is the difference between the largest and smallest units in a dataset. I.e., it is the maximum value in a dataset minus the smallest value in a dataset.

Click to see how to calculate *Range*

We can use the commands `min` and `max` to calculate this.

```
max(fraps) - min(fraps)
```

```
[1] 13
```

The range of our dataset is 13.

NOTE that the *R* function `range` will not return the statistic of range, it will only return the maximum and minimum values from the dataset. This is **not** the correct answer when asked for the range for any homework or test questions.

```
range(fraps)
```

```
[1] 2 15
```

4.4.4 Median

The median is also known as the 50th percentile, and is the midpoint of the data when ordered from least to greatest. If there are an even number of data points, then it is the average point between the two center points. For odd data, this is the $\frac{n+1}{2}$ -th observation. For even data, since we need to take an average, this is the $\frac{\frac{n}{2} + (\frac{n}{2} + 1)}{2}$.

Click to see how to calculate the *Median*

R has a default method for calculating median - `median`.

4.4.4.1 Odd number of values

```
# use default  
median(fraps)
```

```
[1] 4
```

Now, to calculate by hand:

```
length(fraps)
```

```
[1] 9
```

We have an odd length.

```
# order gets the order  
order(fraps)
```

```
[1] 1 3 7 2 4 6 5 9 8
```

```
# [] tells R which elements to put where  
frap_order <- fraps[order(fraps)]  
  
frap_order
```

```
[1] 2 2 2 3 4 7 8 13 15
```



```
# always use parentheses
# make sure the math maths right!
(length(frap_order)+1)/2
```

```
[1] 5
```

Which is the fifth element in the vector?

```
frap_order[5]
```

```
[1] 4
```

4.4.4.2 Even number of values

Given that we are trying to find the “middle” number, this is more complicated when using a vector that has an even number of values. Let’s calculate this by hand:

```
# remove first element
even_fraps <- fraps[-1]

even_fraps_order <- even_fraps[order(even_fraps)]

even_fraps_order
```

```
[1] 2 2 3 4 7 8 13 15
```

```
median(even_fraps)
```

```
[1] 5.5
```

Now, by hand: $\frac{\frac{n}{2} + (\frac{n}{2} + 1)}{2}$.

```
n <- length(even_fraps_order)

# get n/2 position from vector
m1 <- even_fraps_order[n/2]
# get n/2+1 position
m2 <- even_fraps_order[(n/2)+1]
```

```
# add these values, divide by two for "midpoint"
med <- (m1+m2)/2

med
```

```
[1] 5.5
```

As we can see, these values are equal!

4.4.5 Other quartiles and quantiles

We also use the 25th percentile (25%, 0.25) and the 75th percentile (75%, 0.75) to understand data distributions. These are calculated similar to the above, but the bottom quartile is only $\frac{1}{4}$ of the way between values and the 75th quartile is $\frac{3}{4}$ of the way between values. We can use the *R* function `quantile` to calculate these. **Before implementing, please read this whole section - *R* has nine different methods for calculating the quantiles!**

Click to see how to find *Quartiles* and *Quantiles*

R has a default function for finding quantiles which is called `quantile`.

```
quantile(frap_order)
```

| | | | | |
|----|-----|-----|-----|------|
| 0% | 25% | 50% | 75% | 100% |
| 2 | 2 | 4 | 8 | 15 |

We can specify a quantile as well:

```
quantile(frap_order, 0.75)
```

```
75%
8
```

We can also calculate these metrics by hand. Let's do it for the even dataset, since this is more difficult.

```
quantile(even_fraps_order)
```

| 0% | 25% | 50% | 75% | 100% |
|------|------|------|------|-------|
| 2.00 | 2.75 | 5.50 | 9.25 | 15.00 |

Note that the 25th and 75th percentiles are also between two different values. These can be calculated as a quarter and three-quarters of the way between their respective values.

```
# 75th percentile

n <- length(even_fraps_order)

# get position
p <- 0.75*(n+1)

# get lower value
# round down
m1 <- even_fraps_order[trunc(p)]

# get upper value
# round up
m2 <- even_fraps_order[ceiling(p)]

# position between
# fractional portion of rank
frac <- p-trunc(p)

# calculate the offset from lowest value
val <- (m2 - m1)*frac

# get value
m1 + val
```

```
[1] 11.75
```

Wait... why does our value differ?

R , by default, calculates quantiles using what is called **Type 7**, in which the quantiles are calculated by $p_k = \frac{k-1}{n-1}$, where n is the length of the vector and k refers to the quantile being used. However, in our class and in many situations, we use **Type 6** interpretation of $p_k = \frac{k}{n+1}$. Let's try using **Type 6**:

```
quantile(even_fraps_order, type = 6)
```

| 0% | 25% | 50% | 75% | 100% |
|------|------|------|-------|-------|
| 2.00 | 2.25 | 5.50 | 11.75 | 15.00 |

Now we have the same answer as we calculated by hand!

This is a classic example of how things in *R* (and in statistics in general!) can depend on interpretation and are not always “hard and fast” rules.

In this class, we will be using Type 6 interpretation for the quantiles - you will have to specify this in the quantile function EVERY TIME! If you do *not* specify Type 6, you will get the questions incorrect and you will get answers that do not agree with the book, with Excel, or what you calculate by hand.

4.4.6 Mode

The mode is the value which occurs most frequently within a dataset.

Click to see manual method

There is no default code to find the mode in *R*. We can do this by hand, by counting the number of occurrences of each value as shown below.

This method is based on the custom function for mode outlined by [Statology](#).

```
# unique counts
u <- unique(fraps)
u
```

```
[1]  2  3  4  8  7 15 13
```

```
# which elements match
match(fraps,u)
```

```
[1] 1 2 1 3 4 5 1 6 7
```

```
# count them
tab <- match(fraps,u) |>
  tabulate()
tab
```

```
[1] 3 1 1 1 1 1 1
```

Get the highest value.

```
u[tab==max(tab)]
```

```
[1] 2
```

Notice this uses `==`. This is a logical argument that means “is equal to” or “is the same as”. For example:

```
2 == 2
```

```
[1] TRUE
```

These values are the same, so `TRUE` is returned.

```
2 == 3
```

```
[1] FALSE
```

These values are unequal, so `FALSE` is returned. *R* will read `TRUE` as 1 and `FALSE` as ZERO, such that:

```
sum(2==2)
```

```
[1] 1
```

and

```
sum(2==3)
```

```
[1] 0
```

This allows you to find how many arguments match your condition quickly, and even allows you to subset based on these indices as well. Keep in mind you can use greater than `<`, less than `>`, greater than or equal to `<=`, less than or equal to `>=`, is equal to `==`, and is not equal to `!=` to identify numerical relationships. Other logical arguments include:

- `&`: both conditions must be TRUE to match (e.g., `c(10,20) & c(20,10)`). Try the following as well: `fraps < 10 & fraps > 3`.
- `&&`: and, but works with single elements and allows for better parsing. Often used with `if`. E.g., `fraps < 10 && fraps > 3`. This will not work on our multi-element `frap` vector.
- `|`: or, saying at least one condition must be true. Try: `fraps > 10 | fraps < 3`.
- `||`: or, but for a single element, like `&&` above.
- `!`: not, so “not equal to” would be `!=`.

However, the library *DescTools* has many missing basic statistical functions, including **Mode()**. Use this method in your code to reduce clutter and make the code run faster.

Click to see DescTools method

```
#Load the package
#Must be loaded each time it you reboot R
library(DescTools)

#Run the function
Mode(fraps)
```

```
[1] 2
attr("freq")
[1] 3
```

```
###Note the capitalization of mode###
```

The output of this function is three lines. In our case it reads:

```
[1] 2
attr("freq")
[1] 3
```

The **first number** is the mode; in our case, two is our mode. The second line is the computer telling us that our second number is the frequency of our first number. In our case, that means our mode, which is two, appears three times. The number 2 occurs thrice.

If there are two or more modes, these numbers will also appear in the first line of the output.

4.4.7 Variance

When we are dealing with datasets, the variance is a measure of the total spread of the data. The variance is calculated using the following:

$$\sigma^2 = \frac{\sum (x_i - \bar{x})^2}{n - 1}$$

Essentially, this means that for every value of x , we are finding the difference between that value and the mean and squaring it, summing all of these squared differences, and dividing them by the number of samples in the dataset minus one.

Click to see how to calculate *Variance*

R has a built in function for variance - `var`.

```
var(frap_order) |>
  round(2)
```

```
[1] 24.44
```

Click to see how to calculate *Variance* by hand

Let's calculate the `frap` dataset variance by hand.

```
frap_order
```

```
[1] 2 2 2 3 4 7 8 13 15
```

Now to find the differences.

```
diffs <- frap_order - mean(frap_order)

diffs
```

```
[1] -4.2222222 -4.2222222 -4.2222222 -3.2222222 -2.2222222  0.7777778  1.7777778
[8]  6.7777778  8.7777778
```

Note that R is calculating the same thing for the entire vector! Since these are differences from the mean, they should sum to zero.

```
sum(diffs)
```

```
[1] 3.552714e-15
```

This is not quite zero due to rounding error, but is essentially zero as it is 0.0000000000000036.

```
# square differences
diffs_sq <- diffs^2

diffs_sq
```

```
[1] 17.8271605 17.8271605 17.8271605 10.3827160  4.9382716  0.6049383  3.1604938
[8] 45.9382716 77.0493827
```

Now we have the squared differences. We need to sum these and divide by $n - 1$.

```
n <- length(frap_order)

var_frap <- sum(diffs_sq)/(n-1)

var_frap
```

```
[1] 24.44444
```

Let's check this against the built-in variance function in *R*.

```
var(frap_order)
```

```
[1] 24.44444
```

They are identical! We can check this using a logical argument.

```
var_frap == var(frap_order)
```

```
[1] TRUE
```

Seeing as this is `TRUE`, we calculated it correctly.

4.4.8 Standard deviation

The standard deviation (σ) is the square root of the variance, and is used to help describe the spread of datasets.

[Click to see how to calculate standard deviation](#)

```
sqrt(var_frap)
```

```
[1] 4.944132
```

There is also the built in `sd` function in *R*:

```
sd(frap_order)
```

```
[1] 4.944132
```

We can test the manual method against the built in `sd` function in *R*:

```
sqrt(var_frap) == sd(frap_order)
```

```
[1] TRUE
```

As you can see, we calculated this correctly!

4.4.9 Standard error

The standard error is used to help understand the spread of data and to help estimate the accuracy of our measurements for things like the mean. The standard error is calculated thusly:

$$SE = \frac{\sigma}{\sqrt{n}}$$

[Click to see how to calculate standard error](#)

There is not a built in function for standard error in *R*, but we can use *DescTools*.

```
MeanSE(frap_order) #This is from DescTools!
```

```
[1] 1.648044
```

Remember, the standard error is used to help reflect our *confidence* in a specific measurement (*e.g.*, how certain we are of the mean, and what values we believe the mean falls between). We want our estimates to be as precise as possible with as little uncertainty as possible. Given this, does having more samples make our estimates more or less confident? Mathematically, what happens as our sample size *increases*?

4.4.10 Coefficient of variation

The coefficient of variation, another measure of data spread and location, is calculated by the following:

$$CV = \frac{\sigma}{\mu} * 100$$

Click to see how to calculate coefficient of variation

If we were to calculate this by hand, it would be done with this function:

```
cv <- function(x){  
  sigma <- sd(x)  
  mu <- mean(x)  
  val <- sigma/mu*100  
  return(val)  
}  
  
cv(frap_order) |>  
  round(2)
```

```
[1] 79.46
```

```
#Remember that we will need to round values.
```

However, we can also complete this with a *DescTools* function.

```
(CoefVar(frap_order)*100) |>  
  round(2)
```

```
[1] 79.46
```

4.4.11 Outliers

[Click to see manual method](#)

Outliers are any values that are outside of the 1.5 times the interquartile range. We can calculate this for our example dataset as follows:

```
lowquant <- quantile(frap_order,0.25,type = 6) |> as.numeric()
hiquant <- quantile(frap_order,0.75,type = 6) |> as.numeric()
iqr <- hiquant - lowquant
```

[Click to see IQR method](#)

We can also calculate the interquartile range using IQR. **Remember**, you must use `type = 6`!

```
iqr <- IQR(frap_order, type = 6)
```

Now, to calculate the “whiskers”.

```
lowbound <- lowquant - (1.5*iqr)
hibound <- hiquant + (1.5*iqr)

# low outliers?
# select elements that match
# identify using logical "which"
frap_order[which(frap_order < lowbound)]
```

```
numeric(0)
```

```
# high outliers?
# select elements that match
# identify using logical "which"
frap_order[which(frap_order > hibound)]
```

```
numeric(0)
```

We have no outliers for this particular dataset.

Click to see DescTools method

This can also be done with *DescTools*.

```
Outlier(frap_order)
```

```
numeric(0)
```

Again, there are no outliers as this code is effectively the same. If there are, it will return all outliers, but not indicate if they are high or low. This can be solved by comparing the returned outlier to the mean. If it is above the mean, it is a high outlier, and vice versa.

4.5 Homework: Descriptive Statistics

Now that we've covered these basic statistics, it's your turn!

4.5.1 Homework instructions

Please create an *RMarkdown* document that will render as an `.html` file. You will submit this file to show your coding and your work. Please refer to the [Introduction to R](#) for refreshers on how to create an `.html` document in *RMarkdown*. You will need to do the following for each of these datasets:

- mean
- median
- range
- interquartile range
- variance
- standard deviation
- coefficient of variation
- standard error
- whether there are any “outliers”

Please show all of your work for full credit.

4.5.2 Data for homework problems

Click for homework data

For each question, calculate the mean, median, range, interquartile range, variance, standard deviation, coefficient of variation, standard error, and whether there are any “outliers”.

Please also write your own short response to the *Synthesis question* posed, which will involve thinking about the data and metrics you just analyzed.

4.5.2.1 1: UNK Nebraskans

Every year, the university keeps track of where students are from. The following are data on the number of students admitted to [UNK from the state of Nebraska](#):

```
# create dataset in R
nebraskans <- c(5056,5061,5276,5244,5209,
               5262,5466,5606,5508,5540,5614)

years <- 2023:2013

nebraskans_years <- cbind(years,nebraskans) |>
  as.data.frame()

nebraskans_years
```

| | years | nebraskans |
|----|-------|------------|
| 1 | 2023 | 5056 |
| 2 | 2022 | 5061 |
| 3 | 2021 | 5276 |
| 4 | 2020 | 5244 |
| 5 | 2019 | 5209 |
| 6 | 2018 | 5262 |
| 7 | 2017 | 5466 |
| 8 | 2016 | 5606 |
| 9 | 2015 | 5508 |
| 10 | 2014 | 5540 |
| 11 | 2013 | 5614 |

Using these data, please calculate the mean, median, range, interquartile range, variance, standard deviation, coefficient of variation, standard error, and whether there are any “outliers” for the number of UNK students from Nebraska.

In order to do this, we will need to first select the column that denotes the number of Nebraskans from the dataset. Remember, we need to save this as an object in *R* to do the analyses. Here is a demonstration of getting the column to look at the mean, so that you can repeat this for the other questions. This relies heavily on the use of `$`, used to get the data from a specific column.

```
# $ method
nebraskans <- nebraskans_years$nebraskans

nebraskans
```

```
[1] 5056 5061 5276 5244 5209 5262 5466 5606 5508 5540 5614
```

Now we can get the `mean` of this vector.

```
mean(nebraskans) |>
  round(0) # don't forget to round!
```

```
[1] 5349
```

Synthesis question: Do you think that there are any really anomalous years? Do you feel data are similar between years? *Note* we are not looking at trends through time but whether any years are outliers.

4.5.2.2 2: Piracy in the Gulf of Guinea

The following is a dataset looking at oceanic conditions and other variables associated with pirate attacks within the region between 2010 and 2021 (Moura et al. 2023). Using these data, please calculate the mean, median, range, interquartile range, variance, standard deviation, coefficient of variation, standard error, and whether there are any “outliers” for distance from shore for each pirate attack (column `Distance_from_Coast`).

```
pirates <- read_csv("https://figshare.com/ndownloader/files/42314604")
```

New names:

Rows: 595 Columns: 40

-- Column specification

----- Delimiter: "," chr
(18): Period, Season, African_Season, Coastal_State, Coastal_Zone, Navi... dbl

```
(20): ...1, Unnamed: 0, Year, Month, Lat_D, Lon_D, Distance_from_Coast,... dtm
(2): Date_Time, Date
i Use `spec()` to retrieve the full column specification for this data. i
Specify the column types or set `show_col_types = FALSE` to quiet this message.
* `` -> `...1`
```

Synthesis question: Do you notice any patterns in distance from shore? What may be responsible for these patterns? *Hint:* Think about what piracy entails and also what other columns are available as other variables in the above dataset.

4.5.2.3 3: Patient ages at presentation

The following is a dataset on skin sores in different communities in Australia and Oceania, specifically looking at the amount of time that passes between skin infections (Lydeamore et al. 2020a). This file includes multiple different datasets, and focuses on data from children in the first five years of their life, on household visits, and on data collected during targeted studies (Lydeamore et al. 2020b).

```
ages <- read_csv("https://doi.org/10.1371/journal.pcbi.1007838.s006")
```

```
Rows: 17150 Columns: 2
-- Column specification -----
Delimiter: ","
chr (1): dataset
dbl (1): time_difference

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

Let's see what this file is like real fast. We can use the command `dim` to see the **rows** and **columns**.

```
dim(ages)
```

```
[1] 17150      2
```

As you can see, this file has only two columns but 17,150 rows! For the column `time_difference`, please calculate the mean, median, range, interquartile range, variance, standard deviation, coefficient of variation, standard error, and whether there are any “outliers”.

Synthesis question: Folks will often think about probabilities of events being “low but never zero”. What does that mean in the context of these data? What about these data make you feel like probabilities may decrease through time but never become zero?

5 Diagnosing data visually

5.1 The importance of visual inspection

Inspecting data visually can give us a lot of information about whether data are normally distributed and about whether there are any major errors or issues with our dataset. It can also help us determine if data meet model assumptions, or if we need to use different tests more appropriate for our datasets.

5.2 Sample data and preparation

Before we start, we must load our *R* libraries.

```
library(tidyverse)
```

5.3 Histograms

A histogram is a frequency diagram that we can use to visually diagnose data and their distributions. We are going to examine a histogram using a random string of data. *R* can generate random (though, actually pseudorandom) strings of data on command, pulling them from different distributions. These distributions are pseudorandom because we can't actually program *R* to be random, so it starts from a wide variety of pseudorandom points.

5.3.1 Histograms on numeric vectors


Click to see how to make a default histogram

The following is how to create default histograms on data. If you need to create custom bin sizes, please see the notes under [Cumulative frequency plot](#) for data that are not already in frequency format.

```
# create random string from normal distribution
# this step is not necessary for data analysis in homework
set.seed(8675309)

x <- rnorm(n = 1000, # 1000 values
           mean = 0,
           sd = 1)


# make histogram
hist(x)
```



NOTE that a histogram can only be made on a vector of values. If you try to make a histogram on a data frame, *you will get an error and it will not work*. You have to specify which column you wish to use with the `$` operator. (For example, for dataframe `xy` with columns `x` and `y`, you would use `hist(xy$x)`).

We can increase the number of bins to see this better.

```
hist(x,breaks = 100)
```



The number of bins can be somewhat arbitrary, but a value should be chosen based off of what illustrates the data well. *R* will auto-select a number of bins in some cases, but you can also select a number of bins. Some assignments will ask you to choose a specific number of bins as well.

5.3.2 Histograms on frequency counts

Click to see how to make a histogram with frequency data

Say, for example, that we have a dataset where everything is already shown as frequencies. We can create a frequency histogram using `barplot`.


```
count_table <- matrix(nrow = 4, ncol = 2, byrow = T,
                      data = c("Cat 1", 4,
                                "Cat 2", 8,
                                "Cat 3", 7,
                                "Cat 4", 3)) |>
  as.data.frame()

colnames(count_table) <- c("Category", "Count")

# ensure counts are numeric data
count_table$Count <- as.numeric(count_table$Count)

# manually create histogram
barplot(count_table$Count, # response variable, counts for histogram
```

```
axisnames = T, # make names on plot
names.arg = count_table$Category) # make these the names
```



5.3.3 ggplot histograms

Click to see how to make *fancy* histograms (optional)


The following is an *optional* walkthrough on how to make really fancy plots.

We can also use the program **ggplot**, part of the **tidyverse**, to create histograms.

```
# ggplot requires data frames
x2 <- x |> as.data.frame()
colnames(x2) <- "x"

ggplot(data = x2, aes(x = x)) +
  geom_histogram()
```


``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.



ggplot is nice because we can also clean up this graph a little.

```
ggplot(x2,aes(x=x)) + geom_histogram() +  
  theme_minimal()
```

``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.



We can also do a histogram of multiple values at once in *R*.

```
x2$cat <- "x"

y <- rnorm(n = 1000,
           mean = 1,
           sd = 1) |>
  as.data.frame()

colnames(y) <- "x"
y$cat <- "y"


xy <- rbind(x2,y)

head(xy)
```

```
      x cat
1 -0.9965824 x
2  0.7218241 x
3 -0.6172088 x
4  2.0293916 x
5  1.0654161 x
6  0.9872197 x
```

```
ggplot(xy, aes(x = x, fill = cat)) +  
  geom_histogram()
```


``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.



We can also make this look a little nicer.

```
ggplot(xy, aes(x = x, colour = cat)) +  
  geom_histogram(fill = "white", alpha = 0.5, # transparency  
                position = "identity") +  
  theme_classic()
```


``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.



We can show these a little differently as well.

```
ggplot(xy, aes(x = x, fill = cat))+  
  geom_histogram(position = "identity", alpha = 0.5) +  
  theme_minimal()
```

``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.




There are lots of other commands you can incorporate as well if you so choose; I recommend checking [sites like this one](#) or using ChatGPT.

5.4 Boxplots

Click to see how to make boxplots


We can also create boxplots to visualize the spread of the data. Boxplots include a bar for the median, a box representing the interquartile range between the 25th and 75th percentiles, and whiskers that extend $1.5 \cdot IQR$ beyond the 25th and 75th percentiles. We can create a boxplot using the command `boxplot`.

```
# using pre-declared variable x  
  
boxplot(x)
```



We can set the axis limits manually as well.

```
boxplot(x, # what to plot
        ylim = c(-4, 4), # set y limits
        pch = 19) # make dots solid
```



On the above plot, *outliers* for the dataset are shown as dots beyond the ends of the “whiskers”.

5.5 Skewness

[Click to read about skewness](#)

Skew is a measure of how much a dataset “leans” to the positive or negative directions (*i.e.*, to the “left” or to the “right”). To calculate skew, we are going to use the `moments` library.

```
# don't forget to install if needed!
library(moments)

skewness(x)
```


```
[1] -0.07158066
```

Generally, a **value between -1 and $+1$ for skewness is “acceptable”** and not considered overly skewed. Positive values indicate “right” skew and negative values indicate a “left” skew. If something is too skewed, it may violate assumptions of normality and thus need *non-parametric* tests rather than our standard parametric tests - something we will cover later!

Let’s look at a skewed dataset. We are going to artificially create a skewed dataset from our `x` vector.

```
# create more positive values
x3 <- c(x,
      x[which(x > 0)]*2,
      x[which(x > 0)]*4,
      x[which(x > 0)]*8)

hist(x3)
```



```
skewness(x3) |>
  round(2)
```

```
[1] 2.18
```

As we can see, the above is a heavily skewed dataset with a positive (“right”) skew.

5.6 Kurtosis


Click to read about kurtosis

Kurtosis refers to how sharp or shallow the peak of the distribution is (*platykurtic* vs. *leptokurtic*). Remember - *platykurtic* are *plateaukurtic*, wide and broad like a plateau, and *leptokurtic* distributions are sharp. Intermediate distributions that are roughly normal are *mesokurtic*.

Much like skewness, **kurtosis values of > 2 and < -2 are generally considered extreme**, and thus not mesokurtic. This threshold can vary a bit based on source, but for this class, we will use a threshold of ± 2 for both skewness and kurtosis.

Let's see the kurtosis of `x`. **Note** that when doing the equation, a normal distribution actually has a kurtosis of 3; **thus, we are doing kurtosis -3 to “zero” the distribution and make it comparable to skewness.**

```
hist(x)
```



```
# non-zeroed  
kurtosis(x) |> round(2)
```

```
[1] 3.05
```

```
# zeroed  
# put whole thing to be rounded in brackets  
(kurtosis(x)-3) |> round(2)
```

```
[1] 0.05
```

Note when rounding a calculation, the calculation must be in brackets or it will only round the last value.


As expected, our values drawn from a normal distribution are not overly skewed. Let's compare these to a more kurtic distribution:

```
xk <- x^3  
kurtosis(xk)-3
```

```
[1] 29.12246
```

What does this dataset look like?

```
hist(xk,breaks = 100)
```



As we can see, this is a very *leptokurtic* distribution.


5.7 Cumulative frequency plot

A *cumulative frequency plot* shows the overall spread of the data as a cumulative line over the entire dataset. This is another way to see the spread of the data and is often complementary to a histogram.

Click to see how to make a cumulative frequency plot if data is not in histogram/frequency format


The use of the Empirical Cumulative Distribution Function **ecdf** can turn a variable into what is needed to create a cumulative frequency plot.

```
plot(ecdf(x)) #Creating a cumulative frequency plot
```



```
plot(ecdf(x),  
      xlab = "Data Values", #Labeling the x-axis  
      ylab = "Cumulative Probability", #Labeling the y-axis  
      main = "ECDF of X") #Main title for the graph
```

ECDF of X



Click to see how to make a cumulative frequency plot if data is in histogram/frequency format

If you have a list of frequencies (say, *for river discharge over several years*), you only need to do the `cumsum` function. For example:

```
y <- c(1, 2, 4, 8, 16, 8, 4, 2, 1)

sum_y <- cumsum(y)

print(y)
```

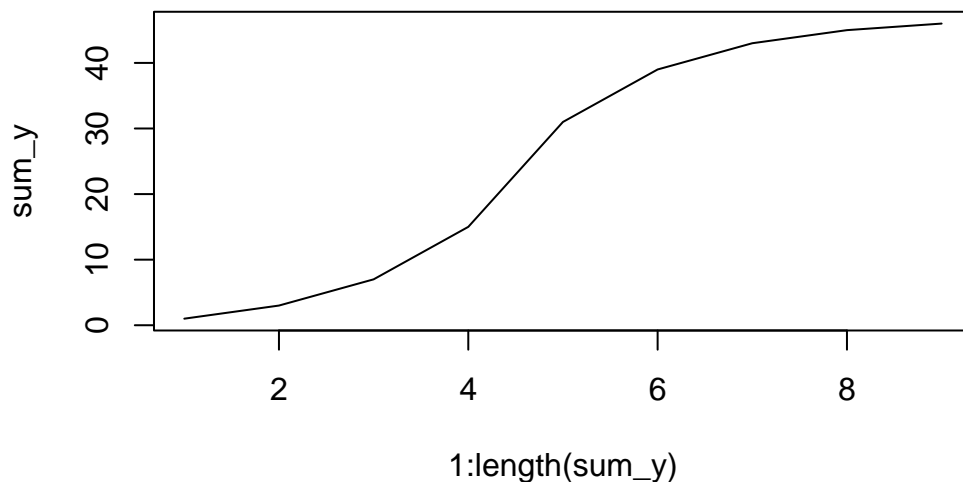
```
[1] 1 2 4 8 16 8 4 2 1
```

```
print(sum_y)
```

```
[1] 1 3 7 15 31 39 43 45 46
```

Now we can see we have our cumulative sums. Let's plot these. **NOTE** that this method will *not* have the x variables match the dataset you started with, it will only plot the curve based on the number of values given.

```
plot(x = 1:length(sum_y), # get length of sum_y, make x index
     y = sum_y, # plot cumulative sums
     type = "l") # make a line plot
```



5.8 Homework: Visualizing data

Directions:

Please complete all computer portions in an **rmarkdown** document knitted as an html. Upload any “by hand” calculations as images in the HTML or separately on *Canvas*.

5.8.1 Helpful hint

Click for a useful hint!


HINT: For 3.5, consider just making a vector of the values of interest for a histogram.

For example, see the following. For reference:

- `c` means “concatenate”, or place things together in an object.

```
# numeric vector data for counts
y <- c(17,24,16)

# manually create a histogram using barplot
barplot(y,
  # axis names must be true
  axisnames = T,
  # input names here
  # each category as a separate quoted character string
  names.arg = c("Cat 1", "Cat 2", "Cat 3"))
```

5.8.2 Problem 1

The Research Institute Nature and Forest in Brussels, Belgium, conducted a survey of fish from the southern North Sea to help with the identification of fish from seabird diets (Verstraete et al. 2020). We will be looking at the lengths one specific taxon - the Atlantic Herring *Clupea harengus*.

Copy and run the code below to load in your data. Use the object `clupea.lengths` for this problem.

```
# download fish data
fishes <- read_csv("https://zenodo.org/records/4066594/files/reference_data.csv")

# isolate Clupea harengus
clupea <- fishes |>
  filter(scientific_name == "Clupea harengus")

# get vector of lengths
# remove missing data
clupea.lengths <- clupea$fish_length_mm |>
  na.omit()
```

1. Create a histogram of these data.
2. Create a cumulative frequency plot of these data.
3. Calculate the kurtosis of these data, and make a conclusion about the kurtosis of the dataset.
4. Calculate the skewness of these data, and make a conclusion about the skewness of the dataset.
5. Create a boxplot of these data.

5.8.3 Problem 2

The Fremont Bridge in Seattle is one of the only bridges that connects the north and south sides of the city, crossing the Fremont Cut - a navigational canal between the Puget Sound and Lake Union & Lake Washington. As such, it is an important transportation corridor. We will be looking at counts of bicyclists, taken hourly for the east (predominately northbound) side over a seven year period (Weber 2019). This dataset has 56,904 records - thus, this is not something that would be easily done without the use of a program like *R*!

```
bicycles <- read_csv("https://zenodo.org/records/2648564/files/Fremont_Bridge_Hourly_Bicycle")
bike.count <- bicycles$`Fremont Bridge East Sidewalk` |>
  na.omit()
```

1. Create a histogram of these data.
2. Create a cumulative frequency plot of these data.
3. Calculate the kurtosis of these data, and make a conclusion about the kurtosis of the dataset.
4. Calculate the skewness of these data, and make a conclusion about the skewness of the dataset.
5. Create a boxplot of these data.

5.8.4 Problem 3

Take the object `bike.count` from the previous problem and perform a *log transformation* of the data. To do this, we will take the natural log of the dataset plus one. We must add 1 to every value in the dataset, because the natural log of 0 is negative infinity; adding one makes every value in the dataset quantifiable as the minimum value for any count data is 0.

Perform the following action on your computer, replacing `x` with your `bike.data` object. Name your object `log.bikes` instead of `log.x`.

```
log.x <- log(x + 1)
```

After performing the log transformation, complete the following.

1. Create a histogram of these data.
2. Create a cumulative frequency plot of these data.
3. Calculate the kurtosis of these data, and make a conclusion about the kurtosis of the dataset.
4. Calculate the skewness of these data, and make a conclusion about the skewness of the dataset.

5. Create a boxplot of these data.
6. How did the log transformation change these data?

5.8.5 Problem 4

Is there pseudoreplication in the bicycle dataset? Why or why not?

5.9 Acknowledgment

Thanks to Hernan Vargas & Riley Grieser for help in formatting this page. Additional comments provided by BIOL 305 classes.

6 Normality & hypothesis testing

6.1 Normal distributions

A *standard normal distribution* is a mathematical model that describes a commonly observed phenomenon in nature. When measuring many different kinds of datasets, the data being measured often becomes something that resembles a standard normal distribution. This distribution is described by the following equation:

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

This equation is fairly well defined by the *variance* (σ^2), the overall spread of the data, and by the *standard deviation* (σ), which is defined by the square root of the variance.




Figure 6.1: A standard normal distribution, illustrating the percentage of area found within each standard deviation away from the mean. By Ainali on Wikipedia; CC-BY-SA 3.0.

Standard normal distributions have a mean, median, and mode that are *equal*. The standard normal distribution is a *density function*, and we are interested in the “area under the curve” (AUC) to understand the relative probability of an event occurring. At the mean/median/mode, the probability on either side of the distribution is 50%. When looking at a normal distribution distribution, it is impossible to say the probability of a specific event occurring, but it is possible to state the probability of an event *as extreme or more extreme than the event observed* occurring. This is known as the p value.

6.1.1 Example in nature

An example of naturally occurring normal distributions

In order to see an example of the normal distribution in nature, we are going to examine the BeeWalk survey database from the island of Great Britain (Comont 2020). We are not interested in the bee data at present, however, but in the climatic data from when the surveys were performed.

```
beewalk <- read_csv("https://figshare.com/ndownloader/files/44726902")
```

```
Rows: 306550 Columns: 49
```

```
-- Column specification -----
```

```
Delimiter: ","
```

```
chr (30): Website.ID, Website.RecordKey, SiteName, Site.section, ViceCounty,...
```

```
dbl (19): RecordKey, established, Precision, Transect.lat, Transect.long, tr...
```

```
i Use `spec()` to retrieve the full column specification for this data.
```

```
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

Note that this is another massive dataset - 306,550 rows of data!

The dataset has the following columns:

```
colnames(beewalk)
```

| | | |
|----------------------------|------------------------|---------------------|
| [1] "RecordKey" | "Website.ID" | "Website.RecordKey" |
| [4] "SiteName" | "Site.section" | "ViceCounty" |
| [7] "established" | "GridReference" | "Projection" |
| [10] "Precision" | "Transect.lat" | "Transect.long" |
| [13] "transect.OS1936.lat" | "Transect.OS1936.long" | "transect_length" |
| [16] "section_length" | "section_grid_ref" | "H1" |
| [19] "H2" | "H3" | "H4" |

```

[22] "habitat_description" "L1"                "L2"
[25] "land_use_description" "start_time"        "end_time"
[28] "sunshine"           "wind_speed"        "temperature"
[31] "TaxonVersionKey"    "species"           "latin"
[34] "queens"             "workers"           "males"
[37] "unknown"            "Comment"           "transect_comment"
[40] "flower_visited"     "StartDate"         "EndDate"
[43] "DateType"          "Year"              "Month"
[46] "Day"               "Sensitive"         "Week"
[49] "TotalCount"

```

We are specifically interested in `temperature` to determine weather conditions. Let's see what the mean of this variable is.

```
mean(beewalk$temperature)
```

```
[1] NA
```

Hmmm... we are getting an NA value, indicating that not every cell has data recorded. Let's view `summary`.

```
summary(beewalk$temperature)
```

| Min. | 1st Qu. | Median | Mean | 3rd Qu. | Max. | NA's |
|------|---------|--------|-------|---------|-------|-------|
| 0.00 | 16.00 | 19.00 | 18.65 | 21.00 | 35.00 | 16151 |

As we can see, 16,151 rows do not have temperature recorded! We want to remove these NA rows, which we can do by using `na.omit`.

```

beewalk$temperature |>
  na.omit() |>
  mean() |>
  round(2) # don't forget to round!

```


```
[1] 18.65
```

Now we can record the mean.

Let's visualize these data using a histogram. *Note* I do not use `na.omit` as the `hist` function automatically performs this data-cleaning step!

```
hist(beewalk$temperature,breaks = 5)
```


Histogram of beewalk\$temperature



Even with only five breaks, we can see an interesting, normal-esque distribution in the data. Let's refine the bin number.

```
hist(beewalk$temperature,breaks = 40)
```

Histogram of beewalk\$temperature



With forty breaks, the pattern becomes even more clear. Let's see what a *standard normal distribution* around these data would look like.


```
# save temperature vector without NA values
temps <- beewalk$temperature |> na.omit()

mu <- mean(temps)
t.sd <- sd(temps)

# sample random values
normal.temps <- rnorm(length(temps), # sample same size vector
                      mean = mu,
                      sd = t.sd)

hist(normal.temps, breaks = 40)
```

Histogram of normal.temps



As we can see, our normal approximation of temperatures is not too dissimilar from the distribution of temperatures we actually see!

Let's see what kind of data we have for temperatures:

```
# load moments library
library(moments)

skewness(temps)
```

```
[1] 0.02393257
```


Data do not have any significant skew.

```
kurtosis(temps)-3
```

```
[1] 0.3179243
```

Data do not show any significant kurtosis.

6.1.2 Effect of sampling


How sample size effects statistical calculations

Oftentimes, we will see things approach the normal distribution as we collect more samples. We can model this by subsampling our temperature vector.

```
# make reproducible
set.seed(1839)


sub.temps <- sample(temps,
                    size = 10,
                    replace = FALSE)

hist(sub.temps, main = "10 samples")
```




With only ten values sampled, we do not have much of a normal distribution. Let's up this to 100 samples.

```
sub.temps <- sample(temps,  
                    size = 100,  
                    replace = FALSE)  
  
hist(sub.temps, main = "100 samples", breaks = 10)
```



Now we are starting to see more of a normal distribution! Let's increase this to 1000 temperatures.

```
sub.temps <- sample(temps,  
                    size = 1000,  
                    replace = FALSE)  
  
hist(sub.temps, main = "1000 samples", breaks = 40)
```



Now the normal distribution is even more clear. As we can also see, the more we sample, the more we approach the true means and distribution of the actual dataset. Because of this, we can perform experiments and observations of small groups and subsamples and make inferences about the whole, given that most systems naturally approach statistical distributions like the normal!

6.2 Hypothesis testing

Since we can define specific areas under the curve within these distributions, we can look at the percentage of area within a certain bound to determine how likely a specific outcome would be. Thus, we can begin to test what the *probability of observing an event* is within a theoretical, probabilistic space. A couple of important conceptual ideas:

1. We may not be able to know the probability of a specific event, but we can figure out the probability of events more extreme or less extreme as that event.
2. If the most likely result is the mean, then the further we move away from the mean, the less likely an event becomes.
3. If we look *away* from the mean at a certain point, then the area represents the chances of getting a result *as extreme or more extreme than what we observe*. This probability is known as the *p* value.

Once we have a *p* value, we can make statements about the event that we've seen relative to the overall nature of the dataset, but we do not have sufficient information to declare if this result is *statistically significant*.

6.2.1 Critical Values - α

In order to determine if something is significant, we compare things to a *critical value*, known as α . This value is traditionally defined as 0.05, essentially stating that we deem an event as significant if 5% or fewer of observed or predicted events are as extreme or more extreme than what we observe.

Your value should always set your α critical value before you do your experiments and analyses.

Our critical value of α represents our criterion for *rejecting the null hypothesis*. We set our α to try to minimize the chances of error.

Type I Error is also known as a **false-positive**, and is when we **reject the null hypothesis when the null is true**.

Type II Error is also known as a **false-negative**, and is when we **support the null hypothesis when the null is false**.

By setting an α , we are creating a threshold of probability at which point we can say, with confidence, that results are different.

6.2.2 Introduction to p values

Let's say that we are looking at a dataset defined by a standard normal distribution with $\mu = 0$ and $\sigma = 1$. We draw a random value, x , with $x = 1.6$. What is the probability of drawing a number this extreme or more extreme from the dataset?

First, let's visualize this distribution:

[Click here to view plot creation code.](#)

```
###THIS WILL TAKE A WHILE TO RUN###

# create gigantic normal distribution dataset
# will be essentially normal for plotting
# rnorm gets random values
x <- rnorm(100000000)


# convert to data frame
x <- as.data.frame(x)
# rename column
colnames(x) <- c("values")

# thank you stack overflow for the following
```

```
# Creating density plot
p = ggplot(x,
           aes(x = values)
           ) +
  # generic density plot, no fill
  geom_density(fill="lightblue")

# Building shaded area
# create new plot object
p2 <- p + # add previous step as a "backbone"
  # rename axes
  geom_vline(xintercept = 1.6) +
  xlab("Test Statistic") +
  ylab("Frequency (Probability)") +
  # make it neat and tidy
  theme_classic()

# plot it
# can use ggsave function to save
plot(p2)
```



Above, the solid black line represents x , with the illustrated standard normal distribution being filled in blue.

Let's see how much of the area represents values *as extreme or more extreme* as our value x .


[Click here](#) for code to create figure.

```
### THIS WILL TAKE A WHILE TO RUN ###

# Getting the values of plot
# something I wasn't familiar with before making this!
d <- ggplot_build(p)$data[[1]]

# Building shaded area
# create new plot object
p2 <- p + # add previous step as a "backbone"
  # add new shaded area
  geom_area(data = subset(d, x < 1.6), # select area
            # define color, shading intensity, etc.
            aes(x=x,y=y), fill = "white", alpha = 1) +
  # add value line
  geom_vline(xintercept = 1.6, colour = "black",
             linetype = "dashed",linewidth = 1) +
  # rename axes
  xlab("Test Statistic") +
  ylab("Frequency (Probability)") +
  # make it neat and tidy
  theme_classic()

# plot it
# can use ggsave function to save
plot(p2)
```



Now we can see that it is only a portion of the distribution *as extreme or more extreme* than the value we placed on the graph. The area of this region is our p value. This represents the *probability of an event as extreme or more extreme occurring* given the random variation observed in the dataset or in the distribution approximating the dataset. This is the value we compare to α - our threshold for rejecting the null hypothesis - to determine whether or not we are going to reject the null hypothesis.

Let's look at the above graph again, but let's visualize a two-tailed α around the mean with a 95% confidence interval. First, we need to get the Z scores for our $\alpha = 0.05$, which we calculate by taking $\frac{\alpha}{2}$ to account for the two tails. (Two tails essentially meaning we reject the null mean if we see things *greater than* or *less than* our expected value to a significant extent). We can calculate Z scores using `qnorm`.

```
low_alpha <- qnorm(0.025) # looks left
hi_alpha <- qnorm(0.975) # looks left

print(paste0("Our Z scores are: ",
             round(low_alpha,2),
             " & ",
             round(hi_alpha,2)))
```

```
[1] "Our Z scores are: -1.96 & 1.96"
```

The above values make sense, given the distribution is symmetrical. Our above dashed line is as $Z = 1.96$, which means we should have a $p = 0.05$, so the dashed line should be *closer* to the mean than our cutoffs.


[Click here for code to create figure.](#)

```
### THIS WILL TAKE A WHILE TO RUN ###

# Getting the values of plot
# something I wasn't familiar with before making this!
d <- ggplot_build(p)$data[[1]]

# Building shaded area
# create new plot object
p2 <- p + # add previous step as a "backbone"
  # add new shaded area
  geom_area(data = subset(d, x < 1.96), # select area
            # define color, shading intensity, etc.
            aes(x=x,y=y), fill = "white", alpha = 1) +
  # add value line
  geom_vline(xintercept = 1.96, colour = "black",
             linetype = "dashed",linewidth = 1) +
  geom_vline(xintercept = low_alpha, colour = "red",
             linetype = "solid",linewidth = 1) +
  geom_vline(xintercept = hi_alpha, colour = "red",
             linetype = "solid",linewidth = 1) +
  # rename axes
  xlab("Test Statistic") +
  ylab("Frequency (Probability)") +
  # make it neat and tidy
  theme_classic()

# plot it
# can use ggsave function to save
plot(p2)
```

Exactly as we calculated, we see that our $p < \alpha$ and thus we do not see an area (in blue) less than the area that would be further than the mean as defined by the red lines.

6.2.3 Calculating a Z score

When we are trying to compare our data to a normal distribution, we need to calculate a Z score for us to perform the comparison. A Z score is essentially a measurement of the number of standard deviations we are away from the mean on a standard normal distribution. The equation for a Z score is:

$$Z = \frac{\bar{x} - \mu}{\frac{\sigma}{\sqrt{n}}}$$

Where \bar{x} is either a sample mean or a sample value, μ is the population mean, σ is the population standard deviation and n is the number of individuals in your sample (1 if comparing to a single value).

We can calculate this in *R* using the following function:

```
zscore <- function(xbar, mu, sd.x, n = 1){
  z <- (xbar - mu)/(sd.x/sqrt(n))
  return(z)
}
```

NOTE that is the above isn't working for you, *you have a mistake somewhere in your code*. Try comparing - character by character - what is listed above to what you have.

Let's work through an example, where we have a sample mean of 62 with 5 samples compared to a sample mean of 65 with a standard deviation of 3.5.

```
Z <- zscore(xbar = 62, # sample mean
            mu = 65, # population mean
            sd.x = 3.5, # population standard deviation
            n = 5) # number in sample

print(Z)
```

```
[1] -1.91663
```

Now, we can calculate the p value for this Z score.

```
pnorm(Z)
```

```
[1] 0.0276425
```

After rounding, we get $p = 0.03$, a p that is significant as it is less than our overall α .

6.2.4 Calculating the p value

We have two different methods for calculating a p value:

6.2.4.1 Comparing to a z table

We can compare the z value we calculate to a z table, such as the one at ztable.net. On this webpage, you can scroll and find tables for positive and negative z scores. *Note* that normal distributions are symmetrical, so you can also transform from negative to positive to get an idea of the area as well. Given that a normal distribution is centered at 0, a z score of 0 will have a p value of 0.50.

On the z tables, you will find the tenths place for your decimal in the rows, and then go across to the columns for the hundredths place. For example, go to the website and find the p value for a z score of -1.33 . You should find the cell marked 0.09176. *Note* the website uses naked decimals, which we do not use in this class.

For values that aren't on the z table, we can approximate its position between different points on the z table or, if it is extremely unlikely, denote that $p < 0.0001$.

6.2.4.2 Using *R*

In *R*, we can calculate a p value using the function `pnorm`. This function uses the arguments of `p` for our p value, `mean` for the mean of our distribution, `sd` for the standard deviation of our distribution, and also information on whether we want to log-transform p or if we are testing a specific hypothesis (lower tail, upper tail, or two-tailed). The function `pnorm` defaults to a standard normal distribution, which would be a z score, but it can also perform the z transformations for us if we define the mean and standard deviation.

For example, if we have a z score of -1.33 :

```
pnorm(-1.33)
```

```
[1] 0.09175914
```

As we can see, we get the same result as our z table, just with more precision!

There are other functions in this family as well in *R*, including `dnorm` for quantiles, `qnorm` for determining the z score for a specific p value, and `rnorm` for getting random values from a normal distribution with specific dimensions. For now, we will focus on `pnorm`.

6.2.5 Workthrough Example

When this class was being designed, [Hurricane Milton](#) was about to make contact with Florida. Hurricane Milton is considered one of the strongest hurricanes of all time, so we can look at historical hurricane data to determine just how powerful this storm really was. We can get information on maximum wind speeds of all recorded Atlantic hurricanes as of 2024 from Wikipedia.

[Click here](#) for details on creating the file for this walkthrough.

```
# https://stackoverflow.com/questions/59481608/how-to-remove-the-characters-from-a-string-and-
# cocatenating wikipedia files.
hurricanes1 <- list.files("~/Downloads/",pattern = "List_of_Category_1_Atlantic*")

speedlist_cat1 <- NULL
for(j in 1:length(hurricanes1)){
  test <- read_csv(paste0("~/Downloads/", hurricanes1[j]))
  test <- test[-1,] # remove extra header
  for(i in 1:nrow(test)){
    speed <- test$`Peak intensity...3`[i]
```

```

    speed2 <- strsplit(speed, " ")[[1]][1]
    speed2 <- gsub("\\D+", "", speed2)
    speedlist_cat1 <- c(speedlist_cat1, speed2)
  }
}

speedlist_cat1 <- as.numeric(speedlist_cat1)

# cocatenating wikipedia files.
hurricanes2 <- list.files("~/Downloads/", pattern = "List_of_Category_2_Atlantic*")

speedlist_cat2 <- NULL
for(j in 1:length(hurricanes2)){
  test <- read_csv(paste0("~/Downloads/", hurricanes2[j]))
  test <- test[-1,] # remove extra header
  for(i in 1:nrow(test)){
    speed <- test$`Peak intensity...3`[i]
    speed2 <- strsplit(speed, " ")[[1]][1]
    speed2 <- gsub("\\D+", "", speed2)
    speedlist_cat2 <- c(speedlist_cat2, speed2)
  }
}

speedlist_cat2 <- as.numeric(speedlist_cat2)

# cocatenating wikipedia files.
hurricanes3 <- list.files("~/Downloads/", pattern = "List_of_Category_3_Atlantic*")

speedlist_cat3 <- NULL
for(j in 1:length(hurricanes3)){
  test <- read_csv(paste0("~/Downloads/", hurricanes3[j]))
  test <- test[-1,] # remove extra header
  for(i in 1:nrow(test)){
    speed <- test$`Peak intensity...3`[i]
    speed2 <- strsplit(speed, " ")[[1]][1]
    speed2 <- gsub("\\D+", "", speed2)
    speedlist_cat3 <- c(speedlist_cat3, speed2)
  }
}

speedlist_cat3 <- as.numeric(speedlist_cat3)
speedlist_cat3 <- na.omit(speedlist_cat3) |>

```

```

as.numeric()

# cocatenating wikipedia files.
hurricanes4 <- list.files("~/Downloads/",pattern = "List_of_Category_4_Atlantic*")

speedlist_cat4 <- NULL
for(j in 1:1){
  test <- read_csv(paste0("~/Downloads/", hurricanes4[j]))
  test <- test[-1,] # remove extra header
  for(i in 1:nrow(test)){
    speed <- test$`Peak intensity...3`[i]
    speed2 <- strsplit(speed," ")[[1]][1]
    speed2 <- gsub("\\D+", "", speed2)
    speedlist_cat4 <- c(speedlist_cat4,speed2)
  }
}

speedlist_cat4 <- as.numeric(speedlist_cat4)

# cocatenating wikipedia files.
hurricanes5 <- list.files("~/Downloads/",pattern = "List_of_Category_5_Atlantic*")

speedlist_cat5 <- NULL
for(j in 1:1){
  test <- read_csv(paste0("~/Downloads/", hurricanes5[j]))
  test <- test[-1,] # remove extra header
  for(i in 1:nrow(test)){
    speed <- test$`Peak intensity...4`[i]
    speed2 <- strsplit(speed," ")[[1]][1]
    speed2 <- gsub("\\D+", "", speed2)
    speedlist_cat5 <- c(speedlist_cat5,speed2)
  }
}

speedlist_cat5 <- as.numeric(speedlist_cat5)

speedlist <- c(speedlist_cat1,speedlist_cat2,speedlist_cat3,
              speedlist_cat4,speedlist_cat5)

speedlist <- as.data.frame(speedlist)

colnames(speedlist) <- "Hurricane_Windspeed"

```

```
write_csv(speedlist,path = "~/Documents/GitHub/biol105_unk/assignments/hurricane_speeds.csv")
```

```
hurricanes <- read_csv("https://raw.githubusercontent.com/jacobccooper/biol305_unk/main/assignm
```

```
Rows: 889 Columns: 1
```

```
-- Column specification -----
```

```
Delimiter: ","
```

```
dbl (1): Hurricane_Windspeed
```

```
i Use `spec()` to retrieve the full column specification for this data.
```

```
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

Above, we have loaded a .csv of one column that has all the hurricane speeds up to 2024. Hurricane Milton is the last row - the most recent hurricane. Let's separate this one out. We will use [,], which defines [rows,columns] to subset data.

```
milton <- hurricanes$Hurricane_Windspeed[nrow(hurricanes)]
```


```
other_hurricanes <- hurricanes[-nrow(hurricanes),]
```

We want to compare the windspeed of Milton (180 mph) to the overall distribution of hurricane speeds. We can visualize this at first.

```
# all hurricanes
```

```
hist(hurricanes$Hurricane_Windspeed)
```

Histogram of hurricanes\$Hurricane_Windspeed



Windspeeds are more towards the lower end of the distribution, with strong storms being rarer.

For the sake of this class, we will assume we can use a normal distribution for these data, but if we were doing an official study we would likely need to use a non-parametric test (we will cover these later, but they cover non-normal data).

```
mu <- mean(other_hurricanes$Hurricane_Windspeed)
mu
```

```
[1] NA
```

Hmmm... we need to use `na.omit` to be sure we do this properly.

```
other_hurricanes_windspeed <- na.omit(other_hurricanes$Hurricane_Windspeed)
mu <- mean(other_hurricanes_windspeed)
mean(other_hurricanes_windspeed)
```

```
[1] 102.5254
```

Next, we need the standard deviation.

```
sd.hurricane <- sd(other_hurricanes_windspeed)
sd.hurricane
```

```
[1] 23.08814
```

Now, we can calculate our Z value.

```
Z <- (milton - mu)/sd.hurricane
Z
```

```
[1] 3.355603
```

How significant is this?

```
pnorm(Z)
```

```
[1] 0.999604
```

This is greater than 0.5, so we need to do $1 - p$ to figure things out.

```
1 - pnorm(Z)
```

```
[1] 0.000395961
```

This rounds to 0.0004, which means that this is an *extremely* strong hurricane.

6.2.5.1 Non-normality, for those curious

We can do a Shapiro-Wilk test of normality to see if this dataset is normal.

```
shapiro.test(other_hurricanes_windspeed)
```

```
Shapiro-Wilk normality test
```

```
data:  other_hurricanes_windspeed  
W = 0.88947, p-value < 2.2e-16
```

A $p < 0.05$ indicates that these data are *non-normal*.

We can do a Wilcoxon-Test since these data are extremely non-normal.

```
wilcox.test(other_hurricanes_windspeed,  
            milton)
```

```
Wilcoxon rank sum test with continuity correction
```

```
data:  other_hurricanes_windspeed and milton  
W = 6.5, p-value = 0.08678  
alternative hypothesis: true location shift is not equal to 0
```

Using non-normal corrections, we find that this is *not* an extremely strong hurricane, but it is near the upper end of what we would consider “normal” under historical conditions. Still an extremely bad hurricane!

6.3 Confidence Intervals

Because we can figure out the probability of an event occurring, we can also calculate *confidence intervals*. A *confidence interval* provides a range of numbers around a value of interest that indicates where we believe the mean of a population lies and our confidence that it lies within that range. **Note** that nothing is ever 100% certain, but this helps us determine where a mean is and demonstrates our confidence in our results.

Specifically, if the tails of the distribution, our α , are 0.05, then we have an area of 0.95 around the mean where we do not reject results. Another perspective on this area is that we can say with 95% certainty that a mean is within a certain area, and that if values fall within that confidence area then we do not reject the null hypothesis that the means are equal.

We will cover several different ways to calculate confidence intervals, but for normal distributions, we use the following equation, with the 0.95 confidence interval shown as an example:

$$CI = \bar{x} \pm Z_{1-\frac{\alpha}{2}} \cdot \frac{\sigma}{\sqrt{n}}$$

This interval gives us an idea of where the mean should lie. For example, if we are looking at the aforementioned `beewalk` temperature data, we can calculate a 0.95 confidence interval around the mean.

```
temps <- na.omit(beewalk$temperature)

xbar <- mean(temps)
n <- length(temps)
sdtemp <- sd(temps)
# Z for 0.95 as P, so for 0.975, 0.025
# get value from P
Z <- qnorm(0.975)

CI <- Z*(sdtemp/sqrt(n))

CI
```

```
[1] 0.01458932
```

We have a very narrow confidence zone, because we have so many measurements. Let's round everything and present it in a good way.

If I want numbers to show up *in text* in RMarkdown, I can add code to a line of plain text using the following syntax:

```
# DO NOT RUN
# Format in plaintext
`r xbar`
```

Typing that into the `plaintext` should render as the following: ``r xbar``. Then I can also type my answer as follows:

The 95% Confidence Interval for the mean for this temperature dataset is ``r round(xbar,2)` \pm `r round(CI,2)``.

Figure 6.2: This is the “coded” version of the text below. Compare the above window to the text below this image.

Typing that into the plain text should render as the following: 18.645963. Then I can also type my answer as follows:

The 95% Confidence Interval for the mean for this temperature dataset is:

```
print(paste(round(xbar, 2),
            "+/-",
            round(CI, 2)))
```

```
[1] "18.65 +/- 0.01"
```

Note that the above is rounded to two decimal places to illustrative purposes ONLY, and should be rounded to one decimal place if it was a homework assignment because the original data has only one decimal place.

6.4 Homework: Normal distributions

Submit one `html` file, as created in *RStudio*, with you answers to the following questions. For maximum clarity, create headings to separate your problems. (Remember, a header can be invoked by placing `#` in front of a line of text. For example: the header here is written as `# Homework: Chapter 8`).

6.4.1 Z to p

For each of the following, report the corresponding p value for the z score.

- A. 2.3
- B. 6.4
- C. -1.5
- D. -9.2
- E. 3.2

6.4.2 P to z

For the following, report the corresponding z score for the stated p value. For this question, assume every value is “looking left” (`lower.tail = T`).

- A. 0.96
- B. 0.23
- C. 0.14
- D. 0.0036
- E. 0.47
- F. 0.62

6.4.3 Confidence intervals

For the following scenarios, calculate the confidence interval for the mean.

- A. $\bar{x} = 7.6$, $\sigma = 1.3$, $n = 5$, $CI = 99\%$
- B. $\bar{x} = -4.9$, $\sigma = 0.7$, $n = 6$, $CI = 95\%$
- C. $\bar{x} = 2.1$, $\sigma = 0.4$, $n = 10$, $CI = 90\%$
- D. $\bar{x} = 0.1$, $\sigma = 3$, $n = 12$, $CI = 95\%$

6.4.4 Italian climate - spring

As part of a study on how climate affects the weights of Alpine Ibex, researchers recorded the maximum temperature in the spring at their study site (Brambilla et al. 2024). The researchers find another site with an average spring maximum temperature of 6.2°C, and claim that it is warmer than their study site. Are they correct?

Evaluate their claims by:

1. Calculating a z -score for their temperature (6.2).
2. Calculating a p -value for their temperature (6.2).
3. Reporting a 95% confidence interval for the mean for the temperatures reported in their dataset.
4. In a single sentence, state whether you find this temperature to be significantly warmer than the mean or not, stating what evidence you are using to make this conclusion.

The data for this problem will be in the object `spring_temps` below:

```
# load weather data
weather <- read_csv("https://zenodo.org/records/12749273/files/20240505_env_var_dataset.csv")

# isolate spring temperatures
spring_temps <- weather$SpringTmax |>
  # remove NA values
  na.omit() |>
  # ensure data are in correct numeric format
  as.numeric()
```

6.4.5 Italian climate - Summer

Using the same dataset (**weather**) as the previous question, isolate the data for the summer (column **SummTmax**). The researchers find the average temperature of another site is 13.6°C, and claim this is warmer than their research site. Are they correct? As with the previous question, complete the following:

1. Calculating a z -score for their temperature (6.2).
2. Calculating a p -value for their temperature (6.2).
3. Reporting a 95% confidence interval for the mean for the temperatures reported in their dataset.
4. In a single sentence, state whether you find this temperature to be significantly warmer than the mean or not, stating what evidence you are using to make this conclusion.

Note for this question, you will have to isolate the column and perform all of the analysis steps yourself.

7 Transforming data

7.1 Importance of normality Data

For this walkthrough, we will use two datasets - one with normal data, and one with non-normal data.

Our normal dataset is drawn from measurements of birds in western Ecuador, collected by several friends of Dr. Cooper, where we will be specifically looking at measurements of the Plain Antvireo *Dysithamnus mentalis* (Lele et al. 2022).

Our non-normal dataset counts the number of job vacancies offering an hourly wage by region in the United Kingdom for different time periods (Urban Big Data Centre 2025).

```
library(tidyverse)

## Bird data

birds <- read_csv("https://zenodo.org/records/6511860/files/Lele_et_al._2022_BITR-21-373_raw")

antvireo <- birds |>
  filter(SPECIES == "PLAIN ANTVIREO") |>
  # remove abnormal birds
  filter(BILL.LENGTH > 5)

antvireo.bill <- antvireo$BILL.LENGTH |>
  na.omit()

## Job advert data

jobs <- read_csv("https://zenodo.org/records/14771706/files/adzuna_month_ttwa_vacancies_with")

# take a subset
# whole dataset is too large!
hourly.jobs <- jobs$vacancies_offering_an_hourly_wage[1:100]
```

7.2 Testing if data are normal


There are two major methods we can use to see if data are normally distributed.

7.2.1 Histograms

One of the first things that you should do is look at a histogram of your data. Histograms will help you spot any large data irregularities, and can help you get an idea of whether you should expect your data to be non-normally distributed.

First, let's look at our `antvireo.bill` data:

```
hist(antvireo.bill)
```




These data appear relatively normal.

Next, let's look at the `hourly.jobs` data.

```
hist(hourly.jobs)
```

Histogram of hourly.jobs



These data appear highly non-normal; most values are low, and we have an extreme right skew.


7.2.2 QQ Plots

Visual method for assessing normality

One way to see if data are normal is to use a [QQ plot](#). These plot data quantiles to theoretical quantiles to see how well they align, with a perfectly normal distribution having a completely linear QQ plot. Let's look at these with our `antvireo.bill` data.

```
qqnorm(antvireo.bill)
```

Normal Q-Q Plot




As we can see above, the data are roughly linear, which means are data appear normal. The “stairsteps” are from the accuracy in measuring temperature, which was likely rounded and thus created a distribution that is not completely continuous.

Now, lets check how our `hourly.jobs` data look:

```
qqnorm(hourly.jobs)
```

Normal Q-Q Plot



As we can see, these data are not very linear, suggesting that the data are highly non-normal.

7.2.3 Shapiro-Wilk test

Statistical method for assessing normality

Another way to test for normality is to use a Shapiro-Wilk test of normality. We will not get into the specifics of this distribution, but this tests the null hypothesis that data originated in a normal distribution, with the alternative hypothesis that the data originated in a non-normal distribution.

NOTE that the Shapiro-Wilk test does not perform well with extremely large datasets, which may be the result of [model overfitting](#).

This test uses an $\alpha = 0.05$, and we *reject the null hypothesis* if our $p < \alpha$, with p representing the probability of observing something as extreme or more extreme than the result we observe. If we reject the null hypothesis, our data are non-normal and require transformation. If we accept the null hypothesis, we can proceed with treating our data as normally distributed.

Let's look at the `antvireo.bill` data:

```
shapiro.test(antvireo.bill)
```

```
Shapiro-Wilk normality test
```

```
data:  antvireo.bill  
W = 0.97513, p-value = 0.2984
```

As we expected from our `qqnorm` plot, our data for the Plain Antvireo bills are normally distributed, and we accept the null hypothesis as $p > \alpha$ with $0.30 > 0.05$.

Next, let's do a Shapiro-Wilks test on our `hourly.jobs` data.

```
shapiro.test(hourly.jobs)
```

```
Shapiro-Wilk normality test
```

```
data:  hourly.jobs  
W = 0.39393, p-value < 2.2e-16
```

As we can see, our $p < 0.001$, therefore we reject the null hypothesis and conclude the `hourly.jobs` data are non-normal.

7.3 Transforming data

There are multiply different transformations that can be performed on different datasets. In this class, we will focus on three transformations:

1. Log transformations, some of the most common transformations
2. Square transformations
3. Square-root transformations

Note: there are many other types of transformations, we just don't go in depth with them here.

When transforming a dataset, you must perform the mathematical function across all values within the dataset. You can then assess whether this dataset is normally distributed and determine whether you can proceed with statistical analyses that you would perform on normal distributions.

The process for all transformations is the same, so this walkthrough will only perform one: the log transformation.

7.3.1 Log transformation example

Walkthrough of a transformation for normality

We know that our `hourly.jobs` dataset is non-normally distributed, so we need to perform a transformation to see if we can achieve normality. Here, we will perform a log transformation. In *R*, the function `log` performs a *natural log* (\ln) by default. Because $\ln(0)$ is *−Infinity*, folks will often add +1 to their entire dataset to avoid zero values. This is not necessary if your entire dataset is non-zero and positive.

```
# check if any values are below zero
# function will return # vals at or below 0
sum(hourly.jobs <= 0)
```

```
[1] 0
```


There are no values less than or equal to zero in this dataset. However, if we wanted to add one to them, we could do so by just taking `hourly.jobs + 1` and creating a new object.

Next, let's perform the log transformation:

```
log.hourly.jobs <- hourly.jobs |>
  log()
```


Very straightforward! How do these data look?

```
hist(log.hourly.jobs)
```



These data appear far more normal. What about on a qqplot?

```
qqnorm(log.hourly.jobs)
```



The line is much straighter, as is expected for a normal distribution.

Now, for the Shapiro-Wilk test:

```
shapiro.test(log.hourly.jobs)
```

Shapiro-Wilk normality test

```
data: log.hourly.jobs  
W = 0.98204, p-value = 0.1911
```

We now have a $p > \alpha$ with $0.19 > 0.05$, indicating that these data are now normally distributed. We can now proceed with our analyses!

Note, however, that you must *back transform* your data after a transformation!

A back transformation of $\ln(x)$ is e^x . In *R*, e is found by entering `exp(1)`. Thus we would do the following if looking at the mean from the logs:

```
log.mean <- mean(log.hourly.jobs)  
  
log.mean |> round(2)
```

```
[1] 5.05
```

And back transformed:

```
transformed.log.mean <- exp(1)^log.mean  
transformed.log.mean |> round(0)
```

```
[1] 156
```

7.4 Homework: Transforming data

This homework assignment is designed to help you become more familiar with manipulating data and with transforming data.

7.4.1 Analyzing data

Using the `bird` dataset above, pick two other species of birds and do the following:

1. Provide information on the mean, median, mode, kurtosis, and skewness
2. Create a histogram plot
3. Determine if the data are normally distributed
4. If data are *not* normally distributed, see if a square, log, or square root transformation will make the data normal.
5. Provide an overall assessment of the data's normality.

7.4.2 Transforming Data

Using the below newly created data object, perform all the same steps as you did for the bird data above.

```
hourly.jobs.homework <- jobs$vacancies_offering_an_hourly_wage[200:300]
```

7.4.3 Submitting assignment

Submit your homework assignment as an `html` file on *Canvas*.

8 Probability distributions

8.1 Probability distributions

We rely on multiple different probability distributions to help us understand what probable outcomes are for a specific scenario. All of the tests that we are performing are comparing our results to what we would expect under perfectly random scenarios. For example, if we are flipping a coin, we are interested in whether the observation we have of the flips on our coin matches our expectation given the probability of getting heads or tails on a perfectly fair coin. While it is possible to get all heads or all tails on a coin flip, it is highly unlikely and may lead us to believe we have an unfair coin. The more trials we perform, the more confident we can be that our coin is atypical.

We perform similar comparisons for other distributions. If we are comparing sets of events, we can look at the probability of those events occurring if events are occurring randomly. If we are comparing counts, we can compare our counts to our expectation of counts if events or subjects are distributed randomly throughout the matrix or whether two sets of counts are likely under the same sets of assumptions.

Remember, for our specific tests, we are setting an α value in advance (traditionally 0.05, or 5%) against which we compare our p value, with p representing the probability of observing an event *as extreme* or *more extreme* than the event we observe given a specific probability distribution.

Previously, we talked about the *normal distribution*, which is used to approximate a lot of datasets in nature. However, several other probability distributions are also useful for biological systems, which are outlined here.

8.2 Review: Basic statistics

If we are looking at things that could be one *or* the other thing, we add the probabilities. For example, the odds of getting a four or a six on a 20 sided die are $\frac{1}{20} + \frac{1}{20} = \frac{2}{20} = \frac{1}{10}$.

If we are looking at things that could be one thing *and* another thing, we multiply the probabilities together. For example, if we roll two six-sided die and we are looking at the odds of a 2 or a 4 on both die, the odds are $(\frac{1}{6} + \frac{1}{6}) \cdot (\frac{1}{6} + \frac{1}{6}) = \frac{1}{3} \cdot \frac{1}{3} = \frac{1}{9}$.

8.3 Binomial distribution

A *binomial distribution* is one in which only two outcomes are possible - often coded as 0 and 1 and usually representing failure and success, respectively. The binomial is described by the following function:

$$p(x) = \binom{n}{x} p^x (1-p)^{n-x}$$

where n = number of trials, x = the number of successes, and p = the probability of a success under random conditions.

In *R*, the binomial distribution is represented by the following functions:

- `dbinom`: the density of a binomial distribution
- `pbinom`: the distribution function, or the probability of a specific observation
- `qbinom`: the value at which a specific probability is found (the *quantile function*)
- `rbinom`: generates random values according to a binomial.

8.3.1 Binomial examples

Binomial example walkthrough

Let's see what this looks like. Let's consider a scenario where we flip a coin 10 times and get 9 heads. How likely is this outcome?

```
x <- pbinom(q = 9, # number successes, 9 heads
            size = 10, # number of trials, 10 flips
            prob = 0.5) # probability with a fair coin

round(x, 4)
```

```
[1] 0.999
```

NOTE that the trailing 0 is dropped, such that the real answer is 0.9990. However, we mentioned before that the p value should be the probability of a result as extreme or more extreme, meaning that it should *always* be less than 0.5. If we are reporting a value of *greater* than 0.5, then we are comparing to the upper tail of the distribution. For a one-tailed α of 0.05, this would mean that we are looking for a value *greater than* 0.95 ($1 - \alpha$).

So, our real p is:

```
1 - round(x,4)
```


```
[1] 0.001
```

Again, the trailing zero is missing. Given that $p < \alpha$, we *reject the null hypothesis* that this is a fair coin.

How does this distribution look?

```
# number of successes
# start at 0 for no heads
x <- 0:10
# cumulative probability to left of outcome
y <- pbinom(x,
            size = 10,
            prob = 0.5,
            lower.tail = T)


# cumulative probability of results to the left
plot(x,
     y,
     type="l") # line plot
```



What about if we always have p less than 0.5 to reflect two tails?


```
# any value greater than 0.5 is subtracted from 1
y[y > 0.5] <- 1 - y[y > 0.5]
```

```
plot(x,
     y,
     type="l")
```




What if we do this with a bigger dataset, like for 50 flips?

```
# number of successes
# start at 0 for no heads
x <- 0:50
# cumulative probability to left of outcome
y <- pbinom(x,
            size = length(x),
            prob = 0.5,
            lower.tail = T)

# any value greater than 0.5 is subtracted from 1
y[y > 0.5] <- 1 - y[y > 0.5]

plot(x,
     y,
     type="l")
```



As we increase the number of flips, we can see that the probability of success forms a *normal distribution* centered on the outcome given the default probability. Thus, as we deviate from our *expected outcome* (initial probability multiple by the number of trials), then our results become less likely.

8.3.2 Binomial exact tests

Binomial exact test walkthrough

We can perform exact binomial tests by using the *R* function `binom.test`. This is a built in function within *R*. This test requires the following arguments:

- **x**: number of successes (success = outcome of interest)
- **n**: number of trials (number of events)
- **p**: probability of success in a typical situation (*i.e.*, for a fair coin, this is 50%)
- **alternative**: the hypothesis to be tested, whether `two.sided`, `greater`, or `less`.
- **conf.level** is the *confidence level* to be returned; default is 95%.

Let's say you flip a coin ten times, randomly assigning one side as a "success" and one side as a "failure". We do ten flips, and get 3 "successes". How likely is this outcome?

```
binom.test(x = 3,    # three successes
           n = 10,   # ten flips
           p = 0.5)  # 50% chance on fair coin
```

Exact binomial test

```
data: 3 and 10
number of successes = 3, number of trials = 10, p-value = 0.3438
alternative hypothesis: true probability of success is not equal to 0.5
95 percent confidence interval:
 0.06673951 0.65245285
sample estimates:
probability of success
                0.3
```

Now let's say we do 1000 flips, and we get 300 successes.

```
binom.test(x = 300,
           n = 1000,
           p = 0.5)
```

Exact binomial test

```
data: 300 and 1000
number of successes = 300, number of trials = 1000, p-value < 2.2e-16
alternative hypothesis: true probability of success is not equal to 0.5
95 percent confidence interval:
 0.2717211 0.3294617
sample estimates:
probability of success
                0.3
```

As we can see, both of these return a confidence interval among other things. If we save the object, we can access these “slots” of data using the \$ character.

```
binom_result <- binom.test(x = 3,
                          n = 10,
                          p = 0.5)

binom_result$p.value |> round(2)
```

```
[1] 0.34
```

```
binom_result$conf.int
```

```
[1] 0.06673951 0.65245285  
attr(,"conf.level")  
[1] 0.95
```

This test is easily implemented, but always double check and make sure you are setting it up correctly.

8.4 Poisson distribution

The *Poisson distribution* is used to reflect random count data. Specifically, the Poisson is used to determine if success events are *overdispersed* (i.e., regularly spaced), *random*, or *underdispersed* (i.e., clustered). The Poisson introduces the variable *lambda* (λ) which represents the mean (μ) and the variance (σ^2), which are equal in a Poisson distribution. A Poisson distribution is described by the following function:

$$p(x) = \frac{\lambda^x e^{-\lambda}}{x!}$$


The Poisson is represented by the following functions in *R* which closely resemble the functions for the normal and binomial distributions:

- `dpois`: the log density function
- `ppois`: log distribution (probability) function
- `qpois`: quantile function
- `rpois`: random values from a Poisson.

8.4.1 Poisson example

Let's look at the probability of 0 to 10 successes when we have our $\lambda = 1$.

```
x <- 0:10  
y <- ppois(x, lambda = 1)  
  
# any value greater than 0.5 is subtracted from 1  
y[y > 0.5] <- 1 - y[y > 0.5]  
  
plot(x, y, type="l")
```



As we can see, the probability of rare events is high, whereas the probability quickly decreases as the number of successes increases.

8.4.2 Poisson test

Much like the Binomial Distribution and its `binom.test`, we can use `poisson.test` to analyze data via a Poisson Distribution. This command uses the arguments:

- `x`: number of events of interest
- `T`: time base (if for an event count)
- `r`: hypothesized rate or ratio
- `alternative` and `conf.level` are the same as for `binom.test`

We will not often use the `poisson.test` in this class, but it is good to be aware of.

8.5 Cumulative Probabilities

With both the Binomial and the Poisson, we can calculate cumulative probabilities. Both of these require the density (`d`) versions of the arguments.

8.5.1 Binomial cumulative

Obtaining cumulative probability in a binomial distribution

Let's say we have ten trials with three successes and a base probability of 0.5. We can calculate the probability to the *left* by using the following:

```
pbinom(q = 3,  
       size = 10,  
       prob = 0.5)
```

```
[1] 0.171875
```

As we can see, this is ~17.19%. Now let's try using `dbinom`. This command gives us the value at an individual bin, given that it is a more discrete distribution for these smaller sample sizes.

```
dbinom(x = 0:3,  
       size = 10,  
       prob = 0.5)
```

```
[1] 0.0009765625 0.0097656250 0.0439453125 0.1171875000
```

Above, we can see the probability of each number of successes three and fewer, for 0, 1, 2, and 3. Let's sum these probabilities.

```
dbinom(x = 0:3,  
       size = 10,  
       prob = 0.5) |>  
sum()
```

```
[1] 0.171875
```

As we can see, we get the same value as for `pbinom`! We can use this method for finding very specific answers, like what is the probability of getting between 3 and 6 successes in ten trials?

```
dbinom(x = 3:6,  
       size = 10,  
       prob = 0.5) |>  
sum() |>  
round(4)
```

```
[1] 0.7734
```

The probability of getting one of these outcomes is 77.34%.

8.5.2 Poisson cumulative probability

Cumulative probability for Poisson distributions

Likewise, we can use `ppois` to get the p value and `dpois` to get the distribution function of specific outcomes. So, let's say we have a scenario with a $\lambda = 0.5$ and we are looking at the probability of 2 successes or greater. In this case, we have an infinite series, which we can't calculate. However, we can calculate the probability of what it *isn't* and then subtract from 1. In this case, we are looking for the probability of *not* having 0 or 1 successes.

```
dpois(x = 0:1,  
      lambda = 0.5)
```

```
[1] 0.6065307 0.3032653
```

Now, let's sum this and subtract it from 1.

```
1 - dpois(x = 0:1, lambda = 0.5) |>  
  sum()
```

```
[1] 0.09020401
```

The probability is only about 9%.

8.6 Homework

Remember - your homework must be completed in an **rmarkdown** document submitted as an **html** file.

For each question, you must determine whether a binomial or a Poisson distribution is most appropriate. (Provide justification on your document for why you chose the distribution you did). After picking the correct distribution, answer the statistical question posed.

Some of the inspiration for these problems is drawn from [Statology](#).

8.6.1 Problem 1

There is a claw machine at the local fair that is notoriously difficult to use. In fact, it is estimated that only 1 in 50 attempts are successful at getting a toy out of the machine. What is the probability that you would try it 50 times and win thrice?

8.6.2 Problem 2

Whooping Cranes *Grus americanus* are some of the most endangered birds in North America, with only 557 birds in the population that migrates through Nebraska. Conversely, approximately 736,000 Sandhill Cranes *Antigone canadensis* migrate through Nebraska. Assuming that census counts of flocks of 5000 cranes have the expected average ratio of cranes, what are the odds of finding a crane flock with two or more Whooping Cranes?

8.6.3 Problem 3

Your friend has a twenty-sided dice that you suspect to be unfair. They have rolled a 20 four times in the past thirty rolls. Is the dice unfair?

8.6.4 Problem 4

You roll three die at a single time. What are the odds of getting exactly two fours?

8.6.5 Problem 5

You notice that in a class of 200 people, the odds of any given student being absent is approximately 3%. What are the odds that 5 people miss a given class?

8.6.6 Problem 6

There are approximately 3 severe storm warnings per 16 week semester at UNK. What are the odds of having a semester with no severe storm warnings?

8.6.7 Problem 7

You are giving mice an experimental treatment that is designed to change their fur color. In a trial of 500 mice, 15 had their fur change color. Using this precedent, are the odds of having 6 mice have their fur change color in another trial of 500 individuals?

8.6.8 Problem 8

What are the odds of rolling no fours on two twenty-sided die?

8.6.9 Problem 9

Dr. Cooper once caught his mom shaking sodas and putting them back at the shelves at a gas station while they were traveling. Assuming his mom was able to shake 10% of the sodas before getting caught, what are the odds of getting three shaken sodas if you buy 20 sodas?

9 2 (Chi-squared) tests

9.1 χ^2 -squared distribution

χ^2 -squared (pronounced “*kai*”, and spelled “*chi*”) is a distribution used to understand if count data between different categories matches our expectation. For example, if we are looking at students in the class and comparing major vs. number of books read, we would expect *no association*, however we may find an association for a major such as English which required reading more literature. The χ^2 introduces a new term *degrees of freedom* (*df*) which reflects the number of individuals in the study. For many tests, *df* are needed to reflect how a distribution changes with respect the number of individuals (and amount of variation possible) within a dataset. The equation for the χ^2 is as follows, with the χ^2 being a special case of the *gamma* (γ or Γ) distribution that is affected by the *df*, which is defined as the number of rows minus one multiplied by the number of columns minus one $df = (rows - 1)(cols - 1)$:

$$f_n(x) = \frac{1}{2^{\frac{n}{2}} \Gamma(\frac{n}{2})} x^{\frac{n}{2}-1} e^{-\frac{x}{2}}$$


The χ^2 -squared distribution is also represented by the following functions, which perform the same things as the previous outlined equivalents for Poisson and binomial:

- `dchisq`
- `pchisq`
- `qchisq`
- `rchisq`

We can view these probabilities as well:

```
x <- 0:10  
  
y <- pchisq(x, df = 9)  
  
# any value greater than 0.5 is subtracted from 1  
y[y > 0.5] <- 1 - y[y > 0.5]
```

```
plot(x,y,type="l")
```



9.1.1 Calculating the test statistic

We evaluate χ^2 tests by calculating a χ^2 value based on our data and comparing it to an expected χ^2 distribution. This *test statistic* can be evaluated by looking at a χ^2 table or by using *R*. **Note** that you need to know the degrees of freedom in order to properly evaluate a χ^2 test. We calculate our test statistic as follows:

$$\chi^2 = \sum \frac{(o - e)^2}{e}$$

where e = the number of expected individuals and o = the number of observed individuals in each category. Since we are squaring these values, we will only have positive values, and thus this will **always be a one-tailed test**.

There are multiple types of χ^2 test, including the following we will cover here:

- χ^2 Goodness-of-fit test
- χ^2 test of independence

9.1.2 χ^2 goodness-of-fit test

A χ^2 goodness-of-fit test looks at a vector of data, or counts in different categories, and asks if the observed frequencies vary from the expected frequencies.

9.1.2.1 χ^2 estimate by hand

Let's say, for example, we have the following dataset:

| Hour | No. Drinks Sold |
|-------|-----------------|
| 6-7 | 3 |
| 7-8 | 8 |
| 8-9 | 15 |
| 9-10 | 7 |
| 10-12 | 5 |
| 12-13 | 20 |
| 13-14 | 18 |
| 14-15 | 8 |
| 15-16 | 10 |
| 16-17 | 12 |

Now, we can ask if the *probability* of selling drinks is the same across all time periods.

```
drinks <- c(3, 8, 15, 7, 5, 20, 18, 8, 10, 12)
```

We can get the expected counts by assuming an equal probability for each time period; thus,
 $Exp(x) = \frac{N}{categories}$.

```
# sum all values for total number
N <- sum(drinks)

# get length of vector for categories
cats <- length(drinks)

# repeat calculation same number of times as length
exp_drinks <- rep(N/cats, cats)

exp_drinks
```

```
[1] 10.6 10.6 10.6 10.6 10.6 10.6 10.6 10.6 10.6 10.6
```

Now, we can do out χ^2 calculation.

```
chi_vals <- ((drinks - exp_drinks)^2)/exp_drinks
```

```
chi_vals
```

```
[1] 5.44905660 0.63773585 1.82641509 1.22264151 2.95849057 8.33584906  
[7] 5.16603774 0.63773585 0.03396226 0.18490566
```

```
sum(chi_vals)
```

```
[1] 26.45283
```

And now, to get the probability.

```
chi_vals |>  
  # get chi statistic  
  sum() |>  
  # get p value  
  pchisq(df = length(chi_vals) - 1,  
         # looking RIGHT  
         lower.tail = F)
```

```
[1] 0.001721825
```

Here, we get $p = 0.002$, indicating that there is not an equal probability for selling drinks at different times of day.

9.1.2.2 χ^2 estimation by code

We can use the test `chisq.test` to perform this analysis as well.

```
chisq.test(drinks)
```

Chi-squared test for given probabilities

```
data: drinks
```

```
X-squared = 26.453, df = 9, p-value = 0.001722
```

As we can see, these values are exactly the same as we just calculated by hand! **Note** that we can define the probability `p` if we want, otherwise it defaults to `p = rep(1/length(x), length(x))`.

9.1.3 χ^2 test of independence

Usually when we use a χ^2 , we are looking at count data. Let's consider the following hypothetical scenario, comparing experience with *R* between non-biology majors (who, in this theoretical scenario, do not regularly use *R*) and Biology majors who are required to take *R* for this class:

Table 9.2: The above table of counts is also known as a **contingency table**. Intuitively, we can see a difference, but we want to perform a statistical test to see just how likely these counts would be if both groups were equally likely. We can calculate this both “by hand” and using built in *R* functions.

| Major | <i>R</i> experience | No <i>R</i> experience |
|--------------------|---------------------|------------------------|
| <i>Non-biology</i> | 3 | 10 |
| <i>Biology</i> | 9 | 2 |

9.1.3.1 χ^2 estimations by hand

First, we can enter the data into *R*.

```
data <- matrix(data = c(3,10,9,2), nrow = 2, ncol = 2, byrow = T)

colnames(data) <- c("R", "No R")
rownames(data) <- c("Non-biology", "Biology")

data
```

```
      R No R
Non-biology 3  10
Biology     9   2
```

Next, we need the *observed* - *expected* values. We determine expected values either through probability ($0.5 \cdot n$ for equal probability for two categories) or via calculating the the expected values (see later section on *expected counts*). In this case, since we are looking at equally likely in each cell, we have an expected matrix as follows:

```
# total datapoints
N <- sum(data)

expected <- matrix(data = c(0.25*N,0.25*N,0.25*N,0.25*N), nrow = 2, ncol = 2, byrow = T)
```

```
colnames(expected) <- c("R", "No R")
rownames(expected) <- c("Non-biology", "Biology")

expected
```

```
      R No R
Non-biology 6   6
Biology     6   6
```

Now we need to find our *observed* - *expected*.

```
o_e <- data - expected

o_e
```

```
      R No R
Non-biology -3   4
Biology      3  -4
```

Note that in *R* we can add and subtract matrices, so there's no reason to reformat these data!

Now, we can square these data.

```
o_e2 <- o_e^2

o_e2
```

```
      R No R
Non-biology 9  16
Biology      9  16
```

Next, we take these and divide them by the expected values and then sum those values.

```
chi_matrix <- o_e2/expected

chi_matrix
```

| | R | No R |
|-------------|-----|----------|
| Non-biology | 1.5 | 2.666667 |
| Biology | 1.5 | 2.666667 |

```
sum(chi_matrix)
```

```
[1] 8.333333
```

Here, we get a χ^2 value of 8.333333. We can use our handy family functions to determine the probability of this event:

```
chi_matrix |>
  # get chi statistic
  sum() |>
  pchisq(df = 1,
         # looking RIGHT
         lower.tail = F)
```

```
[1] 0.003892417
```

Here, we get a p value of:

```
chi_matrix |>
  sum() |>
  pchisq(df = 1, lower.tail = F) |>
  round(3)
```

```
[1] 0.004
```

Alternatively, we can calculate this using *expected counts*. For many situations, we don't know what the baseline probability *should* be, so we calculate the expected counts based on what we do know. Expected counts are calculated as follows:

$$Exp(x) = \frac{\Sigma(row_x) \cdot \Sigma(col_x)}{N}$$

where N is the sum of all individuals in the table. For the above example, this would look like this:


```

data_colsums <- colSums(data)
data_rowsums <- rowSums(data)
N <- sum(data)

expected <- matrix(data = c(data_colsums[1]*data_rowsums[1],
                             data_colsums[2]*data_rowsums[1],
                             data_colsums[1]*data_rowsums[2],
                             data_colsums[2]*data_rowsums[2]),
                    nrow = 2, ncol = 2, byrow = T)

# divide by total number
expected <- expected/N

colnames(expected) <- colnames(data)
rownames(expected) <- rownames(data)

expected

```

| | R | No R |
|-------------|-----|------|
| Non-biology | 6.5 | 6.5 |
| Biology | 5.5 | 5.5 |

Here, we can see our expected number are not quite 50/50! this will give us a different result than our previous iteration.

```

e_o2 <- ((data - expected)^2)/expected

sum(e_o2)

```

```
[1] 8.223776
```

Now we have a χ^2 of:

```

e_o2 |>
  sum() |>
  round(2)

```

```
[1] 8.22
```

As we will see below, is the exact same value as we get for an uncorrected `chisq.test` from *R*'s default output.

9.1.3.2 χ^2 estimations in *R*

We can calculate this in *R* by entering in the entire table and using `chisq.test`.

```
data
```

| | R | No R |
|-------------|---|------|
| Non-biology | 3 | 10 |
| Biology | 9 | 2 |

```
chi_data <- chisq.test(data, correct = F)
```

```
chi_data
```

Pearson's Chi-squared test

```
data: data
X-squared = 8.2238, df = 1, p-value = 0.004135
```

Here, we get the following for your p value:

```
chi_data$p.value |>
  round(3)
```

```
[1] 0.004
```

This is significant with $\alpha = 0.05$.

Note that these values are slightly different. they will be even more different if `correct` is set to `TRUE`. By default, *R*, uses a [Yate's correction for continuity](#). This accounts for error introduced by comparing the discrete values to a continuous distribution.

```
chisq.test(data)
```

Pearson's Chi-squared test with Yates' continuity correction

```
data: data
X-squared = 6.042, df = 1, p-value = 0.01397
```

Applying this correction *lowers* the degrees of freedom, and increases the p value, thus making it harder to get $p < \alpha$.

Note that the Yate's correction is only applied for 2 x 2 contingency tables.

Given the slight differences in calculation between by hand and what the functions of R are performing, *it's important to always show your work.*

9.2 Fisher's exact test

χ^2 tests don't work in scenarios where we have very small count sizes, such as a count size of 1. For these situations with small sample sizes and very small count sizes, we use Fisher's exact test. This test gives us the p value directly - no need to use a table of any kind! Let's say we have a 2x2 contingency table, as follows:

| | |
|-----|-----|
| a | b |
| c | d |

Where row totals are $a+b$ and $c+d$ and column totals are $a+c$ and $b+d$, and $n = a+b+c+d$. We can calculate p as follows:

$$p = \frac{(a+b)!(c+d)!(a+c)!(b+d)!}{a!b!c!d!n!}$$

In R , we can use the command `fisher.test` to perform these calculations. For example, we have the following contingency table, looking at the number of undergrads and graduate students in introductory and graduate level statistics courses:

| | Undergrad | Graduate |
|-------------|-----------|----------|
| Intro Stats | 8 | 1 |
| Grad Stats | 3 | 5 |

```
stats_students = matrix(data = c(8,1,3,5),
                        byrow = T, ncol = 2, nrow = 2)

stats_students
```

```
      [,1] [,2]
[1,]    8    1
[2,]    3    5
```

```
fisher.test(stats_students)
```

Fisher's Exact Test for Count Data

```
data: stats_students
p-value = 0.04977
alternative hypothesis: true odds ratio is not equal to 1
95 percent confidence interval:
 0.7934527 703.0167380
sample estimates:
odds ratio
 11.10917
```

In this situation, $p = 0.05$ when rounded, so we would fail to reject but note that this is borderline.

9.2.1 Chapter 7

Complete problems 7.1, 7.2, 7.3, 7.5, and 7.7 as written in your text books. For these problems, please also state your *null* and *alternative* hypotheses, as well as a conclusion as to whether you *support* or *reject* the null.

Next, do problems 7.9 and 7.11. For these two problems, pick which test is best and then perform said test. State your hypothesis and make a conclusion with respect to your p value.

Be sure to submit your homework as a knitted `html` document.

10 Means testing

10.1 Introduction

Previously, we talked about normal distributions as a method for comparing samples to overall populations or comparing individuals to overall populations. However, sample sizes can introduce some error, and oftentimes we may not have access to an entire population. In these situations, we need a better test that can account for this changing error and the effect of different sample sizes. This is especially important when comparing two samples to each other. We may find a small sample from one population and a small sample for another, and we want to determine if these came from the same overall population as effectively as possible.

10.1.1 Parametric and Non-parametric tests

We have divided tests into *parametric* and *non-parametric* tests below. *Parametric* tests are those that follow a normal distribution; *non-parametric* tests violate this expectation.

Remember, parametric tests are more powerful and preferred in all circumstances. If your data are not parametric, you will first have to see if the data can be transformed. If the data cannot be transformed, then you can proceed with a non-parametric test.

10.1.2 A little history

Why is it called “Student’s t test”?

The distribution that we commonly refer to as a t -distribution is also sometimes known as a “Student’s t -distribution” as it was first published by a man with the pseudonym of “Student”. Student was in fact [William Sealy Gossett](#), an employee of the Guinness corporation who was barred from publishing things by his employer to ensure that trade secrets were not made known to their competitors. Knowing that his work regarding statistics was important, Gossett opted to publish his research anyway under his pseudonym.

10.2 Dataset

For all of the examples on this page, we will be using a dataset on the morphology of canine teeth for identification of predators killing livestock (Courtenay 2019).

```
canines <- read_csv("https://figshare.com/ndownloader/files/15070175")
```

We want to set up some of these columns as “factors” to make it easier to process and parse in *R*. We will look at the column OA for these examples. Unfortunately, it is unclear what exactly OA stands for since this paper is not published at the present time.

```
canines$Sample <- as.factor(canines$Sample)

# we will be examining the column "OA"

canines$OA <- as.numeric(canines$OA)

summary(canines)
```

| Sample | WIS | WIM | WIB |
|---------|----------------|----------------|-----------------|
| Dog :34 | Min. :0.1323 | Min. :0.1020 | Min. :0.03402 |
| Fox :41 | 1st Qu.:0.5274 | 1st Qu.:0.3184 | 1st Qu.:0.11271 |
| Wolf:28 | Median :1.1759 | Median :0.6678 | Median :0.25861 |
| | Mean :1.6292 | Mean :1.0233 | Mean :0.44871 |
| | 3rd Qu.:2.4822 | 3rd Qu.:1.5194 | 3rd Qu.:0.74075 |
| | Max. :4.8575 | Max. :3.2423 | Max. :1.51721 |

| D | RDC | LDC | OA |
|------------------|-----------------|-----------------|---------------|
| Min. :0.005485 | Min. :0.05739 | Min. :0.02905 | Min. :100.7 |
| 1st Qu.:0.034092 | 1st Qu.:0.28896 | 1st Qu.:0.22290 | 1st Qu.:139.2 |
| Median :0.182371 | Median :0.61777 | Median :0.55985 | Median :149.9 |
| Mean :0.250188 | Mean :0.88071 | Mean :0.84615 | Mean :148.4 |
| 3rd Qu.:0.361658 | 3rd Qu.:1.26417 | 3rd Qu.:1.26754 | 3rd Qu.:158.0 |
| Max. :1.697461 | Max. :3.02282 | Max. :3.20533 | Max. :171.5 |

10.3 Parametric tests

10.3.1 *t*-distribution

For these scenarios where we are testing a single sample mean from one or more samples we use a *t*-distributions. A *t*-distribution is a specially altered normal distribution that has been

adjusted to account for the number of individuals being sampled. Specifically, a t -distributions with infinite degrees of freedom is the same as a normal distribution, and our degrees of freedom help create a more platykurtic distribution to account for error and uncertainty. The distribution can be calculated as follows:

$$t = \frac{\Gamma(\frac{v+1}{2})}{\sqrt{\pi\nu}\Gamma(\frac{\nu}{2})} \left(1 + \frac{t^2}{\nu}\right)^{-\frac{(v+1)}{2}}$$

These t -distributions can be visualized as follows:




Figure 10.1: IkamusumeFan - Wikipedia

For all t -tests, we calculate the degrees of freedom based on the number of samples. If comparing values to a single sample, we use $df = n - 1$. If we are comparing two sample means, then we have $df = n_1 + n_2 - 2$.

Importantly, we are testing to see if the means of the two distributions are equal in a t -test. Thus, our hypotheses are as follows:

$H_0 : \mu_1 = \mu_2$ **or** $H_0 : \mu_1 - \mu_2 = 0$

$H_A : \mu_1 \neq \mu_2$ **or** $H_A : \mu_1 - \mu_2 \neq 0$

When asked about hypotheses, remember the above as the statistical hypotheses that are being directly tested.

In *R*, we have the following functions to help with *t* distributions:

- **dt**: density function of a *t*-distribution
- **pt**: finding our *p* value from a specific *t* in a *t*-distribution
- **qt**: finding a particular *t* from a specific *p* in a *t*-distribution
- **rt**: random values from a *t*-distribution

All of the above arguments required the degrees of freedom to be declared. Unlike the normal distribution functions, these can not be adjusted for your data; tests must be performed using `t.test`.

10.3.2 *t*-tests

We have three major types of *t*-tests:

- **One-sample *t*-tests**: a single sample is being compared to a value, or *vice versa*.
- **Two-sample *t*-tests**: two samples are being compared to one another to see if they come from the same population.
- **Paired *t*-tests**: before-and-after measurements of the same individuals are being compared. This is necessary to account for a repeat in the individuals being measured, and different potential baselines at initiation. In this case, we are looking to see if the difference between before and after is equal to zero.

We also have what we call a “true” *t*-test and “Welch’s” *t*-test. The formula for a “true” *t* is as follows:

$$t = \frac{\bar{x}_1 - \bar{x}_2}{s_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}}$$

Where s_p is based on the “pooled variance” between the samples. This can be calculated as follows:

$$s_p = \sqrt{\frac{(n_1 - 1)(s_1^2) + (n_2 - 1)(s_2^2)}{n_1 + n_2 - 2}}$$

Whereas the equation for a “Welch’s” t is:

$$t = \frac{\bar{x}_1 - \bar{x}_2}{\sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}}$$

Welch’s t also varies with respect to the degrees of freedom, calculated by:

$$df = \frac{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}{\frac{(\frac{s_1^2}{n_1})^2}{n_1-1} + \frac{(\frac{s_2^2}{n_2})^2}{n_2-1}}$$

OK, so why the difference?

A t -test works well under a certain set of assumptions, include *equal variance* between samples and roughly equal *sample sizes*. A Welch’s t -test is better for scenarios with *unequal variance* and *small sample sizes*. If sample sizes and variances are equal, the two t -tests should perform the same.

Because of this, some argue that “Welch’s” should be the default t -test, and **in *R*, Welch’s is the default t -test**. If you want to specify a “regular” t -value, you will have to set the option `var.equal = TRUE`. (The default is `var.equal = FALSE`).

In this class, we will default to a Welch’s test in *all* instances.

If you choose to do a Student’s t -test, you must do the following:

- Download the `car` library
- Use the `leveneTest` function to see if variances are equal between populations

We do not cover this in depth here, but be aware of this difference. For more information, see Ruxton (2006).

10.3.3 One-sample t -tests

Let’s look at the values of all of the dog samples in our `canines` dataset.

```
dogs <- canines |>
  filter(Sample == "Dog") |>
  select(Sample, OA)

xbar <- mean(dogs$OA)
sd_dog <- sd(dogs$OA)
n <- nrow(dogs)
```

Now we have stored all of our information on our dog dataset. Let's say that the overall populations of dogs a mean OA score of 143 with a $\sigma = 1.5$. Is our sample different than the overall population?

```
t.test(x = dogs$OA,  
       alternative = "two.sided",  
       mu = 143)
```

One Sample t-test

```
data: dogs$OA  
t = -0.74339, df = 33, p-value = 0.4625  
alternative hypothesis: true mean is not equal to 143  
95 percent confidence interval:  
 138.4667 145.1070  
sample estimates:  
mean of x  
 141.7869
```

As we can see above, we fail to reject the null hypothesis that our sample is different than the overall mean for dogs.

10.3.4 Two-sample *t*-tests

Now let's say we want to compare foxes and dogs to each other. Since we have all of our data in the same data frame, we will have to *subset* our data to ensure we are doing this properly.

```
# already got dogs  
dog_oa <- dogs$OA  
  
foxes <- canines |>  
  filter(Sample == "Fox") |>  
  select(Sample, OA)  
  
fox_oa <- foxes$OA
```

Now, we are ready for the test!

```
t.test(dog_oa, fox_oa)
```

Welch Two Sample t-test

```
data: dog_oa and fox_oa
t = -6.3399, df = 72.766, p-value = 1.717e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -19.62289 -10.23599
sample estimates:
mean of x mean of y
 141.7869  156.7163
```

As we can see, the dogs and the foxes significantly differ in their OA measurement, so we reject the null hypothesis that $\mu_{dog} = \mu_{fox}$.

10.3.5 Paired *t*-tests

I will do a highly simplified version of a paired *t*-test here just for demonstrations sake. Remember that you want to use paired tests when we are looking at the *same individuals* at different points in time.

```
# create two random distributions
# DEMONSTRATION ONLY

# make repeatable
set.seed(867)

t1 <- rnorm(20,0,1)
t2 <- rnorm(20,2,1)
```

Now we can compare these using `paired = TRUE`.

```
t.test(t1, t2, paired = TRUE)
```

Paired t-test

```
data: t1 and t2
```

```
t = -7.5663, df = 19, p-value = 3.796e-07
alternative hypothesis: true mean difference is not equal to 0
95 percent confidence interval:
 -3.107787 -1.760973
sample estimates:
mean difference
 -2.43438
```

As we can see, we reject the null hypothesis that these distributions are equal in this case. Let's see how this changes though if we set `paired = FALSE`.

```
t.test(t1, t2)
```

Welch Two Sample t-test

```
data:  t1 and t2
t = -8.1501, df = 37.48, p-value = 8.03e-10
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -3.039333 -1.829428
sample estimates:
 mean of x  mean of y
-0.07258938  2.36179080
```

This value differs because, in a paired test, we are looking to see if the difference between the distributions is 0, while in the independent (standard) test we are comparing the overall distributions of the samples.

10.4 Non-parametric tests

The following tests should be used when no data transformations have been successful with your dataset.

10.4.1 Wilcoxon tests

When data (and the differences among data) are non-normal, they violate the assumptions of a *t*-test. In these cases, we have to do a Wilcoxon test (also called a Wilcoxon signed rank test). In *R*, the command `wilcox.test` also includes the Mann-Whitney *U* test for unpaired data and the standard Wilcoxon test *W* for paired data.

10.4.2 Mann-Whitney U

For this test, we would perform the following procedures to figure out our statistics:

1. Rank the pooled dataset from smallest to largest, and number all numbers by their ranks
2. Sum the ranks for the first column and the second column
3. Compute U_1 and U_2 , comparing the smallest value to a Mann-Whitney U table.

The equations for these statistics are as follows, where R represents the sum of the ranks for that sample:

$$U_1 = n_1 n_2 + \frac{n_1(n_1 + 1)}{2} - R_1$$

$$U_2 = n_1 n_2 + \frac{n_2(n_2 + 1)}{2} - R_2$$

In R , this looks like so:

```
wilcox.test(t1, t2, paired = FALSE)
```

```
Wilcoxon rank sum exact test
```

```
data: t1 and t2
```

```
W = 11, p-value = 2.829e-09
```

```
alternative hypothesis: true location shift is not equal to 0
```

10.4.3 Wilcoxon signed rank test

For paired samples, we want to do the Wilcoxon signed rank test. This is performed by:

1. Finding the difference between sampling events for each sampling unit.
2. Order the differences based on their *absolute value*
3. Find the sum of the *positive ranks* and the *negative ranks*
4. The *smaller* of the values is your W statistic.

In R , this test looks as follows:

```
wilcox.test(t1, t2, paired = TRUE)
```

Wilcoxon signed rank exact test

data: t1 and t2

V = 0, p-value = 1.907e-06

alternative hypothesis: true location shift is not equal to 0

10.5 Confidence intervals

In t tests, we are looking at the difference between the means. Oftentimes, we are looking at a confidence interval for the difference between these means. This can be determined by:

$$(\bar{x}_1 - \bar{x}_2) \pm t_{crit} \sqrt{\frac{s_p^2}{n_1} + \frac{s_p^2}{n_2}}$$

This is very similar to the CI we calculated with the Z statistic. **Remember** that we can use the following function to find our desired t , which requires degrees of freedom to work:

```
qt(0.975, df = 10)
```

```
[1] 2.228139
```

10.6 Homework: Chapter 9

For Chapter 9, complete problems 9.1, 9.3, 9.4, 9.5, 9.6, 9.7, 9.8, 9.9, and 9.10. For problems 9.3 - 9.8, be sure to state the null and alternative hypotheses and whether the test is one- or two-tailed.

10.7 Homework: Chapter 10

Two-sample means are practiced in Chapter 10. Please see *Canvas* for more information.

11 ANOVA: Part 1

11.1 Introduction

When we are comparing multiple (2+) populations, we perform what is called an *analysis of variance* - or an ANOVA. We opt for this different method because we are trying to minimize error. As you'll recall, we use α to minimize our chances of making an error and coming to an incorrect conclusion regarding our data. In our previous tests (*t*-tests) we are comparing the means between two different populations, such that $H_0 : \mu_1 = \mu_2$. When comparing multiple populations, comparing the means in this direct fashion can increase the probability of introducing error into a system. Consider the following:

```
library(tidyverse)

# needed for summarizing data
library(plyr)

# needed for better Tukey tests
library(agricolae)

# This creates a reproducible example
# rnorm creates random datasets

set.seed(8675309)

for(i in 1:100){
  x <- rnorm(10)
  if(i == 1){
    data <- x |> as.data.frame()
    colnames(data) <- "Response"
    data$Explanatory <- paste0("x",i)
  }else{
    newdat <- x |> as.data.frame()
    colnames(newdat) <- "Response"
    newdat$Explanatory <- paste0("x",i)
    data <- rbind(data,newdat)
  }
}
```




```

}
}

# summarize by group
summary_data <- ddply(data, "Explanatory", summarise,
  N = length(Response),
  mean = mean(Response),
  sd = sd(Response),
  se = sd / sqrt(N))

ggplot(summary_data, aes(x = Explanatory, y = mean, group = Explanatory)) +
  geom_point() +
  geom_errorbar(data = summary_data, aes(ymin = mean - 2*se, ymax = mean+2*se,
    color = Explanatory), width = 0.1) +
  geom_hline(yintercept = 0, col = "black", linewidth = 0.5) +
  ylim(c(-1.5,1.5)) +
  theme_classic() +
  theme(legend.position = "none",
    axis.text.x = element_text(angle = 90, vjust = 0.5, size = 5))

```



As we can see above, with just ten random samples and 100 sampling events, we get some datasets that do not have the mean included within the interquartile range, and thus have means that would be statistically different than what we draw. As we increase the number of draws, we get closer to the mean:

```


set.seed(8675309)

for(i in 1:100){
  x <- rnorm(100)
  if(i == 1){
    data <- x |> as.data.frame()
    colnames(data) <- "Response"
    data$Explanatory <- paste0("x",i)
  }else{
    newdat <- x |> as.data.frame()
    colnames(newdat) <- "Response"
    newdat$Explanatory <- paste0("x",i)
    data <- rbind(data,newdat)
  }
}

# summarize by group
summary_data <- dplyr::ddply(data, "Explanatory", summarise,
  N = length(Response),
  mean = mean(Response),
  sd = sd(Response),
  se = sd / sqrt(N))

ggplot(summary_data, aes(x = Explanatory, y = mean, group = Explanatory)) +
  geom_point() +
  geom_errorbar(data = summary_data, aes(ymin = mean - 2*se, ymax = mean+2*se,
    color = Explanatory), width = 0.1) +
  geom_hline(yintercept = 0, col = "black", linewidth = 0.5) +
  ylim(c(-1.5,1.5)) +
  theme_classic() +
  theme(legend.position = "none",
    axis.text.x = element_text(angle = 90, vjust = 0.5, size = 5))

```



As we can see, even with 100 sample, we still have some chances of having groups that are different! When we do pairwise comparisons, we are compounding the error and the possibility of coming to an incorrect conclusion. Thus, when comparing multiple groups, we use the variances to see if groups come from the same distribution rather than the mean.

11.2 ANOVA: By hand

We are predominately going to be using the default function `aov` to perform ANOVAs in this course. However, you can expand the following walkthrough if you would like to see step-by-step instructions on calculating a one-way ANOVA by hand.

[Click here to see the manual method.](#)

For this walkthrough, we will use the following psuedorandom dataset:

```
set.seed(8675309)

for(i in 1:4){
  x <- rnorm(10)
  if(i == 1){
    x <- rnorm(10, mean = 2)
    data <- x |> as.data.frame()
    colnames(data) <- "Response"
  }
}
```

```

    data$Explanatory <- paste0("x",i)
  }else{
    newdat <- x |> as.data.frame()
    colnames(newdat) <- "Response"
    newdat$Explanatory <- paste0("x",i)
    data <- rbind(data,newdat)
  }
}


# split into "typical" table
expanded_data <- NULL
expanded_data$x1 <- data$Response[which(data$Explanatory=="x1")]
expanded_data$x2 <- data$Response[which(data$Explanatory=="x2")]
expanded_data$x3 <- data$Response[which(data$Explanatory=="x3")]
expanded_data$x4 <- data$Response[which(data$Explanatory=="x4")]

expanded_data <- expanded_data |>
  as.data.frame()

# summarize by group
summary_data <- dplyr::ddply(data, "Explanatory", summarise,
  N = length(Response),
  mean = mean(Response),
  sd = sd(Response),
  se = sd / sqrt(N))

ggplot(summary_data, aes(x = Explanatory, y = mean, group = Explanatory)) +
  geom_point() +
  geom_errorbar(data = summary_data, aes(ymin = mean - 2*se, ymax = mean+2*se,
    color = Explanatory), width = 0.1) +
  geom_hline(yintercept = 0, col = "black", linewidth = 0.5) +
  ylim(c(-3,3)) +
  theme_classic() +
  theme(legend.position = "none",
    axis.text.x = element_text(angle = 90, vjust = 0.5, size = 5))

```



```
expanded_data |> round(2)
```

| | x1 | x2 | x3 | x4 |
|----|------|-------|-------|-------|
| 1 | 0.45 | 1.99 | 0.38 | -0.64 |
| 2 | 3.02 | 0.04 | 1.15 | 0.00 |
| 3 | 2.15 | -0.40 | 1.57 | 0.05 |
| 4 | 1.34 | -0.47 | 0.59 | 0.68 |
| 5 | 1.01 | -0.41 | -0.62 | -0.25 |
| 6 | 3.97 | 0.68 | -0.23 | -0.15 |
| 7 | 1.56 | 0.69 | 0.06 | -0.87 |
| 8 | 1.10 | 0.53 | -0.31 | -1.98 |
| 9 | 1.85 | -0.19 | -0.25 | 0.24 |
| 10 | 1.17 | 0.38 | -0.15 | 0.04 |

Above, we can see the made-up dataset where it appears as though one population differs from the other populations in our measurements. Let's calculate an ANOVA and find out if this is the case!

NOTE throughout this process that I am trying to name variables in a straightforward fashion so as not to lose my way.

11.2.1 Calculate group means and Grand Mean

Let us assume we have a dataset, x , that is k columns and n rows, with N data points in the entire data frame. We first want to take the column means for each group k , such that we have \bar{x}_k . We also need to find the mean of the entire dataset, \bar{x}_n . We can calculate this by taking $\frac{\sum x}{n}$, which we have to calculate by column as follows.

```
# calculate the mean of each group
# each group is in a single column
group_means <- colMeans(expanded_data)

# rounding to two decimal places
group_means |> round(2)
```

```
      x1      x2      x3      x4
1.76  0.28  0.22 -0.29
```

Next, we need to calculate the number of total entries in the dataset. We have written a function to accomplish this incase some rows have different numbers of entries from others.

```
n <- 0

for(i in 1:ncol(expanded_data)){
  # account for unequal row length, if exists
  sample <- expanded_data[,i] |>
    as.numeric() |>
    na.omit()
  n <- n + length(sample)
}

n
```

```
[1] 40
```

Next, we can calculate the `grand_mean` of all of the data.

```
# sum up all the data
dataset_sum <- colSums(expanded_data) |>
  sum()

# divide by the number of data points
```

```
grand_mean <- dataset_sum/n

# display mean
grand_mean |> round(2)
```

```
[1] 0.49
```

11.2.2 Total sum of squares

To calculate the *total sum of squares* (TSS), we need to take the deviations (differences) of each point from the grand mean \bar{x}_n , square them, and then take the sum of them.

```
# calculate deviates
# can calculate across all table at once
grand_deviates_squared <- (expanded_data - grand_mean)^2

# round output for here
grand_deviates_squared |> round(2)
```

| | x1 | x2 | x3 | x4 |
|----|-------|------|------|------|
| 1 | 0.00 | 2.22 | 0.01 | 1.28 |
| 2 | 6.39 | 0.20 | 0.43 | 0.25 |
| 3 | 2.74 | 0.81 | 1.17 | 0.20 |
| 4 | 0.72 | 0.94 | 0.01 | 0.04 |
| 5 | 0.26 | 0.83 | 1.23 | 0.56 |
| 6 | 12.10 | 0.04 | 0.52 | 0.42 |
| 7 | 1.13 | 0.04 | 0.19 | 1.87 |
| 8 | 0.37 | 0.00 | 0.65 | 6.11 |
| 9 | 1.84 | 0.46 | 0.55 | 0.07 |
| 10 | 0.46 | 0.01 | 0.42 | 0.21 |

```
# calculate the sum of all the deviates
ss_total <- rowSums(grand_deviates_squared) |>
  sum()

ss_total |> round(2)
```

```
[1] 47.73
```

11.2.3 Within-group sum of squares

For each data point, we need to calculate its deviation from its own group mean, squaring these deviations and then summing them together. We can't calculate this quite as elegantly as the aforementioned data, but we can write a function that will operate across the table and create a new dataset on our behalf.

```
# replicate dataset
# replace columns with deviate data
group_deviates <- expanded_data

# loop through each column
for(i in 1:ncol(group_deviates)){
  # get the data in each group
  dat <- group_deviates[,i]
  # calculate the group mean
  mu <- mean(dat)
  # calculate the group deviates
  dev.dat <- (dat - mu)^2
  # save into table
  group_deviates[,i] <- dev.dat
}

group_deviates |> round(2)
```

| | x1 | x2 | x3 | x4 |
|----|------|------|------|------|
| 1 | 1.72 | 2.90 | 0.02 | 0.12 |
| 2 | 1.59 | 0.06 | 0.87 | 0.08 |
| 3 | 0.15 | 0.47 | 1.84 | 0.11 |
| 4 | 0.18 | 0.57 | 0.14 | 0.95 |
| 5 | 0.57 | 0.49 | 0.70 | 0.00 |
| 6 | 4.89 | 0.16 | 0.20 | 0.02 |
| 7 | 0.04 | 0.16 | 0.02 | 0.34 |
| 8 | 0.44 | 0.06 | 0.28 | 2.85 |
| 9 | 0.01 | 0.22 | 0.22 | 0.28 |
| 10 | 0.35 | 0.01 | 0.14 | 0.11 |

```
# calculate sum of data table
ss_within <- colSums(group_deviates) |>
  sum()

ss_within |> round(2)
```



```
[1] 24.33
```

11.2.4 Among-group sum of squares

The total sum of squares is equal to the among groups sum of squares and the within groups sum of squares added together; thus, we can solve this part with some easy arithmetic.

```
ss_among <- ss_total - ss_within  
  
ss_among |> round(2)
```

```
[1] 23.4
```

11.2.5 Calculate degrees of freedom

Our degrees of freedom for the “between” group is the number of categories minus one ($K - 1$).

```
ss_among_df <- ncol(expanded_data) - 1  
  
ss_among_df
```

```
[1] 3
```

Our degrees of freedom for the within group are the number of total samples minus the number of categories ($N - K$).

```
ss_within_df <- n - ncol(expanded_data)  
  
ss_within_df
```

```
[1] 36
```

Our degrees of freedom for the total sum of squares is the number of samples minus one ($N - 1$).

```
ss_total_df <- n - 1  
  
ss_total_df
```

```
[1] 39
```

11.2.6 Calculate mean squares

For each category (among and within), the mean square is equal to the sum of squares divided by the degrees of freedom.

```
ms_among <- ss_among/ss_among_df  
ms_among |> round(2)
```

```
[1] 7.8
```

```
ms_within <- ss_within/ss_within_df  
ms_within |> round(2)
```

```
[1] 0.68
```

11.2.7 Get F statistic

We divide the sum of squares among data point by the sum of squares within data points to obtain our F statistic.

```
f_stat <- ms_among/ms_within  
f_stat |> round(2)
```

```
[1] 11.54
```

11.2.8 Get p value

We can use the function `pf` to calculate the p value for any given F . *Note* that this function requires two different degrees of freedom to work correctly, **and we are always looking right** since this is a unidirectional distribution.

```
pf(f_stat,  
   df1 = ss_among_df,  
   df2 = ss_within_df,  
   lower.tail = F)
```

```
[1] 1.894073e-05
```

Given how small our p value is, we want to round this to $p < 0.0001$. As we can see, it is very unlikely that these are the same population.

11.3 ANOVA: By R

For this, we need to use the dataframe where we have all data in a single column and all ID's in the other columns. We then show how we want the ANOVA to operate across the data using the \sim symbol.

```
data_aov <- aov(Response ~ Explanatory, data = data)
summary(data_aov)
```

| | Df | Sum Sq | Mean Sq | F value | Pr(>F) |
|-------------|----|--------|---------|---------|--------------|
| Explanatory | 3 | 23.40 | 7.801 | 11.54 | 1.89e-05 *** |
| Residuals | 36 | 24.33 | 0.676 | | |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

These numbers match the *by hand* methodology from the previous section, showing that the math is the same.

11.4 Post-hoc Tukey test

ANOVA tells us *if* a test is different, but it doesn't tell us *which* test is different. To do this, we have to perform a Tukey test.

11.4.1 Tukey test by hand

To do this by hand, we will need a lot of data from our aforementioned ANOVA test.

[Click here](#) to see the by-hand method.

We need to calculate pairwise differences between each set of means.

```

x1_mean <- mean(data$Response[data$Explanatory == "x1"])
x2_mean <- mean(data$Response[data$Explanatory == "x2"])
x3_mean <- mean(data$Response[data$Explanatory == "x3"])
x4_mean <- mean(data$Response[data$Explanatory == "x4"])

group_means <- c(x1_mean, x2_mean, x3_mean, x4_mean)

# calculate all pairwise differences
pairwise_mean_diffs <- dist(group_means)

pairwise_mean_diffs |> round(2)

```

```

      1      2      3
2 1.48
3 1.54 0.06
4 2.05 0.57 0.51

```

Next, we need a critical Q value against which we can compare. For Tukey, our degrees of freedom are the same as the degrees of freedom for SS_{within} : $N - K$.

```

# set p value
tukey_q <- qtukey(p = 0.95,
                  # get length of categories / columns
                  nmeans = ncol(expanded_data),
                  df = ss_within_df)

tukey_q |> round(2)

```

```
[1] 3.81
```

We need to multiply Q by the pooled variance. This is the same as the average of the variances for each group.

```

var_data <- 0

# calculate pooled variance
for(i in 1:ncol(expanded_data)){
  var_data <- var_data + var(expanded_data[,i])
}

```

```
pooled_var_dat <- sqrt(var_data/n)
pooled_var_dat |> round(2)
```

```
[1] 0.26
```

We can calculate the Tukey critical value by multiplying the pooled variance by Q .

```
tukey_critical <- tukey_q*pooled_var_dat
tukey_critical |> round(2)
```

```
[1] 0.99
```

Remember, we are comparing to the actual value, not the rounded value.

Which mean differences are difference compared to our critical value?

```
pairwise_mean_diffs[pairwise_mean_diffs < tukey_critical] <- 0
pairwise_mean_diffs
```

```
      1      2      3
2 1.477866
3 1.542282 0.000000
4 2.051068 0.000000 0.000000
```

As we can see above, three differences cross our threshold - all associated with x_1 .

When we graph things, we want to label this group as different. We will cover this a little later in the tutorial.

11.4.2 Tukey test in *R*

Tukey tests in the *R* program are a bit more straightforward.

```
TukeyHSD(data_aov)
```

Tukey multiple comparisons of means
95% family-wise confidence level

Fit: aov(formula = Response ~ Explanatory, data = data)

```
$Explanatory
      diff      lwr      upr      p adj
x2-x1 -1.47786579 -2.468021 -0.4877106 0.0015554
x3-x1 -1.54228164 -2.532437 -0.5521265 0.0009388
x4-x1 -2.05106768 -3.041223 -1.0609125 0.0000147
x3-x2 -0.06441585 -1.054571  0.9257393 0.9980525
x4-x2 -0.57320189 -1.563357  0.4169533 0.4141599
x4-x3 -0.50878604 -1.498941  0.4813691 0.5173399
```

As we can see above, only three comparisons have a p value of < 0.05 , and thus only those three are significantly different. All involve x_1 .

We can also use `HSD.test` to get more specific results:

```
tukey_data_aov <- HSD.test(data_aov,
  # what to group by?
  "Explanatory",
  # significance level?
  alpha = 0.05,
  # are data unbalanced
  unbalanced = FALSE,
  # show answer?
  console = TRUE)
```

Study: data_aov ~ "Explanatory"

HSD Test for Response

Mean Square Error: 0.6758192

Explanatory, means

| | Response | std | r | se | Min | Max | Q25 | Q50 |
|----|-----------|-----------|----|-----------|------------|----------|------------|-------------|
| x1 | 1.7620153 | 1.0505466 | 10 | 0.2599652 | 0.4504476 | 3.972459 | 1.1175485 | 1.44911720 |
| x2 | 0.2841495 | 0.7532422 | 10 | 0.2599652 | -0.4729986 | 1.985826 | -0.3497379 | 0.21347543 |
| x3 | 0.2197337 | 0.7019368 | 10 | 0.2599652 | -0.6150452 | 1.574903 | -0.2436023 | -0.04493909 |

```

x4 -0.2890524 0.7345336 10 0.2599652 -1.9769014 0.684072 -0.5394534 -0.07741642
      Q75
x1 2.07491533
x2 0.64579865
x3 0.53544151
x4 0.04323417

```

Alpha: 0.05 ; DF Error: 36
Critical Value of Studentized Range: 3.808798

Minimum Significant Difference: 0.9901551

Treatments with the same letter are not significantly different.

```

      Response groups
x1  1.7620153      a
x2  0.2841495      b
x3  0.2197337      b
x4 -0.2890524      b

```

```
print(tukey_data_aov)
```

```

$statistics
      MSerror Df      Mean      CV      MSD
      0.6758192 36 0.4942115 166.3422 0.9901551

```

```

$parameters
      test      name.t ntr StudentizedRange alpha
      Tukey Explanatory  4      3.808798 0.05

```

```

$means
      Response      std r      se      Min      Max      Q25      Q50
x1  1.7620153 1.0505466 10 0.2599652 0.4504476 3.972459 1.1175485 1.44911720
x2  0.2841495 0.7532422 10 0.2599652 -0.4729986 1.985826 -0.3497379 0.21347543
x3  0.2197337 0.7019368 10 0.2599652 -0.6150452 1.574903 -0.2436023 -0.04493909
x4 -0.2890524 0.7345336 10 0.2599652 -1.9769014 0.684072 -0.5394534 -0.07741642
      Q75
x1 2.07491533
x2 0.64579865
x3 0.53544151
x4 0.04323417

```

```

$comparison
NULL

$groups
      Response groups
x1  1.7620153      a
x2  0.2841495      b
x3  0.2197337      b
x4 -0.2890524      b

attr("class")
[1] "group"

```

This output is nice because it labels groups based on which groups belong together! **This will be important for plotting, and is cumbersome if you have a lot of means.**

11.5 Plotting our ANOVA results

When we plot our ANOVAs, we want to show which mean is different from all of the others. Below, I will show the full pipeline for performing an ANOVA on a dataset and plotting the data at the end using our `data` object.

```

data_aov <- aov(Response ~ Explanatory,data)

summary(data_aov)

```

```

              Df Sum Sq Mean Sq F value    Pr(>F)
Explanatory    3  23.40   7.801    11.54 1.89e-05 ***
Residuals     36  24.33   0.676
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Now, we need to summarize things based on our `plyr` function.

```

# summarize by group
summary_data <- ddply(data, "Explanatory", summarise,
  N = length(Response),
  mean = mean(Response),
  sd = sd(Response),
  se = sd / sqrt(N))

```


Now, we can use the aforementioned Tukey results to assign to groups. Because `agricolae` functions define the groups in a `$groups` slot, we can pull out these data and perform some minimal transformations to make them ready to plot. *This will save us a lot of time and hassle.*

```
# note first group must be EXACT MATCH to your summary_data object
# groups are saved in the Tukey object
# this is true for Tukey later as well

# the following is a function that will make the significant label table
sig.label.maker <- function(tukey_test, group_name){
  sig.labels <- tukey_test$groups |>
    # convert to a data.frame
    as.data.frame() |>
    # create a new column - place rownames into the column
    # converts to a format better for ggplot
    mutate(Explanatorys = rownames(tukey_test$groups)) |>
    # rename column to prevent confusion
    # specify dplyr; default function may be from plyr and not work
    dplyr::rename(significance = groups)
  colnames(sig.labels)[which(colnames(sig.labels) == "Explanatorys")] <- group_name
  return(sig.labels)
}

# pull out the groups slot from tukey
# same for Kruskal later on!
sig.labels <- sig.label.maker(tukey_data_aov, "Explanatory")

sig.labels
```


| | Response | significance | Explanatory |
|----|------------|--------------|-------------|
| x1 | 1.7620153 | a | x1 |
| x2 | 0.2841495 | b | x2 |
| x3 | 0.2197337 | b | x3 |
| x4 | -0.2890524 | b | x4 |

Note that in the above, the row labels and the column `Explanatory` are identical. We moved these data into a column to make it easier to use `ggplot`. If you want letters to be different - capitalized or something else - you will have to do this yourself. Now we can plot our data and add the labels!

```

ggplot(summary_data, # plot summary data
  # Define plotting - x by group, y is mean, grouping by group
  aes(x = Explanatory, y = mean, group = Explanatory)) +
# add points to plot for y values
geom_point() +
# add error bars around points
geom_errorbar(data = summary_data,
  # define error bars
  aes(ymin = mean - 2*se, ymax = mean+2*se,
    # define color, size
    color = Explanatory), width = 0.1) +
# add line at average for the main group
# this is not always known - nor requires!
geom_hline(yintercept = 0, col = "black", linewidth = 0.5) +
# set vertical limits for plot
ylim(c(-3,3)) +
# make it a classic theme - more legible
theme_classic() +
# add text to plot
geom_text(data = sig.labels,
  # make bold
  fontface = "bold",
  # define where labels should go
  aes(x = Explanatory,
    # define height of label
    y = -2,
    # what are the labels?
    label = paste0(significance))) +
xlab("Explanatory Name") +
ylab("Mean") +
# remove legend - not needed here
theme(legend.position = "none",
  # make label text vertical, easier to read
  axis.text.x = element_text(angle = 90,
    # vertical offset of text
    vjust = 0.5,
    # text size
    size = 12))

```



Note that I place the labels below here, but they could also be placed above. You have to position them based on your own judgment.

Addendum: if you see a mean labeled **a** and **b**, it would be statistically indistinguishable from means labeled **a** and means labeled **b**, despite means with only an **a** or a **b** being different from each other.

11.6 Kruskal-Wallis tests

The Kruskal-Wallis test is the non-parametric version of an ANOVA. To demonstrate this, we will be creating a non-normal distribution by pulling random values from a *uniform* distribution, using the random uniform function `runif`. **Note** we are rounding the data here to make it more similar to non-normal datasets you may encounter, and to increase the probability of ties.

```
## COPY THIS WHOLE CHUNK ##
# will load example data

set.seed(8675309)

for(i in 1:4){
  x <- runif(10, min = -1, max = 1) |>
    round(2)
```

```

if(i == 1){
  x <- runif(10, min = 1, max = 2) |>
    round(2)
  data <- x |> as.data.frame()
  colnames(data) <- "Response"
  data$Explanatory <- paste0("x",i)
}else{
  newdat <- x |> as.data.frame()
  colnames(newdat) <- "Response"
  newdat$Explanatory <- paste0("x",i)
  data <- rbind(data,newdat)
}
}


# split into "typical" table
expanded_data <- NULL
expanded_data$x1 <- data$Response[which(data$Explanatory=="x1")]
expanded_data$x2 <- data$Response[which(data$Explanatory=="x2")]
expanded_data$x3 <- data$Response[which(data$Explanatory=="x3")]
expanded_data$x4 <- data$Response[which(data$Explanatory=="x4")]

expanded_data <- expanded_data |>
  as.data.frame()

# summarize by group
summary_data <- dplyr::ddply(data, "Explanatory", summarise,
  N = length(Response),
  mean = mean(Response),
  sd = sd(Response),
  se = sd / sqrt(N))


ggplot(summary_data, aes(x = Explanatory, y = mean, group = Explanatory)) +
  geom_point() +
  geom_errorbar(data = summary_data, aes(ymin = mean - 2*se, ymax = mean+2*se,
    color = Explanatory), width = 0.1) +
  geom_hline(yintercept = 0, col = "black", linewidth = 0.5) +
  ylim(c(-3,3)) +
  theme_classic() +
  theme(legend.position = "none",
    axis.text.x = element_text(angle = 90, vjust = 0.5, size = 5))

```



We can confirm these data are non-normal with Shapiro-Wilk tests and histograms. Demonstrating with the first column only:

```
hist(data[,1])
```



```
shapiro.test(data[,1])
```

Shapiro-Wilk normality test

```
data: data[, 1]
W = 0.93404, p-value = 0.02187
```

The above histogram appears non-normal, and the Shapiro-Wilk also indicates this is non-normal.

11.6.1 By hand

[Click here](#) to see the Kruskal-Wallis by hand.

First, we need to order all the data in the entire dataset. **This is easiest to do if we use the dataset with all data in a single column.**

```
data$rank <- rank(data$Response, ties.method = "average")
# view first couple of rows
head(data)
```

| | Response | Explanatory | rank |
|---|----------|-------------|------|
| 1 | 1.84 | x1 | 39 |
| 2 | 1.58 | x1 | 34 |
| 3 | 1.51 | x1 | 33 |
| 4 | 1.26 | x1 | 32 |
| 5 | 1.75 | x1 | 37 |
| 6 | 1.92 | x1 | 40 |

Now, we need to calculate the sum of the ranks for each category. The below function will take the sum of the ranks for the rows which match the condition of having the group be x1.

```
x1_sum <- sum(data$rank[which(data$Explanatory=="x1")])
x2_sum <- sum(data$rank[which(data$Explanatory=="x2")])
x3_sum <- sum(data$rank[which(data$Explanatory=="x3")])
x4_sum <- sum(data$rank[which(data$Explanatory=="x4")])
```

Now, we need to calculate our test statistic. The test statistic is H , with:

$$H = \frac{12}{N(N+1)} \cdot \sum \frac{R_j^2}{n_j} - 3(N+1)$$

In this equation, R is the sum of the ranks for a given category. This follows a χ^2 distribution with $k - 1$ degrees of freedom, with k referring to categories.

For these, we need to know what n is for each category.

```
n1 <- length(data$ranks[which(data$Explanatory=="x1")])
n2 <- length(data$ranks[which(data$Explanatory=="x2")])
n3 <- length(data$ranks[which(data$Explanatory=="x3")])
n4 <- length(data$ranks[which(data$Explanatory=="x4")])
```

Next, we can calculate the sums of the $\frac{R^2}{n}$ term.

```
r2_sum <- sum(x1_sum^2/n,
              x2_sum^2/n,
              x3_sum^2/n,
              x4_sum^2/n)
```

Now, we can calculate H .

```
N <- sum(n1, n2, n3, n4)
H <- ((12/(N*(N+1)))*r2_sum)-(3*(N+1))
H |> round(2)
```

```
[1] 57.43
```

Now, we can evaluate this with a χ^2 p value.

```
pchisq(q = H,
       df = ncol(data)-1,
       # remember, looking right!
       lower.tail = FALSE)
```

```
[1] 3.382831e-13
```

As we can see, the probability is extremely low with $p < 0.0001$. One distribution is different, and we can proceed with Tukey tests.

11.6.2 Using *R*

We can also use *R* for this test.

```
kruskal_data <- kruskal.test(Response ~ Explanatory, data)

print(kruskal_data)
```

Kruskal-Wallis rank sum test

```
data: Response by Explanatory
Kruskal-Wallis chi-squared = 22.149, df = 3, p-value = 6.073e-05
```

This test is similar, but not quite the same, as what we calculated above. The package *agricolae* provides a more in-depth Kruskal-Wallis test that gives us more metadata. We have to use the wrapper *with* to get this to work properly, per the help page.

```
# use "with", defining your dataset first
data_kruskal <- with(data,
  # define parameters for Kruskal-Wallis test
  # First, variable of interest
  kruskal(Response,
    # second, grouping variable
    Explanatory,
    # do you want group designations returned?
    group = TRUE,
    # what is this being done on?
    # must match dataset name in "with"!!!
    main = "data"))

# display results
# summary WILL NOT show the right thing
print(data_kruskal)
```

```
$statistics
      Chisq Df      p.chisq  t.value      MSD
22.14917   3 6.07312e-05 2.028094 7.252228

$parameters
      test p.adjusted      name.t ntr alpha
```



```

Kruskal-Wallis      none Explanatory    4  0.05

$means
  Response rank      std r   Min  Max    Q25   Q50   Q75
x1    1.633 35.50 0.2415252 10  1.21 1.92  1.5275  1.720 1.8025
x2   -0.075 14.20 0.6424130 10 -0.88 0.69 -0.6325 -0.105 0.5575
x3    0.020 16.15 0.5327080 10 -0.63 0.95 -0.3250 -0.110 0.3775
x4    0.045 16.15 0.6774011 10 -0.92 0.95 -0.3500 -0.135 0.6475

$comparison
NULL

$groups
  Response groups
x1    35.50      a
x3    16.15      b
x4    16.15      b
x2    14.20      b

attr(,"class")
[1] "group"

```

As we can see, the above gives us our group designations under the `$group` section. This is what we need to be able to plot things, as with ANOVA above. I do not repeat those steps here.

11.7 Homework: Chapter 11

Your homework is to complete problems 11.1, 11.2, 11.3, and 11.4. The first two will require ANOVA, and the last two will require Kruskal-Wallis tests. For each pair of tests, you must also complete the problem by hand in addition to using the default *R* methods. For *all* problems:

- State your hypothesis in words
- State your hypothesis mathematically (hint: μ)
- **Answer in complete sentences** and don't forget to round your answers.
- If you reject the null hypothesis, you must plot the results and label them to show which means are different.

12 ANOVA: Part 2

12.1 Two-way ANOVA

Previously, we discussed one-way ANOVAs, where we are looking at a single factor split across three or more groups and trying to determine if the means of these groups are equal (*i.e.*, $H_0 : \mu_1 = \mu_2 = \dots \mu_i$). ANOVA specifically allows us to analyze the variance of these different groups to ascertain which factors are most responsible for the variation we observe in the data. Because of the way ANOVA operates, we can actually test multiple different combinations of variables simultaneously in what we call a *two-way ANOVA*.

Don't forget to load your required packages - some we have used before, like `agricolae`, `plyr`, and `tidyverse`, but others are new for this section: `multcomp` and `nlme`! As a reminder, these packages are designed for the following:

- `agricolae`: originally written as a Master's thesis at the Universidad Nacional de Ingeniería (Lima, Perú), this package is designed to help with agricultural research.
- `plyr`: tools for common problems, including splitting data, applying functions across data, and combining datasets together.
- `tidyverse`: one we are already familiar with; a wrapper for installing `ggplot2`, `dplyr`, `tidyr`, `readr`, `purrr`, `tibble`, `stringr`, and `forcats`.
- `multcomp`: more in depth and better **MULTiple COMP**arisons via linear models and related models.
- `nlme`: a package for fitting Gaussian and **non-linear mixed-effect** models.
- `PMCMRplus`: a math package with post-hoc tests for Friedman's test

```
library(agricolae)
library(plyr)
library(tidyverse)

# NEW PACKAGES NEEDED
# Don't forget to install these on your machine
library(multcomp)
library(nlme)
library(PMCMRplus)
```

12.2 Designs

There are several different designs for two-way ANOVAs, and we will cover some of the most common designed here.

For these examples, we are going to randomly generated examples. I will refer to the variables as **Response** and **Explanatory** for simplicity's sake.

12.2.1 Randomized block design

Randomized block designs look at combinations of variables that could be affecting the results. More specifically, we are looking at two *strata* or *factors* and their effects on a *continuous* response variable.

```
set.seed(8675309)

# random example

# Blocking variable
Blocking_Variable <- c("Group 1", "Group 2", "Group 3")

# explanatory variables
# these are your columns
# these are your primary hypothesis
Explanatory_1 <- c(10.1, 9.4, 11.1)
Explanatory_2 <- c(12, 13.0, 15.4)
Explanatory_3 <- c(11.2, 10.1, 11.9)

# create "data table" as we normally see it
# combine all columns
data_expanded <- cbind(Blocking_Variable,
                        Explanatory_1,
                        Explanatory_2,
                        Explanatory_3) |>
  as.data.frame() # create data frame

data_expanded
```

| | Blocking_Variable | Explanatory_1 | Explanatory_2 | Explanatory_3 |
|---|-------------------|---------------|---------------|---------------|
| 1 | Group 1 | 10.1 | 12 | 11.2 |
| 2 | Group 2 | 9.4 | 13 | 10.1 |
| 3 | Group 3 | 11.1 | 15.4 | 11.9 |

Note that this table is in the format that we most often see, but we need to reshape these data to make it easier for us to perform our analyses. I created the data here as a matrix with named columns and rows; the following code may need to be adjusted if you do things differently.

```
# expand to "long" format
# if not done earlier, convert to data frame

data <- data_expanded |>
  # !by column for aggregating
  # names_to = what to name column aggregation
  # values_to = what the measurements should be called
  pivot_longer(!Blocking_Variable, names_to = "Explanatory_Variables", values_to = "Measurements")

data
```

```
# A tibble: 9 x 3
  Blocking_Variable Explanatory_Variables Measurements
  <chr>            <chr>            <chr>
1 Group 1         Explanatory_1      10.1
2 Group 1         Explanatory_2      12
3 Group 1         Explanatory_3      11.2
4 Group 2         Explanatory_1       9.4
5 Group 2         Explanatory_2      13
6 Group 2         Explanatory_3      10.1
7 Group 3         Explanatory_1      11.1
8 Group 3         Explanatory_2      15.4
9 Group 3         Explanatory_3      11.9
```

Now we can do our ANOVA. *Note* that I put `factor` around the blocking variable.

```
# mark block by factor
# best to always use
data_aov <- aov(Measurements ~ Explanatory_Variables + factor(Blocking_Variable), data)

summary(data_aov)
```

```
              Df Sum Sq Mean Sq F value Pr(>F)
Explanatory_Variables    2 17.182   8.591  14.412 0.0149 *
factor(Blocking_Variable) 2  6.829   3.414   5.728 0.0670 .
Residuals                4  2.384   0.596
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

In this particular example, the blocking variable does not significantly differ, however the explanatory variable does differ.

Remember, the columns represent your *primary hypothesis*. You will only plot your results if your *primary hypothesis is significant*!

Given that our primary null hypothesis is rejected (that is to say, not all means are equal), we need to plot our results.

To determine which mean(s) differ, we will use a Tukey Test. Unfortunately, the `agricolae` function `HSD.test` does not work as well for these multi-directional ANOVAs, so we need to use `TukeyHSD`.

```
tukey_data_aov <- TukeyHSD(data_aov)

tukey_data_aov
```

```
Tukey multiple comparisons of means
 95% family-wise confidence level
```

```
Fit: aov(formula = Measurements ~ Explanatory_Variables + factor(Blocking_Variable), data = d)
```

```
$Explanatory_Variables
```

| | diff | lwr | upr | p adj |
|-----------------------------|-----------|-----------|------------|-----------|
| Explanatory_2-Explanatory_1 | 3.266667 | 1.019919 | 5.513414 | 0.0144034 |
| Explanatory_3-Explanatory_1 | 0.866667 | -1.380081 | 3.113414 | 0.4333834 |
| Explanatory_3-Explanatory_2 | -2.400000 | -4.646748 | -0.1532523 | 0.0406301 |

```
$`factor(Blocking_Variable)`
```

| | diff | lwr | upr | p adj |
|-----------------|-----------|-----------|----------|-----------|
| Group 2-Group 1 | -0.266667 | -2.513414 | 1.980081 | 0.9082398 |
| Group 3-Group 1 | 1.700000 | -0.546747 | 3.946748 | 0.1118461 |
| Group 3-Group 2 | 1.966667 | -0.280081 | 4.213414 | 0.0745795 |

As we can see above, each pairwise comparison is given a p value for the level of difference. We need to manually label these groups based on these p values, with groups being considered different if $p < 0.05$. We can do this as follows, but unfortunately, we need to do it by hand since we don't have a short-form code (*yet*) for this conversion.

```
# change Explanatory_Variables to your data
Treatments <- unique(data$Explanatory_Variables)

sig_labels <- Treatments |>
```

```
as.data.frame() |>
mutate(Significance = rep(NA, length(Treatments)))

# Change Explanatory_Variables to your data
colnames(sig_labels) <- c("Explanatory_Variables", # MUST BE SAME AS DATA
                          "Significance")

tukey_data_aov$Explanatory_Variables
```

| | diff | lwr | upr | p adj |
|-----------------------------|-----------|-----------|------------|------------|
| Explanatory_2-Explanatory_1 | 3.266667 | 1.019919 | 5.513414 | 0.01440339 |
| Explanatory_3-Explanatory_1 | 0.866667 | -1.380081 | 3.113414 | 0.4338343 |
| Explanatory_3-Explanatory_2 | -2.400000 | -4.646748 | -0.1532523 | 0.04063012 |

NOTE that I am going to have to adjust column names and variable names a few times. Based on the above, we can see that `Explanatory_3` and `Explanatory_1` are not different from each other, but everything else is in a different group relative to each other. We can label these by hand.

```
sig_labels$Significance <- c("A", "B", "A")

sig_labels
```

| | Explanatory_Variables | Significance |
|---|-----------------------|--------------|
| 1 | Explanatory_1 | A |
| 2 | Explanatory_2 | B |
| 3 | Explanatory_3 | A |

As we can see above, now only `Explanatory_2` is given a different letter category.

Now, we can plot these different factors.

```
# summarize by group
# slight adjustment from previous
summary_data <- ddply(data, "Explanatory_Variables", summarise,
                      N = length(as.numeric(Measurements)),
                      mean = mean(as.numeric(Measurements)),
                      sd = sd(as.numeric(Measurements)),
                      se = sd / sqrt(N))

summary_data
```


| | Explanatory_Variables | N | mean | sd | se |
|---|-----------------------|---|----------|-----------|-----------|
| 1 | Explanatory_1 | 3 | 10.20000 | 0.8544004 | 0.4932883 |
| 2 | Explanatory_2 | 3 | 13.46667 | 1.7473790 | 1.0088497 |
| 3 | Explanatory_3 | 3 | 11.06667 | 0.9073772 | 0.5238745 |

```
# SET Y LIMITS
# change based on observed data
ylims <- c(0, 20)

# set label height, can change before plotting function
label_height <- 4

ggplot(summary_data, # plot summary data
  # Define plotting - x by group, y is mean, grouping by group
  aes(x = Explanatory_Variables, y = mean)) +
  # add points to plot for y values
  geom_point() +
  # add error bars around points
  geom_errorbar(data = summary_data,
    # define error bars
    aes(ymin = mean - 2*se, ymax = mean+2*se),
    # width of bar
    width = 0.1) +
  # set vertical limits for plot
  ylim(ylims) +
  # make it a classic theme - more legible
  theme_classic() +
  # add text to plot
  geom_text(data = sig_labels,
    # make bold
    fontface = "bold",
    # define where labels should go
    aes(x = Explanatory_Variables,
      # define height of label
      y = 4,
      # what are the labels?
      label = paste0(Significance))) +
  xlab("Treatment") +
  ylab("Mean") +
  # remove legend - not needed here
  theme(legend.position = "none",
    # make label text vertical, easier to read
```

```
axis.text.x = element_text(angle = 90,
                             # vertical offset of text
                             vjust = 0.5,
                             # text size
                             size = 12))
```



12.2.2 Repeated measures

Now, we are going to do a *repeated measures ANOVA*, where we have the same individuals being measured multiple times. Consider the following imaginary dataset:

```
Visit_1 <- c(5.5,6.2,5.8)
Visit_2 <- c(4.6,5.4,5.2)
Visit_3 <- c(3.8,4.0,3.9)

Individuals <- c(paste0("Individual", " ",1:3))

data <- cbind(Individuals,
              Visit_1,
              Visit_2,
              Visit_3) |>
as.data.frame() |>
```



```

pivot_longer(!Individuals,
              names_to = "Visits",
              values_to = "Measurements")

data

```

```

# A tibble: 9 x 3
  Individuals Visits Measurements
  <chr>      <chr>    <chr>
1 Individual 1 Visit_1 5.5
2 Individual 1 Visit_2 4.6
3 Individual 1 Visit_3 3.8
4 Individual 2 Visit_1 6.2
5 Individual 2 Visit_2 5.4
6 Individual 2 Visit_3 4
7 Individual 3 Visit_1 5.8
8 Individual 3 Visit_2 5.2
9 Individual 3 Visit_3 3.9

```

We need to perform the ANOVA again, but we need to account for the factor of which locations are repeated.

```

repeated_aov <- aov(Measurements ~ factor(Visits) + Error(factor(Individuals)), data)

summary(repeated_aov)

```

```

Error: factor(Individuals)
      Df Sum Sq Mean Sq F value Pr(>F)
Residuals  2 0.4867  0.2433

Error: Within
      Df Sum Sq Mean Sq F value    Pr(>F)
factor(Visits)  2  5.687  2.8433  89.79 0.000475 ***
Residuals      4  0.127  0.0317

---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Unfortunately, because of the model this is, we cannot perform a Tukey Test on the “object” that is created from this ANOVA analysis. We can, however, approach this from a different

direction and get our Tukey results ([thanks to Henrik on StackOverflow!](#)). For this to work, we need to install the packages `nlme` and `multcomp`.

```
# ensure data is proper format
data$Individuals <- as.factor(data$Individuals)
data$Visits <- as.factor(data$Visits)
data$Measurements <- as.numeric(data$Measurements)
```

The next part of the code fits a *linear model* to the data. A linear model, which we will cover later in the class, is mathematically very similar to an ANOVA. However, we can data from this model and extract the ANOVA data to understand more about the interactions. We need to use a linear model for this to account with the relationships between the two.

```
# fit a linear mixed-effects model
# similar to ANOVA

lme_data <- lme(Measurements ~ Visits,
               data = data,
               # define repeated section
               random = ~1|Individuals)

# perform ANOVA on model
anova(lme_data)
```

| | numDF | denDF | F-value | p-value |
|-------------|-------|-------|----------|---------|
| (Intercept) | 1 | 4 | 900.1654 | <.0001 |
| Visits | 2 | 4 | 89.7894 | 5e-04 |

As we can see above, we can get the ANOVA results from this linear mixed-effects model fit to the dataset. Now, we need to know post-hoc which sets are different:

```
lme_data |>
  # "general linear hypothesis"
  # define a comparison to make
  # can add corrections like test = adjusted (type = "bonferroni")
  glht(linfct = mcp(Visits = "Tukey")) |>
  # return a summary of the above
  summary()
```

Simultaneous Tests for General Linear Hypotheses

Multiple Comparisons of Means: Tukey Contrasts

```
Fit: lme.formula(fixed = Measurements ~ Visits, data = data, random = ~1 |  
  Individuals)
```

Linear Hypotheses:

| | Estimate | Std. Error | z value | Pr(> z) |
|------------------------|----------|------------|---------|------------|
| Visit_2 - Visit_1 == 0 | -0.7667 | 0.1453 | -5.277 | <1e-06 *** |
| Visit_3 - Visit_1 == 0 | -1.9333 | 0.1453 | -13.306 | <1e-06 *** |
| Visit_3 - Visit_2 == 0 | -1.1667 | 0.1453 | -8.030 | <1e-06 *** |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
(Adjusted p values reported -- single-step method)

We can see that every visit is different.

```
# manually labeling  
  
sig_levels_repeated <- matrix(data = c("Visit_1", "A",  
                                       "Visit_2", "B",  
                                       "Visit_3", "C"),  
                              byrow = T, ncol = 2) |> as.data.frame()  
  
# make labels match!  
colnames(sig_levels_repeated) <- c("Visits", "Significance")
```

Let's plot these. *Note* that we are not summarizing these the same way, since things are varying based on individual as well.


Note: For reasons I am not certain, you need to put the locations and then `ggplot` uses these colors to define everything. I really don't know why this is happening, so if you have a solution, let me know.

```
summary_data <- ddply(data, "Visits", summarise,  
                      N = length(as.numeric(Measurements)),  
                      mean = mean(as.numeric(Measurements)),  
                      sd = sd(as.numeric(Measurements)),  
                      se = sd / sqrt(N))
```

```

ggplot(summary_data,
       aes(x = Visits, y = mean)) +
  geom_point() +
  geom_errorbar(data = summary_data,
               # define error bars
               aes(ymin = mean - 2*se, ymax = mean+2*se),
               # width of bar
               width = 0.1) +
  # set vertical limits for plot
  ylim(c(0,10)) +
  # make it a classic theme - more legible
  theme_classic() +
  # add text to plot
  geom_text(data = sig_levels_repeated,
           # make bold
           fontface = "bold",
           # define where labels should go
           aes(x = Visits,
              # define height of label
              y = 2,
              # what are the labels?
              label = paste0(Significance))) +
  xlab("Visits") +
  ylab("Measurement") +
  # remove legend - not needed here
  theme(legend.position = "none",
        # make label text vertical, easier to read
        axis.text.x = element_text(angle = 90,
                                     # vertical offset of text
                                     vjust = 0.5,
                                     # text size
                                     size = 12))

```



12.2.3 Factorial ANOVA

Mathematically, a factorial ANOVA is the same as a randomized block ANOVA; please see that section for information on how to run this test.

12.2.4 ANOVA with interaction

Sometimes when we running a model, we want to look for *interactive effects*. Interactive effects are situations where one (or both) variables on their own do not effect the data, but there is a cumulative effect between variables that effects things. Let's look at an example, based on our initial example but with the data altered.

```
set.seed(8675309)

# we are using data from the randomized block ANOVA again

data_expanded
```

| | Blocking_Variable | Explanatory_1 | Explanatory_2 | Explanatory_3 |
|---|-------------------|---------------|---------------|---------------|
| 1 | Group 1 | 10.1 | 12 | 11.2 |
| 2 | Group 2 | 9.4 | 13 | 10.1 |
| 3 | Group 3 | 11.1 | 15.4 | 11.9 |

```

### YOU DO NOT NEED TO DO THIS
### CREATING DATA FOR EXAMPLE

data_expanded$Explanatory_1 <- as.numeric(data_expanded$Explanatory_1)
data_expanded$Explanatory_2 <- as.numeric(data_expanded$Explanatory_2)
data_expanded$Explanatory_3 <- as.numeric(data_expanded$Explanatory_3)

# create some pseudorandom data
# [, -1] excludes first column - group data
data_expanded2 <- cbind(Blocking_Variable,
                        data_expanded[, -1] - 0.75)
data_expanded3 <- cbind(Blocking_Variable,
                        data_expanded[, -1]*1.05)

data_expanded <- rbind(data_expanded,
                        data_expanded2,
                        data_expanded3)

# expand to "long" format
data <- data_expanded |>
  # convert to data frame
  as.data.frame() |>
  # !by column for aggregating
  # names_to = what to name column aggregation
  # values_to = what the measurements should be called
  pivot_longer(!Blocking_Variable, names_to = "Treatments", values_to = "Measurements")

# specifying factor to be safe

interactive_aov <- aov(Measurements ~ factor(Treatments) +
                       factor(Blocking_Variable) +
                       factor(Treatments)*factor(Blocking_Variable),
                       data)

summary(interactive_aov)

```

| | Df | Sum Sq | Mean Sq | F value | Pr(>F) |
|--|----|--------|---------|---------|----------|
| factor(Treatments) | 2 | 53.28 | 26.640 | 59.680 | 1.14e-08 |
| factor(Blocking_Variable) | 2 | 21.18 | 10.588 | 23.719 | 9.01e-06 |
| factor(Treatments):factor(Blocking_Variable) | 4 | 7.39 | 1.848 | 4.141 | 0.015 |
| Residuals | 18 | 8.03 | 0.446 | | |

```

factor(Treatments) ***
factor(Blocking_Variable) ***
factor(Treatments):factor(Blocking_Variable) *
Residuals
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

As we can see above, we have very significant effects for `Treatment` and `Blocking_Variable`, but a less significant effect for the interaction between the two. **Remember - we only need to plot our primary hypothesis.** Note however, that Tukey gives us our differences and *p* values for each set of tests and comparisons:

```
TukeyHSD(interactive_aov)
```

```

Tukey multiple comparisons of means
 95% family-wise confidence level

```

```
Fit: aov(formula = Measurements ~ factor(Treatments) + factor(Blocking_Variable) + factor(Treatments):factor(Blocking_Variable))
```

```

$`factor(Treatments)`
              diff          lwr          upr          p adj
Explanatory_2-Explanatory_1  3.3211111  2.51730768  4.124915  0.0000000
Explanatory_3-Explanatory_1  0.8811111  0.07730768  1.684915  0.0304521
Explanatory_3-Explanatory_2 -2.4400000 -3.24380343 -1.636197  0.0000011

```

```

$`factor(Blocking_Variable)`
              diff          lwr          upr          p adj
Group 2-Group 1 -0.2711111 -1.0749145  0.5326923  0.6710682
Group 3-Group 1  1.7283333  0.9245299  2.5321368  0.0000928
Group 3-Group 2  1.9994444  1.1956410  2.8032479  0.0000159

```

```

$`factor(Treatments):factor(Blocking_Variable)`
              diff          lwr          upr          p adj
Explanatory_2:Group 1-Explanatory_1:Group 1  1.931667e+00  0.02027781
Explanatory_3:Group 1-Explanatory_1:Group 1  1.118333e+00 -0.79305553
Explanatory_1:Group 2-Explanatory_1:Group 1 -7.116667e-01 -2.62305553
Explanatory_2:Group 2-Explanatory_1:Group 1  2.948333e+00  1.03694447
Explanatory_3:Group 2-Explanatory_1:Group 1 -1.776357e-15 -1.91138886
Explanatory_1:Group 3-Explanatory_1:Group 1  1.016667e+00 -0.89472219
Explanatory_2:Group 3-Explanatory_1:Group 1  5.388333e+00  3.47694447
Explanatory_3:Group 3-Explanatory_1:Group 1  1.830000e+00 -0.08138886
Explanatory_3:Group 1-Explanatory_2:Group 1 -8.133333e-01 -2.72472219

```

| | | | | |
|---------------------|-----------------------|---|---------------|-------------|
| Explanatory_1:Group | 2-Explanatory_2:Group | 1 | -2.643333e+00 | -4.55472219 |
| Explanatory_2:Group | 2-Explanatory_2:Group | 1 | 1.016667e+00 | -0.89472219 |
| Explanatory_3:Group | 2-Explanatory_2:Group | 1 | -1.931667e+00 | -3.84305553 |
| Explanatory_1:Group | 3-Explanatory_2:Group | 1 | -9.150000e-01 | -2.82638886 |
| Explanatory_2:Group | 3-Explanatory_2:Group | 1 | 3.456667e+00 | 1.54527781 |
| Explanatory_3:Group | 3-Explanatory_2:Group | 1 | -1.016667e-01 | -2.01305553 |
| Explanatory_1:Group | 2-Explanatory_3:Group | 1 | -1.830000e+00 | -3.74138886 |
| Explanatory_2:Group | 2-Explanatory_3:Group | 1 | 1.830000e+00 | -0.08138886 |
| Explanatory_3:Group | 2-Explanatory_3:Group | 1 | -1.118333e+00 | -3.02972219 |
| Explanatory_1:Group | 3-Explanatory_3:Group | 1 | -1.016667e-01 | -2.01305553 |
| Explanatory_2:Group | 3-Explanatory_3:Group | 1 | 4.270000e+00 | 2.35861114 |
| Explanatory_3:Group | 3-Explanatory_3:Group | 1 | 7.116667e-01 | -1.19972219 |
| Explanatory_2:Group | 2-Explanatory_1:Group | 2 | 3.660000e+00 | 1.74861114 |
| Explanatory_3:Group | 2-Explanatory_1:Group | 2 | 7.116667e-01 | -1.19972219 |
| Explanatory_1:Group | 3-Explanatory_1:Group | 2 | 1.728333e+00 | -0.18305553 |
| Explanatory_2:Group | 3-Explanatory_1:Group | 2 | 6.100000e+00 | 4.18861114 |
| Explanatory_3:Group | 3-Explanatory_1:Group | 2 | 2.541667e+00 | 0.63027781 |
| Explanatory_3:Group | 2-Explanatory_2:Group | 2 | -2.948333e+00 | -4.85972219 |
| Explanatory_1:Group | 3-Explanatory_2:Group | 2 | -1.931667e+00 | -3.84305553 |
| Explanatory_2:Group | 3-Explanatory_2:Group | 2 | 2.440000e+00 | 0.52861114 |
| Explanatory_3:Group | 3-Explanatory_2:Group | 2 | -1.118333e+00 | -3.02972219 |
| Explanatory_1:Group | 3-Explanatory_3:Group | 2 | 1.016667e+00 | -0.89472219 |
| Explanatory_2:Group | 3-Explanatory_3:Group | 2 | 5.388333e+00 | 3.47694447 |
| Explanatory_3:Group | 3-Explanatory_3:Group | 2 | 1.830000e+00 | -0.08138886 |
| Explanatory_2:Group | 3-Explanatory_1:Group | 3 | 4.371667e+00 | 2.46027781 |
| Explanatory_3:Group | 3-Explanatory_1:Group | 3 | 8.133333e-01 | -1.09805553 |
| Explanatory_3:Group | 3-Explanatory_2:Group | 3 | -3.558333e+00 | -5.46972219 |
| | | | upr | p adj |
| Explanatory_2:Group | 1-Explanatory_1:Group | 1 | 3.84305553 | 0.0464895 |
| Explanatory_3:Group | 1-Explanatory_1:Group | 1 | 3.02972219 | 0.5317567 |
| Explanatory_1:Group | 2-Explanatory_1:Group | 1 | 1.19972219 | 0.9173224 |
| Explanatory_2:Group | 2-Explanatory_1:Group | 1 | 4.85972219 | 0.0010200 |
| Explanatory_3:Group | 2-Explanatory_1:Group | 1 | 1.91138886 | 1.0000000 |
| Explanatory_1:Group | 3-Explanatory_1:Group | 1 | 2.92805553 | 0.6439270 |
| Explanatory_2:Group | 3-Explanatory_1:Group | 1 | 7.29972219 | 0.0000003 |
| Explanatory_3:Group | 3-Explanatory_1:Group | 1 | 3.74138886 | 0.0667267 |
| Explanatory_3:Group | 1-Explanatory_2:Group | 1 | 1.09805553 | 0.8457970 |
| Explanatory_1:Group | 2-Explanatory_2:Group | 1 | -0.73194447 | 0.0032271 |
| Explanatory_2:Group | 2-Explanatory_2:Group | 1 | 2.92805553 | 0.6439270 |
| Explanatory_3:Group | 2-Explanatory_2:Group | 1 | -0.02027781 | 0.0464895 |
| Explanatory_1:Group | 3-Explanatory_2:Group | 1 | 0.99638886 | 0.7519790 |
| Explanatory_2:Group | 3-Explanatory_2:Group | 1 | 5.36805553 | 0.0001577 |
| Explanatory_3:Group | 3-Explanatory_2:Group | 1 | 1.80972219 | 0.9999999 |

| | | | | |
|---------------------|-----------------------|---|-------------|-----------|
| Explanatory_1:Group | 2-Explanatory_3:Group | 1 | 0.08138886 | 0.0667267 |
| Explanatory_2:Group | 2-Explanatory_3:Group | 1 | 3.74138886 | 0.0667267 |
| Explanatory_3:Group | 2-Explanatory_3:Group | 1 | 0.79305553 | 0.5317567 |
| Explanatory_1:Group | 3-Explanatory_3:Group | 1 | 1.80972219 | 0.9999999 |
| Explanatory_2:Group | 3-Explanatory_3:Group | 1 | 6.18138886 | 0.0000097 |
| Explanatory_3:Group | 3-Explanatory_3:Group | 1 | 2.62305553 | 0.9173224 |
| Explanatory_2:Group | 2-Explanatory_1:Group | 2 | 5.57138886 | 0.0000766 |
| Explanatory_3:Group | 2-Explanatory_1:Group | 2 | 2.62305553 | 0.9173224 |
| Explanatory_1:Group | 3-Explanatory_1:Group | 2 | 3.63972219 | 0.0947841 |
| Explanatory_2:Group | 3-Explanatory_1:Group | 2 | 8.01138886 | 0.0000000 |
| Explanatory_3:Group | 3-Explanatory_1:Group | 2 | 4.45305553 | 0.0047491 |
| Explanatory_3:Group | 2-Explanatory_2:Group | 2 | -1.03694447 | 0.0010200 |
| Explanatory_1:Group | 3-Explanatory_2:Group | 2 | -0.02027781 | 0.0464895 |
| Explanatory_2:Group | 3-Explanatory_2:Group | 2 | 4.35138886 | 0.0069891 |
| Explanatory_3:Group | 3-Explanatory_2:Group | 2 | 0.79305553 | 0.5317567 |
| Explanatory_1:Group | 3-Explanatory_3:Group | 2 | 2.92805553 | 0.6439270 |
| Explanatory_2:Group | 3-Explanatory_3:Group | 2 | 7.29972219 | 0.0000003 |
| Explanatory_3:Group | 3-Explanatory_3:Group | 2 | 3.74138886 | 0.0667267 |
| Explanatory_2:Group | 3-Explanatory_1:Group | 3 | 6.28305553 | 0.0000070 |
| Explanatory_3:Group | 3-Explanatory_1:Group | 3 | 2.72472219 | 0.8457970 |
| Explanatory_3:Group | 3-Explanatory_2:Group | 3 | -1.64694447 | 0.0001097 |

I do not plot this here, but it would be similar to the other parts of this test.

12.3 Friedman's test

12.3.1 Using R

Friedman's test is a non-parametric alternative to a two-way ANOVA, so as you would guess, it can be painful to implement. We will use an altered version of the same test we've used before:

```
# set seed - make reproducible
set.seed(8675309)

### DO NOT NEED TO REPEAT THIS
### CREATING DATA FOR EXAMPLE

# new set of foods - this time, ten of them
Treatments <- c(paste("Treatment",1:10)) |>
  as.factor()
```

```
# pre-created data frame of locations from earlier
Blocking_Factor <- c(paste("Block", 1:10)) |>
  as.factor()

long_data <- crossing(Blocking_Factor, Treatments)

long_data$Measurements <- NA

for(i in 1:length(unique(long_data$Treatments))){
  subset_rows <- which(long_data$Treatments==long_data$Treatments[i])
  long_data$Measurements[subset_rows] <- runif(n = length(subset_rows),
                                              min = i-2, max = i+2) |>
    round(1)
}

long_data
```

```
# A tibble: 100 x 3
  Blocking_Factor Treatments Measurements
  <fct>           <fct>           <dbl>
1 Block 1        Treatment 1        -0.4
2 Block 1        Treatment 10         3.4
3 Block 1        Treatment 2         1.2
4 Block 1        Treatment 3         5.9
5 Block 1        Treatment 4         6.9
6 Block 1        Treatment 5          7
7 Block 1        Treatment 6         7.6
8 Block 1        Treatment 7         7.6
9 Block 1        Treatment 8         8.1
10 Block 1       Treatment 9         9.8
# i 90 more rows
```

Now that we have our expanded and randomized table, we can get started with our test.

Our calculation for the Friedman's test statistic Q (not to be confused with Tukey's q !) is:

$$Q = \frac{12}{nk(k+1)} \cdot \sum R_j^2 - 3n(k+1)$$

where n is the total number of individuals in each sample in the dataset, k is the number of groups, and R_j^2 is the sum of the ranks.

In this class, we will do this in *R*.

```
friedman_long_data <- friedman.test(y = long_data$Measurements,
                                   groups = long_data$Treatments,
                                   blocks = long_data$Blocking_Factor)

print(friedman_long_data)
```

Friedman rank sum test

data: long_data\$Measurements, long_data\$Treatments and long_data\$Blocking_Factor
 Friedman chi-squared = 79.983, df = 9, p-value = 1.629e-13

Note you will get a different answer if you are switching the blocks and the groups.

We can use the following, from package PMCMRplus, to find the adjacent and non-adjacent groups.

```
# find differences
frdAllPairsConoverTest(y = long_data$Measurements,
                       groups = long_data$Treatments,
                       blocks = long_data$Blocking_Factor,
                       p.adjust.method = "bonf")
```

Pairwise comparisons using Conover's all-pairs test for a two-way balanced complete block

data: y, groups and blocks

| | Treatment 1 | Treatment 10 | Treatment 2 | Treatment 3 | Treatment 4 |
|--------------|-------------|--------------|-------------|-------------|-------------|
| Treatment 10 | 1.00000 | - | - | - | - |
| Treatment 2 | 1.00000 | 1.00000 | - | - | - |
| Treatment 3 | 0.00094 | 0.10167 | 0.80301 | - | - |
| Treatment 4 | 3.8e-09 | 1.6e-06 | 3.4e-05 | 0.19031 | - |
| Treatment 5 | 3.2e-11 | 1.6e-08 | 4.0e-07 | 0.00637 | 1.00000 |
| Treatment 6 | < 2e-16 | 9.2e-16 | 2.5e-14 | 1.5e-09 | 0.00042 |
| Treatment 7 | < 2e-16 | < 2e-16 | 3.6e-16 | 2.0e-11 | 9.3e-06 |
| Treatment 8 | < 2e-16 | < 2e-16 | < 2e-16 | 6.1e-15 | 3.8e-09 |
| Treatment 9 | < 2e-16 | < 2e-16 | < 2e-16 | < 2e-16 | 1.1e-12 |
| | Treatment 5 | Treatment 6 | Treatment 7 | Treatment 8 | |
| Treatment 10 | - | - | - | - | |

| | | | | |
|-------------|---------|---------|---------|---------|
| Treatment 2 | - | - | - | - |
| Treatment 3 | - | - | - | - |
| Treatment 4 | - | - | - | - |
| Treatment 5 | - | - | - | - |
| Treatment 6 | 0.01886 | - | - | - |
| Treatment 7 | 0.00063 | 1.00000 | - | - |
| Treatment 8 | 4.0e-07 | 0.34617 | 1.00000 | - |
| Treatment 9 | 1.4e-10 | 0.00094 | 0.02676 | 1.00000 |

P value adjustment method: bonferroni

As we can see, some pairs are inseparable and others are separable. We can now plot as for the other problems.

12.4 Homework: Chapter 12

For problems 12.1, 12.2, 12.3, 12.4, and 12.5, state your hypotheses in sentence form and mathematically. Then, identify the appropriate ANOVA and perform the analysis. If you reject the null, complete a Tukey test and plot your results, showing letters denoting each group. **Note** that 12.4 requires a Friedman's test, but all other problems require some form of ANOVA.

Next, for problems 12.7, 12.8, and 12.9, identify the appropriate test and justify your reasoning. State the null and alternative hypothesis in word form and mathematically, and perform your analysis. If you perform an ANOVA *and* you reject the null hypothesis, plot your results and label the groups by letter.

Remember, only plot the results if you reject your *primary hypothesis*.

13 Correlation & regression

```
# load required libraries
```

```
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
```

```
v dplyr      1.1.4      v readr      2.1.5
```

```
v forcats    1.0.0      v stringr    1.5.1
```

```
v ggplot2    3.5.2      v tibble     3.3.0
```

```
v lubridate  1.9.4      v tidyr      1.3.1
```

```
v purrr      1.1.0
```

```
-- Conflicts ----- tidyverse_conflicts() --
```

```
x dplyr::filter() masks stats::filter()
```

```
x dplyr::lag()     masks stats::lag()
```

```
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

13.1 Introduction

When we are comparing *two continuous variables*, we use two forms of tests: *correlation* to understand *if* there is a relationship between two variables, and *linear regression* to determine *what* that relationships is.

Remember - “if” is always correlation, and “what is it” is always linear regression when choosing a test for an exam.

13.2 Correlation

Correlation - denoted by ρ (“rho”) and not to be confused with p - is a measure of how closely related to continuous variables appear to be. This can vary from a purely negative relationship to a purely positive relationship, such that $-1 \leq \rho \leq 1$ with $\rho = 0$ indicating a random relationship between data.


We can visualize these as follows:

```
### ILLUSTRATIVE PURPOSES ONLY
```

```
# create two random uniform distributions  
y <- runif(10)  
x <- runif(10)
```

For example, two things compared to themselves have a $\rho = 1$.

```
plot(y, y, pch = 19)
```




```
cor(y, y)
```

```
[1] 1
```

As we can see above, the correlation is 1.

Plotting by the negative will be a correlation of -1.

```
plot(y, -y, pch = 19)
```




```
cor(y, -y)
```

```
[1] -1
```

Lastly, two random variables plotted against each other should have $\rho \approx 0$.

```
plot(y, x,
      pch = 19,
      asp = 1) # aspect ratio
```



```
cor(y, x) |>
  round(2)
```

```
[1] 0.29
```

13.2.1 Pearson's

Pearson's correlation coefficient is our value for *parametric tests*. We often denote our correlation coefficient as r and not ρ for this particular test. It is calculated as follows:

$$r = \frac{\sum xy - \left(\frac{\sum x \sum y}{n}\right)}{\sqrt{\left(\sum x^2 - \frac{(\sum x)^2}{n}\right)\left(\sum y^2 - \frac{(\sum y)^2}{n}\right)}}$$


where x is variable 1, y is variable 2, and n is the total number of data point pairs.

In this class, we will be using *R* to calculate r , which is done using the command `cor`. To ensure we are using the correct method, we need to set `method = "pearson"`.

```
set.seed(8675309)

### EXAMPLE DATA
x <- c(1,2,5,3,4,5,8,7,9,6,10,12,15,20,25)
y <- c(2,5,4,3,8,6,4,2,8,9,15,13,10,18,19)

plot(x, y, pch = 19)
```



```
cor(x, y, method = "pearson") |>
  round(2)
```

```
[1] 0.86
```


As we can see, these data are fairly positively correlated. As x increases, so does y . But how *significant* is this relationship?

Well, we can calculate two things - the *amount of variation explained*, which is r^2 , and the *significance of the relationships*, which is determined via a t test and the equation $t = r\sqrt{\frac{n-2}{1-r^2}}$. This is a two-tailed distribution, with $df = n - 2$.

We can write a function to perform all of these options:

```
biol305_cor <- function(x=NA, y=NA, method = "pearson"){
  if(is.data.frame(x)==T){
    if(ncol(x)==2){
      r <- cor(x[,1], x[,2], method = method)
    }else{
      r <- cor(x, method = method)
    }

    r2 <- r
    r[r==1|r== -1] <- 0

    n <- 2*nrow(x)
  }else{
    r <- cor(x, y, method = method)

    n <- 2*length(x)
  }

  t_val <- r*sqrt((n-2)/(1-r^2))

  p <- pt(t_val, df = n - 2)

  p[p > 0.5] <- 1 - p[p > 0.5]
  p[p > 0.005] <- round(p[p > 0.005],2)
  p[p > 0.0005] <- round(p[p > 0.0005],3)
  p[p > 0.00005] <- round(p[p > 0.00005],4)
  p[p < 0.00005] <- "< 0.0001"

  if(is.data.frame(x)==T){
    print("Correlation:")
    print(round(r, 2))
    if(ncol(x) == 2){
      print(paste0("Degrees of freedom: ", n - 2))
      print(paste0("t value: ", round(t_val, 2)))
    }
  }
}
```

```

        print(paste0("P value: ", p))
    }
    if(ncol(x) > 2){
        print(paste0("Degrees of freedom: ", n - 2))
        print("")
        print("t value: ")
        print(round(t_val, 2))
        print("")
        print("P value: ")
        print(p)
    }
} else{
    print(paste0("Correlation: ", round(r, 2)))
    print(paste0("Degrees of freedom: ", n - 2))
    print(paste0("t value: ", round(t_val, 2)))
    print(paste0("P value: ", p))
}
}

```

Let's test our function.

```
biol305_cor(x, y, method = "pearson")
```

```

[1] "Correlation: 0.86"
[1] "Degrees of freedom: 28"
[1] "t value: 8.93"
[1] "P value: < 0.0001"

```

There we go! Our function printed out everything that we need.

13.2.2 Spearman's

Spearman's correlation is one of the non-parametric methods for our correlation tests. We can use this for ranked data or for non-parametric datasets. We do this the exact same way, except we change `method = "spearman"`.

13.2.3 Other non-parametric methods

To be expanded upon, but not necessary for the class at present.

13.3 Regression

Regression is used when we want to know *what* the relationship is between two variables. Regression operates similar to ANOVA and correlation, providing us with the nature of the relationship, the strength of the relationship, and gives us values for calculating the relationship. **For this class, we are only focusing on linear regression for relationships between linear variables.**

The equation for a regression line is often written as $y_i = \alpha + \beta x_i + e_i$, where α is the y intercept, β is the slope, and e is the error around each point. We will not perform regression calculations by hand in this class.

13.3.1 Parametric

We will use our previous established x and y datasets that are strongly positively correlated for this example. The equation for calculating a linear relationship is `lm`, which stands for “linear model”. This uses equations like ANOVA, but can also use two vectors of data.

```
xy_linear <- lm(y ~ x)
summary(xy_linear)
```

Call:

```
lm(formula = y ~ x)
```

Residuals:

| Min | 1Q | Median | 3Q | Max |
|---------|---------|---------|--------|--------|
| -5.1189 | -1.4838 | -0.5423 | 1.9757 | 5.7460 |

Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------|----------|------------|---------|--------------|
| (Intercept) | 2.1370 | 1.2825 | 1.666 | 0.12 |
| x | 0.7117 | 0.1169 | 6.086 | 3.87e-05 *** |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.964 on 13 degrees of freedom

Multiple R-squared: 0.7402, Adjusted R-squared: 0.7202

F-statistic: 37.04 on 1 and 13 DF, p-value: 3.867e-05

As we can see, this returned an ANOVA table to use that tells us the value and significance of our *intercept* as well as the value and significance of the the slope (here, shown as *x*; it will always show the explanatory variable in this slot for the name).

Looking at the above, we can see that the slope is not significantly non-zero with a $p = 0.12$, but that the slope is significantly non-zero with $p < 0.0001$. We also have our R^2 values returned, which is similar to the r we got for correlation. Indeed, our correlation was $r = 0.86$, with $r^2 = 0.74$, which is very similar to the **Multiple R-squared** shown in the above ANOVA table.

R has a built in function within **ggplot** that will add a linear model to our plot and will show error regions as well. First, we need to make sure our data are in a **data.frame**.

```
data <- data.frame(x, y)
```


```
data
```

| | x | y |
|----|----|----|
| 1 | 1 | 2 |
| 2 | 2 | 5 |
| 3 | 5 | 4 |
| 4 | 3 | 3 |
| 5 | 4 | 8 |
| 6 | 5 | 6 |
| 7 | 8 | 4 |
| 8 | 7 | 2 |
| 9 | 9 | 8 |
| 10 | 6 | 9 |
| 11 | 10 | 15 |
| 12 | 12 | 13 |
| 13 | 15 | 10 |
| 14 | 20 | 18 |
| 15 | 25 | 19 |

Next, we can plot the data in **ggplot**.

```
ggplot(data, aes(x = x, y = y)) +  
  geom_point() +  
  stat_smooth(method = "lm") +  
  theme_minimal()
```

```
`geom_smooth()` using formula = 'y ~ x'
```



Just like that, we have created a plot of our linear regression. **Note** that you should always plot the lines only within the extent of the data; this is harder to do in other programs, but *R* does it for us!

R can also allow us to predict different values using our linear model:

```
# must be in data frame format
test_data <- data.frame(1:25)


colnames(test_data) <- "x" # must be same as explanatory in data

predict(xy_linear, test_data)
```

| | | | | | | | |
|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 |
| 2.848692 | 3.560399 | 4.272105 | 4.983811 | 5.695517 | 6.407223 | 7.118929 | 7.830635 |
| 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 |
| 8.542341 | 9.254047 | 9.965753 | 10.677460 | 11.389166 | 12.100872 | 12.812578 | 13.524284 |
| 17 | 18 | 19 | 20 | 21 | 22 | 23 | 24 |
| 14.235990 | 14.947696 | 15.659402 | 16.371108 | 17.082814 | 17.794521 | 18.506227 | 19.217933 |
| 25 | | | | | | | |
| 19.929639 | | | | | | | |

Let's visualize these points for illustration's sake.

```
plot(x, y, pch = 19)
points(1:25, as.numeric(predict(xy_linear, test_data)),
      pch = 19, col = "red")
```



As we can see above, our points follow the line from the plot. By using this format, however, we can make predictions of value for any point we want.

13.3.2 Non-parametric

We are not doing non-parametric linear regression in this class.

13.4 Homework

13.4.1 Chapter 13

13.4.2 Chapter 14

14 Multivariate methods

14.1 Multivariate analyses

Multivariate analyses are commonly used to help understand the variation that exists in large, multidimensional datasets. These analyses can help us understand relationships among individuals and help us determine how many groups (and what groups) exist in a dataset.

For this particular walkthrough, we will be using data on the Northern Double-collared Sunbird *Cinnyris reichenowi sensu lato* complex from equatorial Afromontane regions (Cooper et al. 2021a, b).

```
library(tidyverse)

cinnyris <- read_csv("https://raw.githubusercontent.com/jacobccooper/biol305_unk/main/datasets/cinnyris")

male_cinnyris <- cinnyris |>
  mutate(Genus = as.factor(Genus)) |>
  mutate(Species = as.factor(Species)) |>
  mutate(Subspecies = as.factor(Subspecies)) |>
  filter(Age == "Adult") |>
  filter(Sex == "Male") |>
  select(-`Collection Name`, -Locality, -Country, -Age,
         -`Notes / Existing Damage Precluding Measurements`) |>
  na.omit()
```

14.2 Principal components

[Click here](#) for an overview of PCA

A very common method of visualizing variation in population is using principal components analyses (PCA). Principal components help collapse multidimensional data into a few dimensions that explain the most variation possible across the dataset.

For principal components, we want to use the library **vegan**.

```
library(vegan)
```

Loading required package: permute

The function for principal components in `vegan` is `rda`, which will only work on a whole dataframe with no non-numeric data. Thus, we must create a new object of only numerical data and then re-merge it with the original dataframe. The row order is maintained, so this will not be an issue.

```
male.cinnyris.data <- male_cinnyris |>
  select(-Genus, -Species, -Subspecies,
         -Collection, -Catalog, -Locality2, -Sex)

cinnyris.pca <- rda(male.cinnyris.data)

cinnyris.pca
```

Call: `rda(X = male.cinnyris.data)`

| | Inertia | Rank |
|---------------|---------|------|
| Total | 23 | |
| Unconstrained | 23 | 6 |


Inertia is variance

Eigenvalues for unconstrained axes:

| PC1 | PC2 | PC3 | PC4 | PC5 | PC6 |
|--------|-------|-------|-------|-------|-------|
| 15.620 | 4.443 | 1.940 | 0.877 | 0.076 | 0.040 |

Above, we can see that we now have a complex object that contains our PCA. We can use the function `biplot` to see how our data look and how our variables contributed to the principal components.


```
biplot(cinnyris.pca)
ordihull(cinnyris.pca,
         group = male_cinnyris$Subspecies)
```

We can also do a `ggplot` to color our data and see how it works. Let's combine our datasets so we can do this.

```
cinnyris.male.pca <- cbind(male_cinnyris,
                           # extract PC values
                           cinnyris.pca$CA$u)

ggplot(cinnyris.male.pca, aes(x = PC1, y = PC2)) +
  geom_point() +
  theme_classic()
```




To the naked eye, these data seem to fall into two major clusters. This can be examined again, this time assigning colors to one of our explanatory variables: subspecies.

```
ggplot(cinnyris.male.pca, aes(x = PC1, y = PC2, colour = Subspecies)) +
  geom_point() +
  stat_ellipse() +
  theme_classic()
```

Too few points to calculate an ellipse

Warning: Removed 1 row containing missing values or values outside the scale range (`geom_path()`).



As we can see, many of these birds fall into one of the two groups. Only one population - *genderuensis* - seems to really bridge both groups.

The following shows how each variable contributed to each principal component:

```
cinnyris.pca$CA$v
```

| | PC1 | PC2 | PC3 |
|------------------|-------------|--------------|-------------|
| Right.wing.chord | -0.61936877 | 0.498711666 | 0.60193981 |
| Tail.length | -0.69190828 | -0.709947486 | -0.12739287 |

| | | | |
|---|---------------|--------------|-------------|
| Culmen.length | -0.33598591 | 0.419621089 | -0.62835637 |
| Bill.depth..base.of.feathers.on.mandible. | -0.01633872 | 0.004353418 | -0.01231269 |
| Bill.width..base.of.feathers.on.maxilla. | -0.05865394 | 0.061714372 | -0.09962101 |
| Left.Tarsus | -0.14509132 | 0.259522190 | -0.46532983 |
| | PC4 | PC5 | |
| Right.wing.chord | 0.07304690 | -0.001241684 | |
| Tail.length | 0.02988534 | -0.008716903 | |
| Culmen.length | -0.53855805 | -0.161236784 | |
| Bill.depth..base.of.feathers.on.mandible. | -0.05951994 | 0.216257889 | |
| Bill.width..base.of.feathers.on.maxilla. | -0.08219923 | 0.962867346 | |
| Left.Tarsus | 0.83272259 | 0.006646209 | |
| | PC6 | | |
| Right.wing.chord | -0.0002699604 | | |
| Tail.length | -0.0062803042 | | |
| Culmen.length | -0.0125622441 | | |
| Bill.depth..base.of.feathers.on.mandible. | 0.9742958451 | | |
| Bill.width..base.of.feathers.on.maxilla. | -0.2212610835 | | |
| Left.Tarsus | 0.0399226008 | | |

We can also see the percent variance explained by each axis, which is contained within the eigenvectors of the PCA:

```
cinnyris.pca$CA$eig/sum(cinnyris.pca$CA$eig)
```

| PC1 | PC2 | PC3 | PC4 | PC5 | PC6 |
|-------------|-------------|-------------|-------------|-------------|-------------|
| 0.679234583 | 0.193214694 | 0.084367157 | 0.038143954 | 0.003291025 | 0.001748588 |

The first principal component explains 68% of the variation, and the second 19% of the variation.

14.3 Clustering

14.3.1 K-means clustering

[Click here for K-means clustering](#)

14.3.2 Hierarchical clustering

[Click here for hierarchical clustering](#)

14.4 Homework

14.4.1 Disentangling microbiomes

This is an example problem. These data are from (Couper et al. 2019).

```
# Run this whole chunk to load data
```

```
microbiome <- read_csv("https://zenodo.org/records/4940316/files/MicrobialCommunityAssembly_1  
t() |>  
as.data.frame()
```

New names:

Rows: 36 Columns: 69

-- Column specification

```
----- Delimiter: "," chr  
(69): ...1, PW001, PW005, PW012, PW013, PW014, PW015, PW016, PW017, PW01...  
i Use `spec()` to retrieve the full column specification for this data. i  
Specify the column types or set `show_col_types = FALSE` to quiet this message.  
* `` -> `...1`
```

```
colnames.x <- microbiome[1,]
```

```
colnames(microbiome) <- colnames.x
```

```
microbiome <- microbiome[-1,]
```

15 Pick the test

15.1 Picking the test

A major component of the final will be picking the correct test to run on some data. Here, I cover the specific ways in which you can determine what test to use.

Please use this web page as an *interactive* way to pick the test.

15.2 Exceptions

Below is a walk through for the most common statistical analyses, but keep in mind there are a few “less common” ones that we are using as well:

- **Binomial test** - if we are looking at something with discrete outcomes - like coin tosses, die rolls, etc. - we are doing a binomial test to determine the probability of a specific outcome. You can do this with `binom.test`.
- **Poisson test** - if we are looking at the probability of obtaining certain counts over events - specifically, looking at the probability of *rare* events - we will use a Poisson. You can do this with `poisson.test`.

15.3 Overview of picking the test

Below is an overview of picking basic statistical tests. There are more complex tests, but these are the main ones to consider for exams in this course.

Table 15.1: Which test to pick for which combinations of variables.

| | Categorical Explanatory | Continuous Explanatory |
|------------------------------|---|--|
| Categorical Re-sponse | <ul style="list-style-type: none"> • χ^2 tests, especially if count data for category <ul style="list-style-type: none"> – Yate’s correction auto-applied for 2x2 tables • Fisher test is a 2x2 table • Null is that counts match a known proportion (goodness-of-fit) or counts match each other (independence) | <ul style="list-style-type: none"> • Logistic regression, modeling categorical responses across continuous variables. – Not covered in BIOL 305 at present |

| | Categorical Explanatory | Continuous Explanatory |
|----------------------|---|--|
| Continuous Re-sponse | <ul style="list-style-type: none"> • Single mean to population <ul style="list-style-type: none"> – This is a Z-score comparison, when the population parameters are known <ul style="list-style-type: none"> * $H_0 : \bar{x} = \mu$ – Almost always better to resort to t-test comparison if unsure is a population - t approaches a Z as $df \rightarrow \infty$ <ul style="list-style-type: none"> * $H_0 : \mu_1 = \mu_2$ • Two measurements <ul style="list-style-type: none"> – Remember - all t-tests default to Welch's in R, assuming unequal variance <ul style="list-style-type: none"> * When $\sigma_1^2 = \sigma_2^2$, Welch's t-test is the same as a standard t-test * Usually better to use Welch's, but justify reasoning – One-sample t-test, comparing the two groups <ul style="list-style-type: none"> * $H_0 : \mu_1 = \mu_2$ • Two samples <ul style="list-style-type: none"> – Two-sample t-test, comparing the two samples <ul style="list-style-type: none"> * $H_0 : \mu_1 = \mu_2$ • Paired samples <ul style="list-style-type: none"> – Special case - used when there are repeated measurements of the same sampling units <ul style="list-style-type: none"> * Need to set <code>paired = TRUE</code> * $H_0 : \mu_d = 0$ • ≥ 3 samples <ul style="list-style-type: none"> – ANOVA - analysis of variance across multiple samples, computing variance within samples and between samples <ul style="list-style-type: none"> * For all ANOVA, $H_0 : \mu_1 = \mu_2 = \dots \mu_i$ * Single type of measurement is a one-way ANOVA * When things are “blocked” into categories - like litters - that are unique per row, use a randomized block ANOVA (remember to put <code>factor()</code> around anything you are worried won't be read as categorical!) <ul style="list-style-type: none"> · <code>aov(response ~ explanatory + block, data)</code> | <ul style="list-style-type: none"> • When looking at <i>if</i> there is a relationship, use correlation <ul style="list-style-type: none"> – $H_0 : \rho = 0$ – Evaluated using a t statistic • When looking at <i>what</i> the relationship is, use linear regression <ul style="list-style-type: none"> – Line formula is $y = \beta x + \alpha + \epsilon_i$ – Very similar to your high-school $y = mx + b$ – $H_0 : \beta = 0$ • Remember - linear regression works like an ANOVA, and has similar outputs <ul style="list-style-type: none"> – Check the ANOVA page for more information! |
| | | |

15.4 Another method - checklist

Feel free to go through the following headings to also help you pick a test.

15.4.1 Explanatory Variable

The explanatory variable is continuous and numeric. Pick this one for a continuous measurement variable, for example, such as Longitude, concentration, or other ratio or integer data.

The explanatory variable is discontinuous and categorical. In these cases, the explanatory variable is a condition, like a control and a treatment.

15.4.2 Continuous explanatory variable

15.4.2.1 Continuous response variable

If you have a continuous response variable, then you need to see what the question is asking:

- **Is there a relationship?**

If you are looking at a problem and it is simply asking *if* there is a relationship, you are looking at a **correlation analysis**.

- **What is the relationship?**

If you are asking *what* the relationship is or looking to be able to *predict* a value based on what you know, you are doing a **linear regression**. Note there are other kinds of regression, but in this class, we focus on linear regression.

If we have multiple response variables, we can do multiple regression, which we do not cover here.

15.4.2.2 Discrete response variable

If you are looking at a discrete response variable, such as a state of 1 or 0 in response to a certain amount of stimulus, then you are doing a **logistic regression**. We did not cover this analysis in this class.

15.4.3 Discrete explanatory variable

For discrete explanatory variables, we are often looking at categorical treatments or distinct groups, like species or geographic locations.

15.4.3.1 Continuous response variable

For a continuous response variable, we need to ask ourselves how many categories we are dealing with.

- If we are dealing with **two categories** or **two measurements from the same individual**, we are using a ***t*-test**.

Make sure you check the *t*-test page to understand what kind of *t*-test is being performed. For repeated measurements from the same individuals or populations under different conditions, we have the **paired *t*-test**. Otherwise, we have the **Welch's *t*-test** as the default in *R*. *NOTE* that the default assumes unequal variance; you must set `var.equal = TRUE` to perform a “**true**” *t*-test. Make sure you familiarize yourself with *why* our code is assuming variances are unequal.

If we *know* what the population is, we will use a *Z*-score, but bear in mind that we will almost always use a *t*-test to account for error. A *t*-test with infinite degrees of freedom is the same as a *Z*-score, so it is better to default to a *t*-test.

- If we have **three or more categories or treatments** then we need to perform an **ANOVA**.

If we are simply comparing multiple groups, we are performing a one-way ANOVA. We also need to make sure we aren't doing some sort of factorial ANOVA, repeated-measures ANOVA, or interactive ANOVA. Please read the ANOVA page to ensure you are using the correct format.

Remember to label ANOVA plots with letters to indicate the separate groups.

If we have multiple response variables, we can use a MANOVA; we do not cover that in this class.

15.4.3.2 Discrete response variable

If our response variable *and* our explanatory variable are discrete (i.e., categorical or nominal), then we are doing a χ^2 test. This tests looks at counts in different categories. For example, looking at proportions of men and women who do and do not smoke would be a classic χ^2 test.

Other Online Resources

<https://statkat.com/statistical-technique-selection/tool-for-selecting-a-statistical-technique-step-one.php>

16 Conclusions

16.1 Parting thoughts

In this class, we have covered two major things: (1) the basics of statistics for biological research and (2) the basics of using *R* to solve different computational problems. It is my hope that this class helps you both become a better researcher and also a more efficient researcher and student by using code to help you with your future projects.

16.2

(pronounced *doh-dah-dah-go-huh-ee*) is a traditional Cherokee farewell. It does not mean goodbye, but rather reflects a parting of ways until a group of folks meet again.

I enjoyed getting to know all of you in class, and please feel free to reach out or stop by and say hi if you are ever passing through Kearney in the future or if you need help with something biology related.

17 Functions & Glossary

17.1 Common Commands

The following are common useful commands used in *R*, with examples of their use.

Commands are given as Windows & Linux / Mac.

17.1.1 The basics

- Save your file with **CTRL + S** / **⌘ + S**
- Undo your last action with **CTRL + Z** / **⌘ + Z**
- Cut text with **CTRL + X** / **⌘ + X**
- Copy text with **CTRL + C** / **⌘ + C**
- Paste text with **CTRL + V** / **⌘ + V**
- Select all text in a document with **CTRL + A** / **⌘ + A**

17.1.2 Code basics

- A code chunk can be inserted using **CTRL + SHIFT + I** / **⌘ + shift + I**
- Run a line of code with **CTRL + ENTER** / **⌘ + return**
- Run a code chunk with **CTRL + SHIFT + ENTER** / **⌘ + shift + return**
- Run all code in the document above where you were **CTRL + SHIFT + ALT + P** / **⌘ + shift + ⌘ + P**
- Run all code with **CTRL + ALT + R** / **⌘ + ⌘ + R**

17.1.3 Code formatting

- `<-` - save a value as an object. On Mac, the keyboard shortcut is `⌘ + -`. Windows can be formatted so that the shortcut is `ALT + -`.

```
x <- 10
x
```

```
[1] 10
```

- `|>` - “pipe” a command or output into another command. You can use the shortcut `CTRL+SHIFT+M` / `⌘ + SHIFT + M`.

```
# make repeatable
set.seed(930)

# random string
x <- rnorm(20)

x |>
  # pass to summary
  summary() |>
  # pass summary through round
  round(2)
```

| Min. | 1st Qu. | Median | Mean | 3rd Qu. | Max. |
|-------|---------|--------|-------|---------|------|
| -1.94 | -0.48 | -0.06 | -0.05 | 0.36 | 1.80 |

- `->` - save a value as an object at the end of a pipeline

```
# same pipeline as previous, but more sequential

# get random values
rnorm(20) |>
  # pass to summary
  summary() |>
  # pass summary through round
  round(2) ->
  # save as x
  x

x
```

| Min. | 1st Qu. | Median | Mean | 3rd Qu. | Max. |
|-------|---------|--------|-------|---------|------|
| -2.60 | -0.77 | -0.24 | -0.18 | 0.57 | 1.73 |

- `c` - concatenate, place two values together

```
x <- c(10,11)
x
```

```
[1] 10 11
```

17.1.4 Boolean logic

In order to create arguments in *R*, such as commands like **which**, we need to use Boolean logic commands.

For this example, I will create a string from 1 to 10 to perform the commands on. *Note* that I am using brackets `[]` to call out an index within the object using these commands.

These are as follows:

```
# example string
x <- 1:10
```

- `>`: greater than

```
x[x > 5]
```

```
[1] 6 7 8 9 10
```

- `<`: less than

```
x[x < 5]
```

```
[1] 1 2 3 4
```

- `>=`: greater than or equal to

```
x[x >= 5]
```

```
[1] 5 6 7 8 9 10
```

- `<=`: less than or equal to

```
x[x <= 5]
```

```
[1] 1 2 3 4 5
```

- ==: is equal to

```
x[x == 5]
```

```
[1] 5
```

- |: or

```
x[x < 3 | x > 7]
```

```
[1] 1 2 8 9 10
```

- &: and

```
x[x > 3 & x < 7]
```

```
[1] 4 5 6
```

17.2 Basic statistics

[Click here for basic statistics](#)

For these examples, we will create a random vector of number to demonstrate how they work.

```
# make repeatable  
set.seed(8675309)  
# get random normal values  
x <- rnorm(1000)
```

- mean - get the mean / average of a set of data

```
mean(x)
```

```
[1] -0.02659863
```

```
median(x)
```

```
[1] -0.0006443188
```

```
IQR(x)
```

```
[1] 1.378508
```

```
range(x)
```

```
[1] -3.449557  3.412794
```


```
# calculate range  
abs(range(x)[1] - range(x)[2])
```

```
[1] 6.862351
```

```
# calculate range  
max(x) - min(x)
```

```
[1] 6.862351
```

```
hist(x)
```



17.3 Null hypotheses

[Click here for null hypotheses](#)

17.3.1 χ^2 tests

17.3.1.1 χ^2 goodness-of-fit

H_0 : The frequency of observations fits the expected observation frequencies.

17.3.1.2 χ^2 test of independence (i.e., association)

H_0 : Variable X and Variable Y are not related.

17.3.2 One-sample tests

17.3.2.1 One-sample t -test

H_0 : $\mu = X$

17.3.2.2 Paired t -test

H_0 : $\mu_1 - \mu_2 = 0$ or $\mu_d = 0$

17.3.2.3 Two-sample t -test

H_0 : $\mu_1 = \mu_2$

17.3.3 Multi-sample tests

All ANOVA tests have essentially the same null hypotheses, with the primary null hypothesis being that all means are the same. Secondary hypotheses may exist about *which* means are different; we have to do a post-hoc test to determine which means are different.

H_0 : $\mu_1 = \mu_2 = \mu_3 = \dots = \mu_n$

17.3.4 Correlation

For all correlation tests, the null is that the correlation is 0 (i.e., not correlated positively or negatively). You can get the p value for a test from a t distribution.

As an example, Pearson's correlation:

$$H_0: \rho = 0$$

17.3.5 Regression

$$H_0: \beta = 0$$

where β is the slope.

17.4 Custom functions from class and elsewhere

The following functions are those developed for this class or adapted from code posted to sources like StackOverflow.

For these stats, we are using the following example data:

```
### EXAMPLE DATA ###  
x <- c(1,2,3,5,5,4,6,8,7,9,5)
```

17.4.0.1 Mode calculations

[Click here for mode](#)

```
# Based on Statology function  
# define function to calculate mode  
# works on vectors of data  
find_mode <- function(x) {  
  # get unique values from vector  
  u <- unique(x)  
  # count number of occurrences for each value  
  tab <- tabulate(match(x, u))  
  
  # if no mode, say so  
  if(length(x)==length(u[tab == max(tab)])){  
    print("No mode.")  
  }else{
```

```

    # return the value with the highest count
    u[tab == max(tab)]
  }
}

find_mode(x)

```

```
[1] 5
```

17.4.0.2 Standard error

```

se <- function(x){
  n <- length(x) # calculate n
  s <- sd(x) # calculate standard deviation
  se_val <- s/sqrt(n)
  return(se_val)
}

se(x)

```

```
[1] 0.7385489
```

17.4.0.3 Coefficient of variation

Click [here](#) for coefficient of variation

```

cv <- function(x){
  sigma <- sd(x)
  mu <- mean(x)
  val <- sigma/mu*100
  return(val)
}

cv(x)

```

```
[1] 48.98979
```

17.4.1 Normal distributions

17.4.1.1 Z score

[Click here for Z-score](#)

Remember - in the Z-score code below, if no n is specified, then it will default to $n = 1$.

```
zscore <- function(xbar, mu, sd.x, n = 1){  
  z <- (xbar - mu)/(sd.x/sqrt(n))  
  return(z)  
}  
  
zscore(xbar = 62,  
       mu = 65,  
       sd.x = 3.5,  
       n = 5)
```

```
[1] -1.91663
```

17.4.2 ANOVA

[Click here for ANOVA functions](#)

The following example data are going to be used to illustrate these functions.

```
#### EXAMPLE DATA ####  
set.seed(8675309)  
  
for(i in 1:4){  
  x <- rnorm(10)  
  if(i == 1){  
    x <- rnorm(10, mean = 2)  
    data <- x |> as.data.frame()  
    colnames(data) <- "Response"  
    data$Explanatory <- paste0("x",i)  
  }else{  
    newdat <- x |> as.data.frame()  
    colnames(newdat) <- "Response"  
    newdat$Explanatory <- paste0("x",i)  
    data <- rbind(data,newdat)  
  }  
}
```

```

}

# split into "typical" table
expanded_data <- NULL
expanded_data$x1 <- data$Response[which(data$Explanatory=="x1")]
expanded_data$x2 <- data$Response[which(data$Explanatory=="x2")]
expanded_data$x3 <- data$Response[which(data$Explanatory=="x3")]
expanded_data$x4 <- data$Response[which(data$Explanatory=="x4")]

expanded_data <- expanded_data |>
  as.data.frame()

```

The above is a *one-way ANOVA*. As a reminder, we would calculate the test as follows:

```

# pivot longer does not work for one-way ANOVA, requires blocking factor
# can rbind things with same colnames to make longer

example_aov <- aov(Response ~ Explanatory, data)

summary(example_aov)

```

```

              Df Sum Sq Mean Sq F value    Pr(>F)
Explanatory    3  23.40    7.801    11.54 1.89e-05 ***
Residuals     36   24.33    0.676
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Above, we can see a significant result of the ANOVA. We can follow this up with a Tukey test. **This requires the package agricolae!** Check the ANOVA pages, however, as not all ANOVA can use this agricolae shortcut method.

```

example_tukey <- HSD.test(example_aov,
  # what to group by?
  "Explanatory",
  # significance level?
  alpha = 0.05,
  # are data unbalanced
  unbalanced = FALSE,
  # show answer?
  console = TRUE)

```

```
Study: example_aov ~ "Explanatory"
```

```
HSD Test for Response
```

```
Mean Square Error: 0.6758192
```

```
Explanatory, means
```

| | Response | std | r | se | Min | Max | Q25 | Q50 |
|----|------------|-----------|----|-----------|------------|----------|------------|-------------|
| x1 | 1.7620153 | 1.0505466 | 10 | 0.2599652 | 0.4504476 | 3.972459 | 1.1175485 | 1.44911720 |
| x2 | 0.2841495 | 0.7532422 | 10 | 0.2599652 | -0.4729986 | 1.985826 | -0.3497379 | 0.21347543 |
| x3 | 0.2197337 | 0.7019368 | 10 | 0.2599652 | -0.6150452 | 1.574903 | -0.2436023 | -0.04493909 |
| x4 | -0.2890524 | 0.7345336 | 10 | 0.2599652 | -1.9769014 | 0.684072 | -0.5394534 | -0.07741642 |
| | Q75 | | | | | | | |
| x1 | 2.07491533 | | | | | | | |
| x2 | 0.64579865 | | | | | | | |
| x3 | 0.53544151 | | | | | | | |
| x4 | 0.04323417 | | | | | | | |

```
Alpha: 0.05 ; DF Error: 36
```

```
Critical Value of Studentized Range: 3.808798
```

```
Minimun Significant Difference: 0.9901551
```

```
Treatments with the same letter are not significantly different.
```

| | Response | groups |
|----|------------|--------|
| x1 | 1.7620153 | a |
| x2 | 0.2841495 | b |
| x3 | 0.2197337 | b |
| x4 | -0.2890524 | b |

17.4.2.1 Summarize data (for plotting)

Remember - you need to change "Explanatory" to your explanatory variable (in quotes!) and you need to change **Response** to your response column (no quotes!). The following requires `plyr` to work, but the function itself should call up `plyr` if you do not yet have it loaded.

```
# summarize by group
summary_data <- function(data, explanatory){
  require(plyr)
```

```

ddply(data, paste(explanatory), summarise,
      N = length(Response),
      mean = mean(Response),
      sd = sd(Response),
      se = sd / sqrt(N))
}

example_summary <- summary_data(data = data, explanatory = "Explanatory")

```

Loading required package: plyr

You have loaded plyr after dplyr - this is likely to cause problems.
 If you need functions from both plyr and dplyr, please load plyr first, then dplyr:
 library(plyr); library(dplyr)

Attaching package: 'plyr'

The following objects are masked from 'package:dplyr':

```

  arrange, count, desc, failwith, id, mutate, rename, summarise,
  summarize

```

The following object is masked from 'package:purrr':

```

  compact

```

```

example_summary

```

| | Explanatory | N | mean | sd | se |
|---|-------------|----|------------|-----------|-----------|
| 1 | x1 | 10 | 1.7620153 | 1.0505466 | 0.3322120 |
| 2 | x2 | 10 | 0.2841495 | 0.7532422 | 0.2381961 |
| 3 | x3 | 10 | 0.2197337 | 0.7019368 | 0.2219719 |
| 4 | x4 | 10 | -0.2890524 | 0.7345336 | 0.2322799 |

17.4.2.2 Significant label maker

This command requires a Tukey HSD object from **agricolae**. You can manually create a table like this for some other scenarios; see relevant pages for documentation.

```
# note first group must be EXACT MATCH to your summary_data object
# groups are saved in the Tukey object
# this is true for Tukey later as well

# the following is a function that will make the significant label table
sig.label.maker <- function(tukey_test, group_name){
  sig.labels <- tukey_test$groups |>
    # convert to a data.frame
    as.data.frame() |>
    # create a new column - place rownames into the column
    # converts to a format better for ggplot
    mutate(Explanatorys = rownames(tukey_test$groups)) |>
    # rename column to prevent confusion
    # specify dplyr; default function may be from plyr and not work
    dplyr::rename(Significance = groups)
  colnames(sig.labels)[which(colnames(sig.labels) == "Explanatorys")] <- group_name
  return(sig.labels)
}

# Function requires explanatory groups in quotes
sig_labels <- sig.label.maker(example_tukey, "Explanatory")

sig_labels
```

| | Response | Significance | Explanatory |
|----|------------|--------------|-------------|
| x1 | 1.7620153 | a | x1 |
| x2 | 0.2841495 | b | x2 |
| x3 | 0.2197337 | b | x3 |
| x4 | -0.2890524 | b | x4 |

17.4.2.3 ANOVA plotter

The following function plots ANOVAS if you have a **summary_data** object and a **sig_labels** object, as shown above. **This does not work on ANOVA with interactive components.**


```

anova_plotter <- function(summary_data, explanatory,
                           response, sig_labels,
                           y_lim=NA, label_height=NA,
                           y_lab=NA, x_lab=NA){

  require(tidyverse)

  plot_data_1 <- summary_data[,c(explanatory, response, "se")]
  plot_data_2 <- sig_labels[,c(explanatory,"Significance")]

  colnames(plot_data_1) <- c("explanatory", "response", "se")
  colnames(plot_data_2) <- c("explanatory", "Significance")

  plot_data <- plot_data_1 |>
    full_join(plot_data_2, by = "explanatory")

  if(is.na(y_lim)){
    if(min(plot_data$response) < 0){
      y_lim <- c(min(plot_data$response) -
        4*max(plot_data$se),
        max(plot_data$response) +
        4*max(plot_data$se))
    }else{
      y_lim <- c(0,max(plot_data$response) +
        4*max(plot_data$se))
    }
  }
  if(is.na(label_height)){label_height <- 0.25*max(y_lim)}
  if(is.na(y_lab)){y_lab <- "Response"}
  if(is.na(x_lab)){x_lab <- "Treatment"}

  plot_1 <- ggplot(plot_data,
    aes(x = explanatory, y = response)) +
    geom_point() +
    geom_errorbar(data = plot_data,
      aes(ymin = response - 2*se,
        ymax = response + 2*se,
        width = 0.1)) +
    ylim(y_lim) +
    theme_classic() +
    theme(legend.position = "none",
      axis.text.x = element_text(angle = 90, vjust = 0.5, size = 5)) +
    geom_text(data = plot_data,

```


```

# make bold
fontface = "bold",
# define where labels should go
aes(x = explanatory,
     # define height of label
     y = label_height,
     # what are the labels?
     label = paste0(Significance))) +
xlab(x_lab) +
ylab(y_lab)

print(plot_1)
}


anova_plotter(summary_data = example_summary,
               explanatory = "Explanatory",
               response = "mean", # from summary_data table! What is to be plotted
               sig_labels = sig_labels)

```



Note that in the above, the default label height is not working for us. We can adjust this with `label_height`.

```
anova_plotter(summary_data = example_summary,
               explanatory = "Explanatory",
               response = "mean", # from summary_data table! What is to be plotted
               sig_labels = sig_labels,
               label_height = -1)
```



Much better!

References

- Brambilla, A., A. von Hardenberg, B. Bassano, L. Ranghetti, L. Keller, and M. Festa-Bianchet (2024). Data from: Climate drives body mass changes in a mountain ungulate: Shorter winters lead to heavier alpine ibex. [Online.] Available at <https://doi.org/10.5061/dryad.w9ghx3fz6>.
- Comont, R. (2020). BeeWalk dataset 2008-23. <https://doi.org/10.6084/m9.figshare.12280547.v4>
- Cooper, J. C. (2021). Biogeographic and Ecologic Drivers of Avian Diversity. [Online.] Available at <https://doi.org/10.6082/uchicago.3379>.
- Cooper, J. C., J. D. Maddox, K. McKague, and J. M. Bates (2021b). Data from: Multiple lines of evidence indicate ongoing allopatric and parapatric diversification in an Afromontane sunbird (*Cinnyris reichenowi*). <https://doi.org/10.5061/dryad.34tmpg4j0>
- Cooper, J. C., J. D. Maddox, K. McKague, and J. M. Bates (2021a). [Multiple lines of evidence indicate ongoing allopatric and parapatric diversification in an Afromontane sunbird \(*Cinnyris reichenowi*\)](#). Ornithology 138:ukaa081.
- Couper, L. I., J. Y. Kwan, J. Ma, and A. Swei (2019). Data from: Drivers and patterns of microbial community assembly in a lyme disease vector. [Online.] Available at <https://doi.org/10.5061/dryad.2nv32qh>.
- Courtenay, L. (2019). Measurements on Canid Tooth Scores. <https://doi.org/10.6084/m9.figshare.8081108.v1>
- Lele, A., H. Garrod, E. Ferguson, C. Azahara Prieto Gil, and M. Ellis (2022). Morphological measurements in a coastal Ecuadorian avifauna. [Online.] Available at <https://zenodo.org/records/6511860>.
- Lydeamore, M. J., P. T. Campbell, D. J. Price, Y. Wu, A. J. Marcato, W. Cuningham, J. R. Carapetis, R. M. Andrews, M. I. McDonald, J. McVernon, S. Y. C. Tong, and J. M. McCaw (2020b). Patient ages at presentation. <https://doi.org/10.1371/journal.pcbi.1007838.s006>
- Lydeamore, M. J., P. T. Campbell, D. J. Price, Y. Wu, A. J. Marcato, W. Cuningham, J. R. Carapetis, R. M. Andrews, M. I. McDonald, J. McVernon, S. Y. C. Tong, and J. M. McCaw (2020a). [Estimation of the force of infection and infectious period of skin sores](#)

- in remote Australian communities using interval-censored data. PLOS Computational Biology 16:e1007838.
- Moura, R., N. P. Santos, and A. Rocha (2023). Processed csv file of the piracy dataset. <https://doi.org/10.6084/m9.figshare.24119643.v1>
- Ruxton, G. D. (2006). The unequal variance *t*-test is an underused alternative to Student's *t*-test and the Mann–Whitney *u* test. Behavioral Ecology 17:688–690.
- Urban Big Data Centre (2025). Counts hourly pay Adzuna jobs. [Online.] Available at <https://zenodo.org/records/14771706>.
- Verstraete, H., W. Courtens, R. Daelemans, M. Van de walle, N. Vanermen, P. Desmet, and E. W. M. Stienen (2020). Photos and measurements of otoliths from fish caught in the southern part of the North Sea. [Online.] Available at <https://zenodo.org/records/4066594>.
- Weber, G. (2019). Fremont Bridge Hourly Bicycle Counts by Month October 2012 to present. [Online.] Available at <https://zenodo.org/records/2648564>.